

Exam 1 Review

CSE 101

Winter 2023

Exam Details

- In class
- Randomized assigned seats
- You may use 6 one-sided pages of notes
- No textbook or electronic aids
- No need to provide proofs unless asked for
- 3 Questions in 45 minutes
 - 1st straightforward implementation of algorithm
 - 2nd requires some thought
 - 3rd can be quite tricky

This Review

- Brief outline of topics that might show up on the exam
- To see anything in more depth use other review options.

Other Review Options

- Lecture podcasts / slides
- Textbook
- OH questions
- Old exams from problem archive

Exam Topics

- Chapter 3
 - Graph basics
 - Explore/DFS
 - Connected components
 - Pre/Post orderings
 - DAGs
 - Topological sort
 - Strongly connected components
- Chapter 4
 - Shortest path definitions
 - BFS
 - Dijkstra
 - Priority queues
 - Bellman-Ford
 - Negative weight cycles
 - Shortest paths in DAGs


Graph Definition

Definition: A *graph* $G = (V, E)$ consists of two things:

- A collection V of *vertices*, or objects to be connected.
- A collection E of *edges*, each of which connects a pair of vertices.

Drawing Graphs

- Draw vertices as points
- Draw edges as line segments or curves connecting those points



$V = \{A, B, C, D, E\}$

$E = \{AB, AC, AD, BD, CE, DE\}$

Explore

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

```
      w.prev  $\leftarrow$  v
```


Result

Theorem: If all vertices start unvisited, `explore(v)` marks as visited exactly the vertices reachable from `v`.

Depth First Search

`explore` only finds the part of the graph reachable from a single vertex. If you want to discover the entire graph, you may need to run it multiple times.

```
DepthFirstSearch(G)
```

```
    Mark all  $v \in G$  as unvisited
```

```
    For  $v \in G$ 
```

```
        If not  $v.visited$ , explore(v)
```

Runtime of DFS

```
explore(v)
  v.visited ← true
  For each edge (v,w)
    If not w.visited
      explore(w)
```

Run once
per vertex $O(|V|)$
total

Run once per
neighboring
vertex $O(|E|)$
total

```
DFS(G)
  Mark all  $v \in G$  as unvisited
  For  $v \in G$ 
    If not v.visited, explore(v)
```

$O(|V|)$

Final runtime:
 $O(|V| + |E|)$

Connected Components

Theorem: The vertices of a graph G can be partitioned into *connected components* so that v is reachable from w if and only if they are in the same connected component.

Computing CCs with DFS

```
ConnectedComponents (G)
```

```
  CCNum  $\leftarrow$  0
```

```
  For  $v \in G$ 
```

```
     $v$ .visited  $\leftarrow$  false
```

```
  For  $v \in G$ 
```

```
    If not  $v$ .visited
```

```
      CCNum++
```

```
      explore( $v$ )
```

```
explore( $v$ )
```

```
   $v$ .visited  $\leftarrow$  true
```

```
   $v$ .CC  $\leftarrow$  CCNum
```

```
  For each edge ( $v, w$ )
```

```
    If not  $w$ .visited
```

```
      explore( $w$ )
```

Runtime $O(|V|+|E|)$.

Pre- and Post- Orders

- Keep track of what DFS does & in what order.
- Have a “clock” and note time whenever:
 - Algorithm visits a new vertex for the first time.
 - Algorithm finishes processing a vertex.
- Record values as $v.pre$ and $v.post$.

Computing Pre- & Post- Orders

```
PreAndPost (G)
```

```
  clock ← 1
```

```
  For v ∈ G
```

```
    v.visited ← false
```

```
  For v ∈ G
```

```
    If not v.visited
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited ← true
```

```
  v.pre ← clock
```

```
  clock++
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

```
  v.post ← clock
```

```
  clock++
```

Runtime $O(|V|+|E|)$.

What do these orders tell us?

Prop: For vertices v, w consider intervals $[v.pre, v.post]$ and $[w.pre, w.post]$. These intervals:


1. Contain each other if v is an ancestor/descendant of w in the DFS tree.
2. Are disjoint if v and w are cousins in the DFS tree.
3. Never interleave
($v.pre < w.pre < v.post < w.post$)

Directed Graphs

Often an edge makes sense both ways, but sometimes streets are one directional.

Definition: A directed graph is a graph where each edge has a direction. Goes *from* v to w .

Draw edges with arrows to denote direction.




Dependency Graphs

Definition: A topological ordering of a directed graph is an ordering of the vertices so that for each edge (v,w) , v comes before w in the ordering.

Cycles

Definition: A cycle in a directed graph is a sequence of vertices $v_1, v_2, v_3, \dots, v_n$ so that there are edges $(v_1, v_2), (v_2, v_3), \dots, (v_{n-1}, v_n), (v_n, v_1)$



DAGs

Definition: A Directed Acyclic Graph (DAG) is a directed graph which contains no cycles.

Theorem: A directed graph G has a topological ordering if and only if it is a DAG.

Topological Sort

TopologicalSort (G)

Run DFS (G) w/ pre/post numbers

Return the vertices in *reverse*
postorder

Runtime: $O(|V|+|E|)$.

Correctness

Proposition: If G is a DAG with an edge $v \rightarrow w$ then $w.\text{post} < v.\text{post}$.

Strongly Connected Components


Definition: In a directed graph G , two vertices v and w are in the same Strongly Connected Component (SCC) if v is reachable from w *and* w is reachable from v .

Lemma: You can actually partition the vertices into components in this way.

Metagraph

Definition: The metagraph of a directed graph G is a graph whose vertices are the SCCs of G , where there is an edge between C_1 and C_2 if and only if G has an edge between some vertex of C_1 and some vertex of C_2 .

Example



Result

Theorem: The metagraph is any directed graph is a DAG.


Problem: Given a directed graph G compute the SCCs of G and its metagraph.

Observation

If v in sink SCC, `explore(v)` finds *exactly* v 's component.



Result

Proposition: Let C_1 and C_2 be SCCs of G with an edge from C_1 to C_2 . If we run DFS on G , the largest postorder number of any vertex in C_1 will be larger than the largest postorder number in C_2 .



Reverse Graph

Definition: Given a directed graph G , the reverse graph of G (denoted G^R) is obtained by reversing the directions of all of the edges of G .



SCC Algorithm

SCCs (G)

Run DFS (G^R) record postorders

Mark all vertices unvisited

For $v \in V$ in reverse postorder

 If not $v.visited$

 explore(v) mark component

Just 2 DFSs! Runtime $O(|V|+|E|)$.

Goal

Problem: Given a graph G with two vertices s and t in the same connected component, find the *best* path from s to t .

What do we mean by best?

- Least expensive
- Best scenery
- Shortest
- For now: fewest edges

Observation

If there is a length $\leq d$ s-v path, then there is some w adjacent to v with a length $\leq (d-1)$ s-w path.

Algorithm Idea

For each d create a list of all vertices at distance d from s .

- For $d=0$, this list is just $\{s\}$.
- For larger d , we want all new vertices adjacent to vertices at distance $d-1$.

Breadth-First Search

BFS(G, s)

For $v \in V$, $\text{dist}(v) \leftarrow \infty$

Initialize Queue Q

$Q.\text{enqueue}(s)$

$\text{dist}(s) \leftarrow 0$

While (Q nonempty)

$u \leftarrow \text{front}(Q)$

 For $(u, v) \in E$

 If $\text{dist}(v) = \infty$

$\text{dist}(v) \leftarrow \text{dist}(u) + 1$

$Q.\text{enqueue}(v)$

$v.\text{prev} \leftarrow u$

$O(|V|)$

$O(|V|)$ iterations

$O(|E|)$ total iterations

Total runtime:
 $O(|V| + |E|)$

Edge Lengths

The number of edges in a path is not always the right measure of distance. Sometimes, taking several shorter steps is preferable to taking a few longer ones.

We assign each edge (u,v) a non-negative length $\ell(u,v)$. The length of a path is the sum of the lengths of its edges.

Problem: Shortest Paths

Problem: Given a Graph G with vertices s and t and a length function ℓ , find the shortest path from s to t .

Algorithm

Distances (G, s, ℓ)

$\text{dist}(s) \leftarrow 0$

While (not all distances found)

Find minimum over $(v, w) \in E$


with v discovered w not

of $\text{dist}(v) + \ell(v, w)$

$\text{dist}(w) \leftarrow \text{dist}(v) + \ell(v, w)$

$\text{prev}(w) \leftarrow v$

Why does this work?



Correctly Assigned Distances

Runtime

- This is too slow.
- **Problem:** Every iteration we have to check every edge.
- **Idea:** Most of the comparison doesn't change much iteration to iteration. Use to save time.
- **Use data structure:** In particular a priority queue.

Priority Queue

A Priority Queue is a datastructure that stores elements sorted by a key value.

Operations:

- Insert – adds a new element to the PQ.
- DecreaseKey – Changes the key of an element of the PQ to a specified *smaller* value.
- DeleteMin – Finds the element with the smallest key and removes it from the PQ.

Dijkstra's Algorithm

```
Dijkstra( $G, s, \ell$ )
```

```
  Initialize Priority Queue  $Q$ 
```

```
  For  $v \in V$ 
```

```
     $\text{dist}(v) \leftarrow \infty$ 
```

```
     $Q.\text{Insert}(v)$ 
```

```
   $\text{dist}(s) \leftarrow 0$ 
```

```
  While( $Q$  not empty)
```

```
     $v \leftarrow Q.\text{DeleteMin}()$ 
```

```
    For  $(v, w) \in E$ 
```

```
      If  $\text{dist}(v) + \ell(v, w) < \text{dist}(w)$ 
```

```
         $\text{dist}(w) \leftarrow \text{dist}(v) + \ell(v, w)$ 
```

```
         $Q.\text{DecreaseKey}(w)$ 
```

$O(|V|)$ times

$O(|V|)$ times

$O(|E|)$ times

Runtime:

$O(|V|)$ Inserts +

$O(|V|)$ DelMins +

$O(|E|)$ DecKeys

Summary of Priority Queues

	Insert/DecreaseKey	DeleteMin	Dijkstra
List	$O(1)$	$O(n)$	$O(V ^2 + E)$
Binary Heap	$O(\log(n))$	$O(\log(n))$	$O(\log V (V + E))$
d -ary Heap	$O\left(\frac{\log(n)}{\log(d)}\right)$	$O\left(\frac{d \log(n)}{\log(d)}\right)$	$O\left(\frac{\log V }{\log(d)} (d V + E)\right)$
Fibonacci Heap	$O(1)^*$	$O(\log(n))^*$	$O(V \log V + E)$

Negative Edge Weights

- So far we have talked about the case of non-negative edge weights.
 - The usual case (distance & time usually cannot be negative).
 - However, if “lengths” represent other kinds of costs, sometimes they can be negative.
- Problem statement same. Find path with smallest sum of edge weights.

Negative Weight Cycles


Definition: A negative weight cycle is a cycle where the total weight of edges is negative.

- If G has a negative weight cycle, then there are probably no shortest paths.
 - Go around the cycle over and over.
- **Note:** For undirected G , a single negative weight edge gives a negative weight cycle by going back and forth on it.

Fundamental Shortest Paths Formula

For $w \neq s$,

$$\text{dist}(w) = \min_{(v,w) \in E} \text{dist}(v) + \ell(v,w).$$




- System of equations to solve for distances.
- When $\ell \geq 0$, Dijkstra gives an order to solve in.
- With $\ell < 0$, might be no solution.

Algorithm Idea

Instead of finding shortest paths (which may not exist), find shortest paths of length at most k .

For $w \neq s$,

$$\text{dist}_k(w) = \min_{(v,w) \in E} \text{dist}_{k-1}(v) + \ell(v, w).$$



Bellman-Ford

```
Bellman-Ford( $G, s, \ell$ )
```

```
   $\text{dist}_0(v) \leftarrow \infty$  for all  $v$ 
```

```
  //cant reach
```

```
   $\text{dist}_0(s) \leftarrow 0$ 
```

What value of k
do we use?

```
  For  $k = 1$  to  $n$ 
```

```
    For  $w \in V$ 
```

```
       $\text{dist}_k(w) \leftarrow \min(\text{dist}_{k-1}(v) + \ell(v, w))$ 
```

```
       $\text{dist}_k(s) \leftarrow \min(\text{dist}_k(s), 0)$ 
```

```
      // s has the trivial path
```

$O(|E|)$

Analysis

Proposition: If $n \geq |V| - 1$ and if G has no negative weight cycles, then for all v ,
 $\text{dist}(v) = \text{dist}_n(v)$.

- If there is a negative weight cycle, there probably is no shortest path.
- If not, we only need to run our algorithm for $|V|$ rounds, for a final runtime $O(|V| |E|)$.

Detecting Negative Cycles

If there are no negative weight cycles, Bellman-Ford computes shortest paths (and they might not exist otherwise).

How do we know whether or not there are any?

Cycle Detection

Proposition: For any $n \geq |V| - 1$, there are no negative weight cycles reachable from s if and only if for every $v \in V$

$$\text{dist}_n(v) = \text{dist}_{n+1}(v)$$

- Detect by running one more round of Bellman-Ford.
- Need to see if *any* v 's distance changes.

Shortest Paths in DAGs

Runtime
 $O(|V| + |E|)$

ShortestPathsInDAGs (G, s, ℓ)

TopologicalSort (G) } $O(|V| + |E|)$

For $w \in V$ in topological order

If $w = s$, $\text{dist}(w) \leftarrow 0$

Else

$O(|E|)$ total

$O(|V|)$ $\text{dist}(w) \leftarrow \min(\text{dist}(v) + \ell(v, w))$ }
total

$\backslash \backslash$ $\text{dist}(v)$ for all upstream v
already computed

Shortest Path Algorithms Summary

Unit Weights: Breadth First Search
 $O(|V| + |E|)$

Non-negative Weights: Dijkstra
 $O(|V| \log |V| + |E|)$

Arbitrary Weights: Bellman-Ford $O(|V| |E|)$

Arbitrary Weights, graph is a DAG:

Shortest-Paths-In-DAGs $O(|V| + |E|)$