

Announcements

- Homework 0 Due today
- Homework 1 online due next Friday
- Discussion section notes/podcast online
- Remember FinAid survey
- No class on Monday

Last Time

- Graphs
- explore/DFS
- Connected components

Graph Definition

Definition: A *graph* $G = (V, E)$ consists of two things:

- A collection V of *vertices*, or objects to be connected.
- A collection E of *edges*, each of which connects a pair of vertices.

Runtime of DFS

```
explore(v)
```

```
  v.visited ← true
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

Run once per
vertex

$O(|V|)$ total

Run once per
neighboring
vertex

$O(|E|)$ total

```
DFS(G)
```

```
  Mark all  $v \in G$  as unvisited
```

```
  For  $v \in G$ 
```

```
    If not v.visited, explore(v)
```

$O(|V|)$

Final runtime: $O(|V| + |E|)$

Runtime of DFS

```
explore(v)
```

```
  v.visited ← true
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

Run once per
vertex

$O(|V|)$ total

Run once per
neighboring
vertex

$O(|E|)$ total

Only Counting work from
Non-recursive calls!

```
DFS(G)
```

```
  Mark all  $v \in G$  as unvisited
```

```
  For  $v \in G$ 
```

```
    If not v.visited, explore(v)
```

$O(|V|)$

Final runtime: $O(|V| + |E|)$

Connected Components

Theorem: The vertices of a graph G can be partitioned into *connected components* so that v is reachable from w if and only if they are in the same connected component.

Today

- Computing connected components
- Pre- / Post- orders
- Directed graphs
- Topological orderings

Problem: Computing Connected Components

Problem: Given a graph G , compute its connected components.

Problem: Computing Connected Components

Problem: Given a graph G , compute its connected components.

Easy: For each v , run `explore(v)` to find vertices reachable from it. Group together into components.

Problem: Computing Connected Components

Problem: Given a graph G , compute its connected components.

Easy: For each v , run `explore(v)` to find vertices reachable from it. Group together into components.

Runtime: $O(|V|(|V| + |E|))$.

Problem: Computing Connected Components

Problem: Given a graph G , compute its connected components.

Easy: For each v , run $\text{explore}(v)$ to find vertices reachable from it. Group together into components.

Runtime: $O(|V|(|V|+|E|))$.

Better: Run $\text{explore}(v)$ to find the component of v . Repeat on unclassified vertices.

DFS lets us do this!

```
ConnectedComponents (G)
```

```
  For  $v \in G$ 
```

```
     $v.\text{visited} \leftarrow \text{false}$ 
```

```
  For  $v \in G$ 
```

```
    If not  $v.\text{visited}$ 
```

```
      explore( $v$ )
```

```
explore( $v$ )
```

```
   $v.\text{visited} \leftarrow \text{true}$ 
```

```
  For each edge  $(v, w)$ 
```

```
    If not  $w.\text{visited}$ 
```

```
      explore( $w$ )
```

DFS lets us do this!

```
ConnectedComponents (G)
```

```
  CCNum  $\leftarrow$  0
```

```
  For  $v \in G$ 
```

```
     $v$ .visited  $\leftarrow$  false
```

```
  For  $v \in G$ 
```

```
    If not  $v$ .visited
```

```
      CCNum++
```

```
      explore( $v$ )
```

```
explore( $v$ )
```

```
   $v$ .visited  $\leftarrow$  true
```

```
  For each edge ( $v, w$ )
```

```
    If not  $w$ .visited
```

```
      explore( $w$ )
```

DFS lets us do this!

```
ConnectedComponents (G)
```

```
  CCNum  $\leftarrow$  0
```

```
  For  $v \in G$ 
```

```
     $v$ .visited  $\leftarrow$  false
```

```
  For  $v \in G$ 
```

```
    If not  $v$ .visited
```

```
      CCNum++
```

```
      explore( $v$ )
```

```
explore( $v$ )
```

```
   $v$ .visited  $\leftarrow$  true
```

```
   $v$ .CC  $\leftarrow$  CCNum
```

```
  For each edge ( $v, w$ )
```

```
    If not  $w$ .visited
```

```
      explore( $w$ )
```

DFS lets us do this!

```
ConnectedComponents (G)
```

```
  CCNum  $\leftarrow$  0
```

```
  For v  $\in$  G
```

```
    v.visited  $\leftarrow$  false
```

```
  For v  $\in$  G
```

```
    If not v.visited
```

```
      CCNum++
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  v.CC  $\leftarrow$  CCNum
```

```
  For each edge (v,w)
```

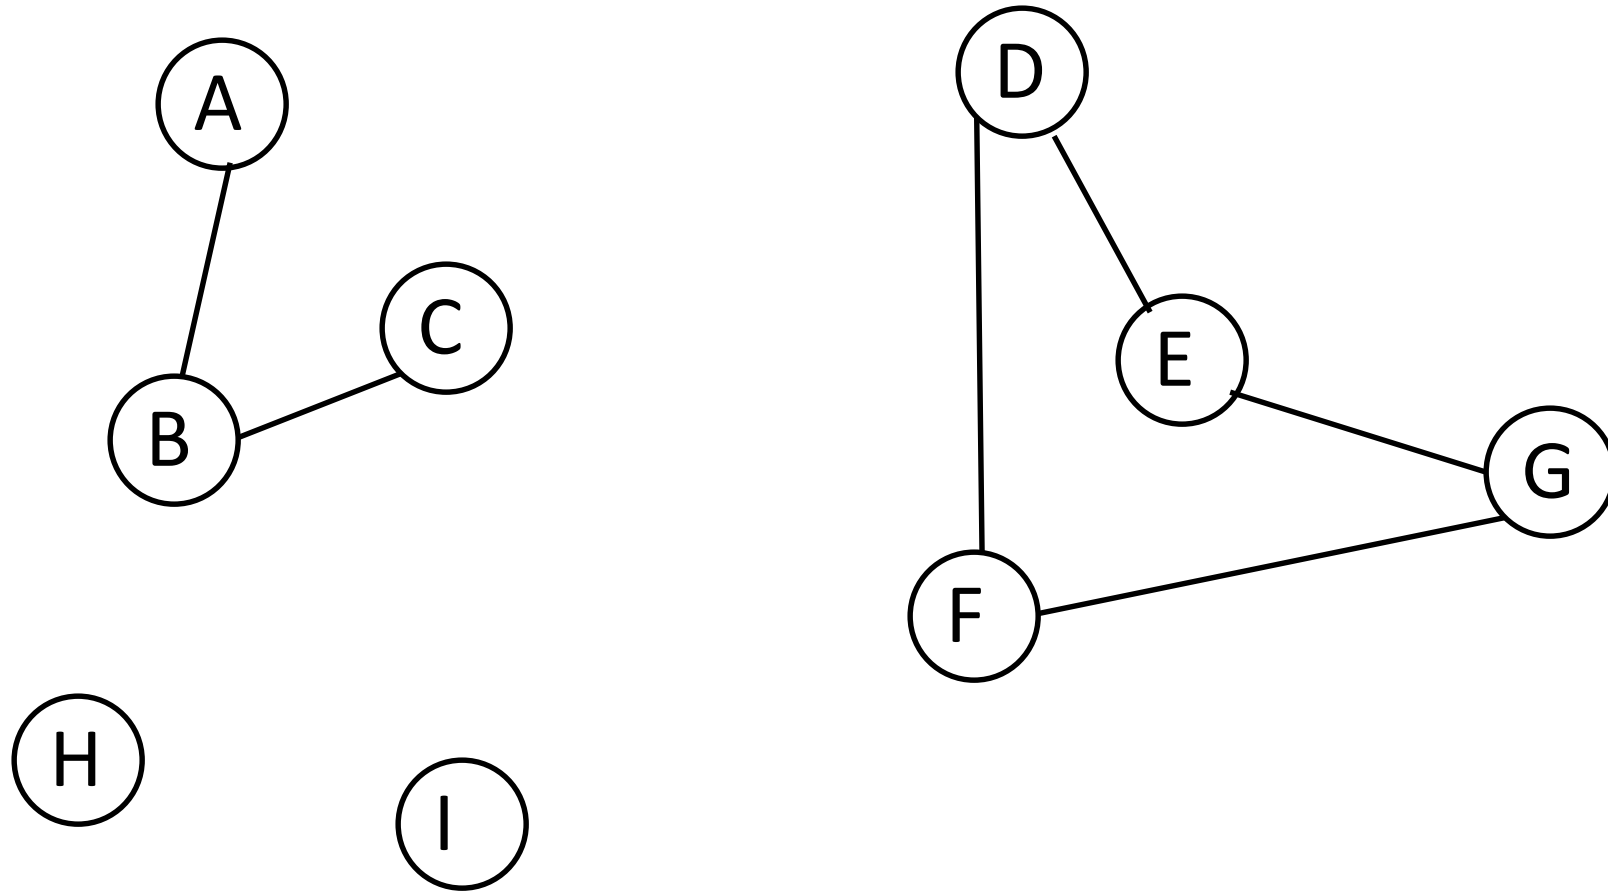
```
    If not w.visited
```

```
      explore(w)
```

Runtime $O(|V|+|E|)$.

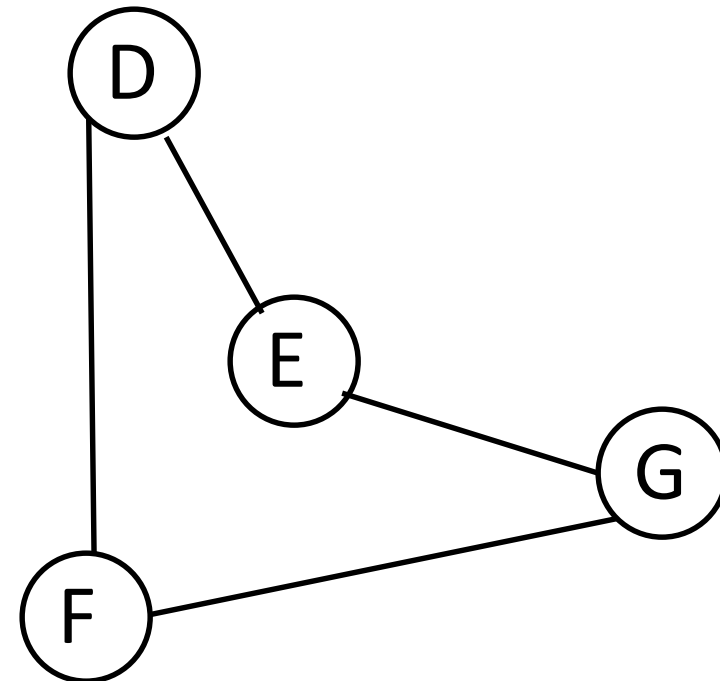
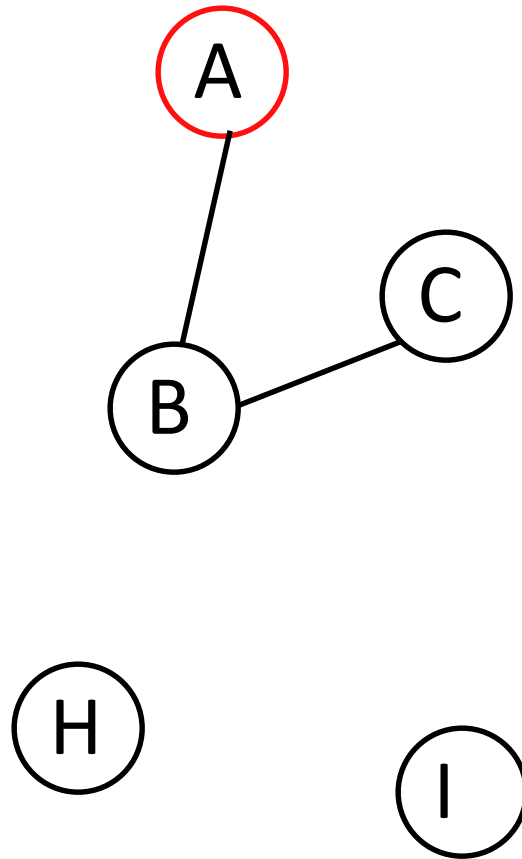
Example

CCNum:



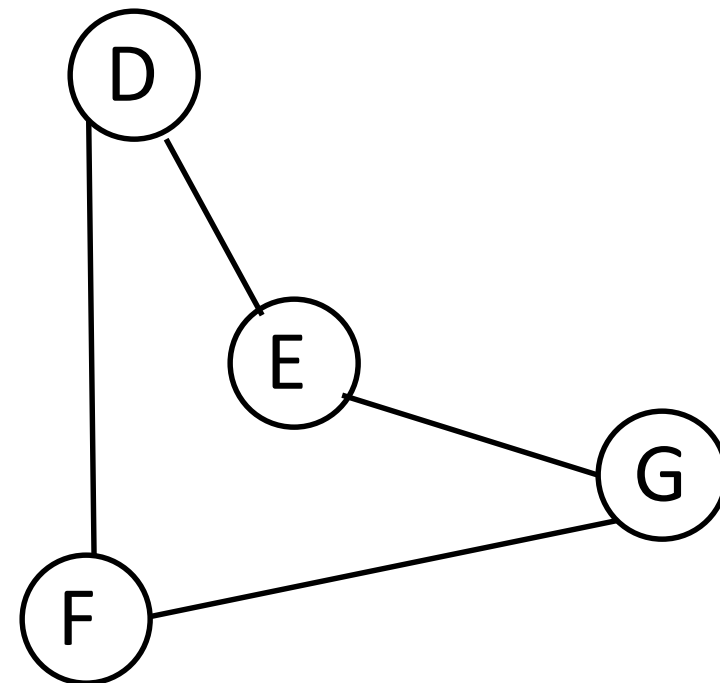
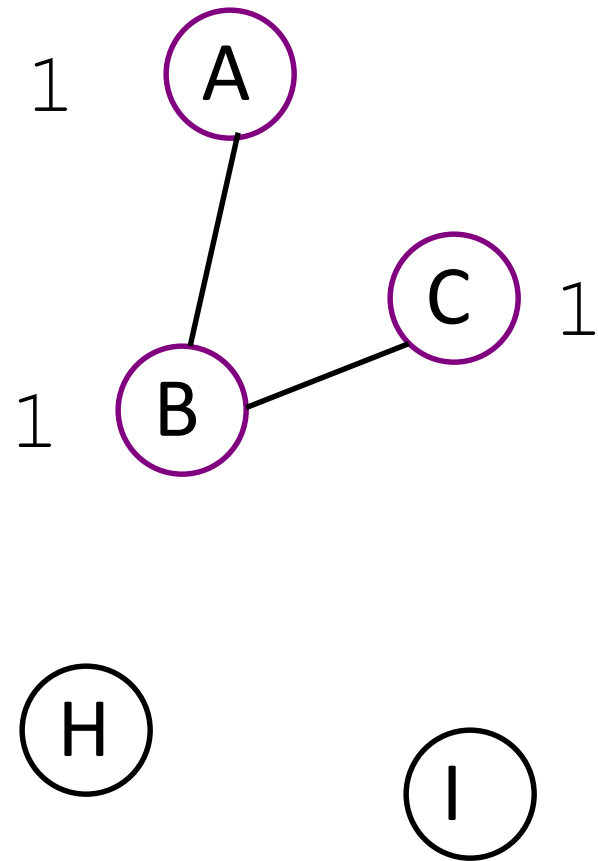
Example

CCNum: 1



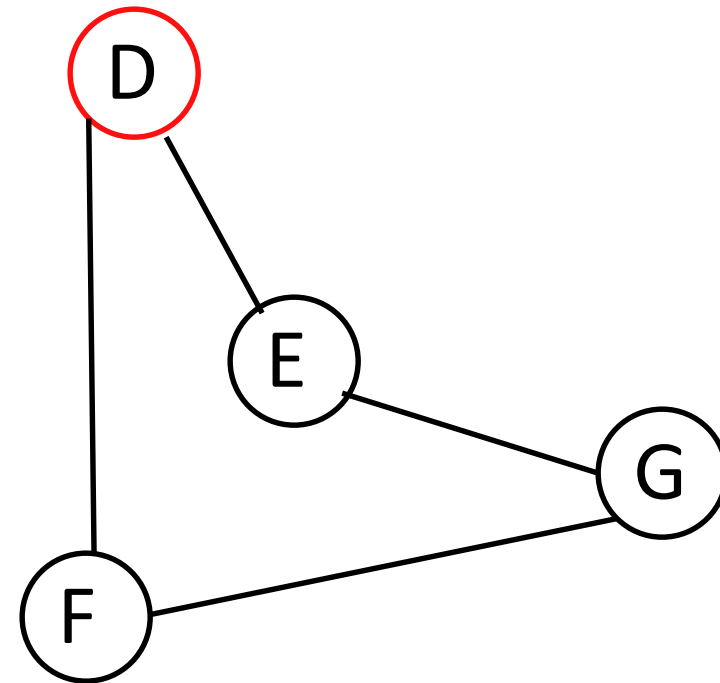
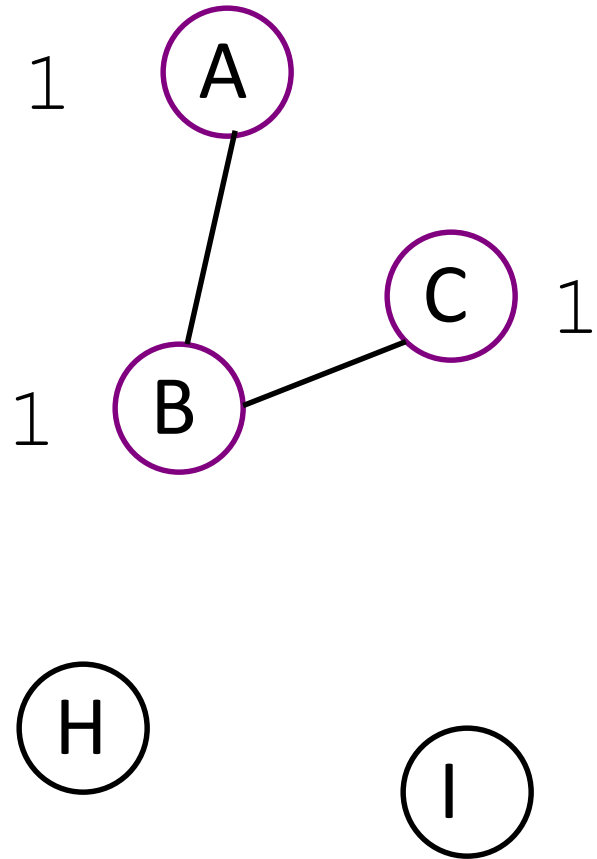
Example

CCNum: 1



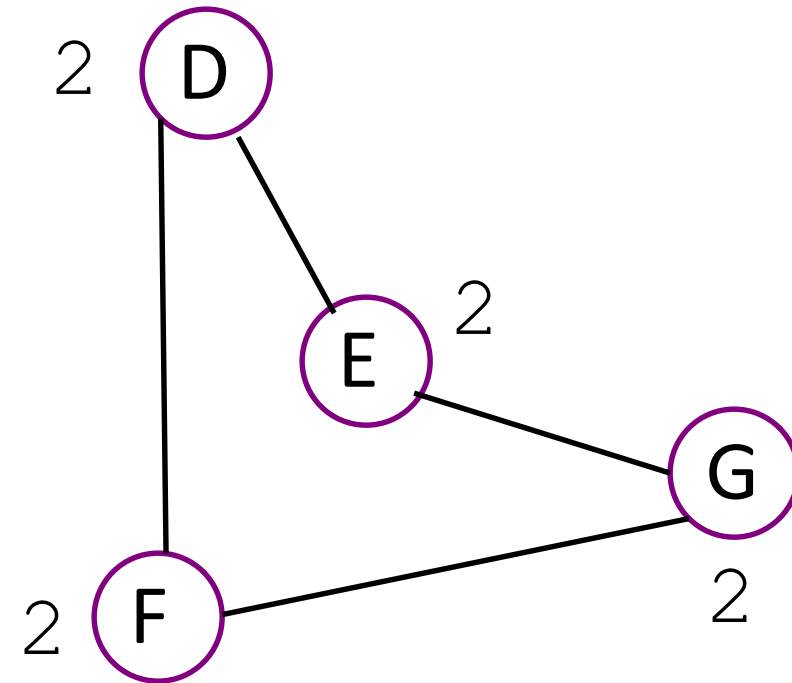
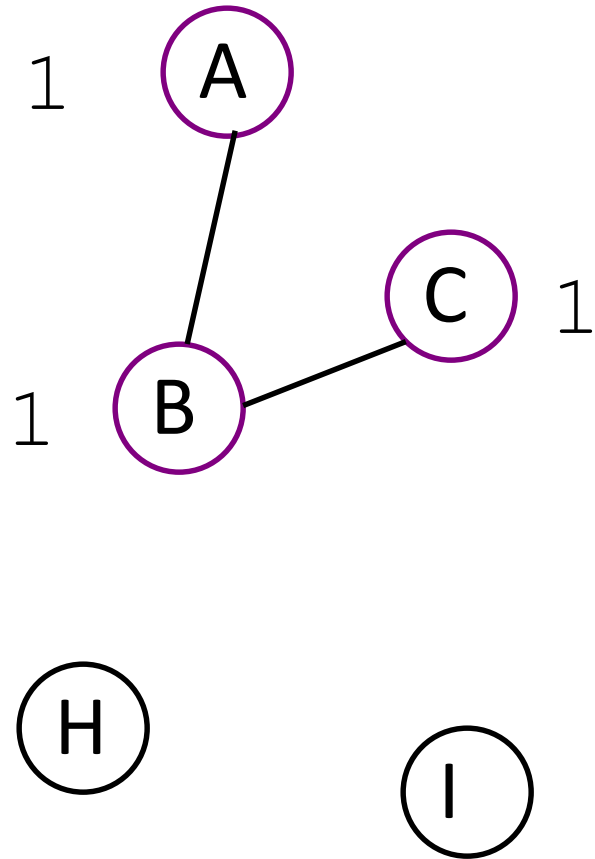
Example

CCNum: 2



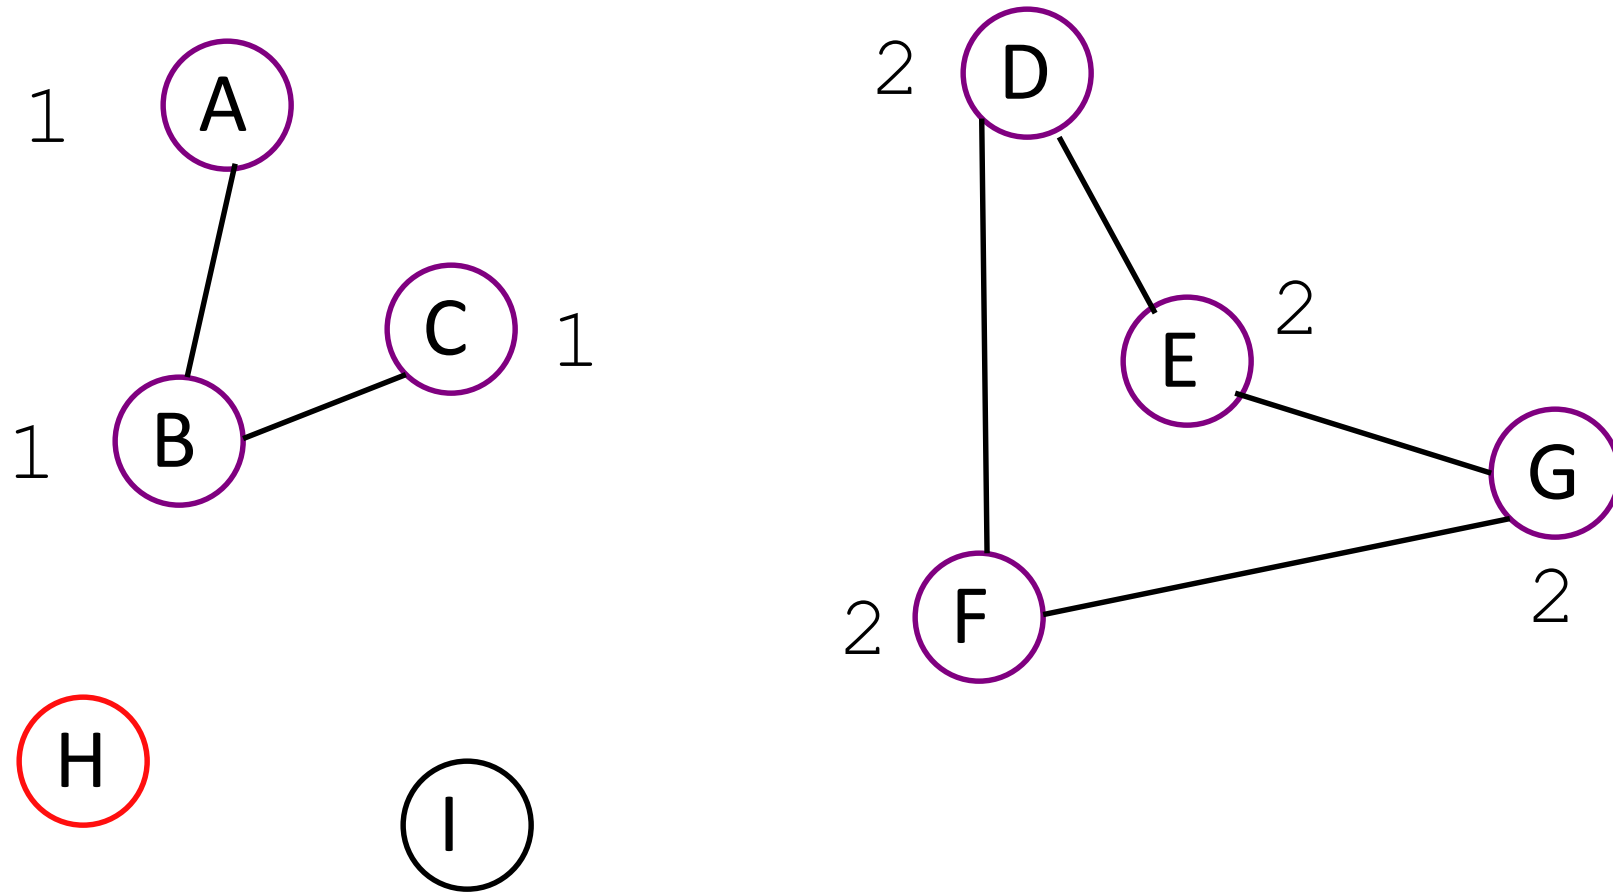
Example

CCNum: 2



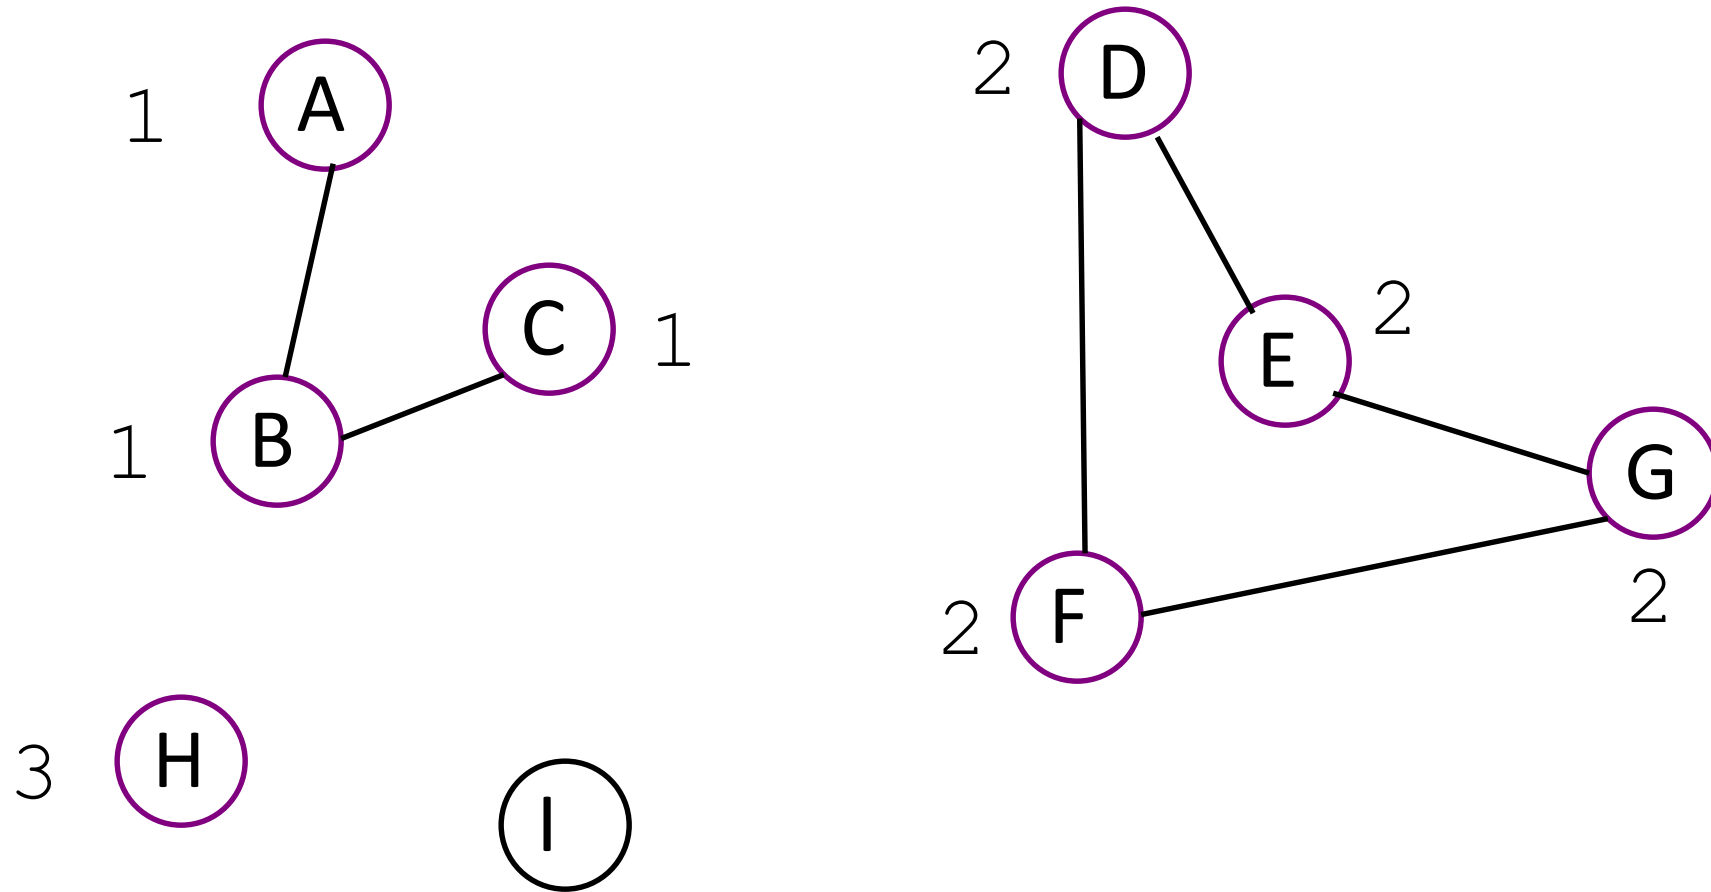
Example

CCNum: 3



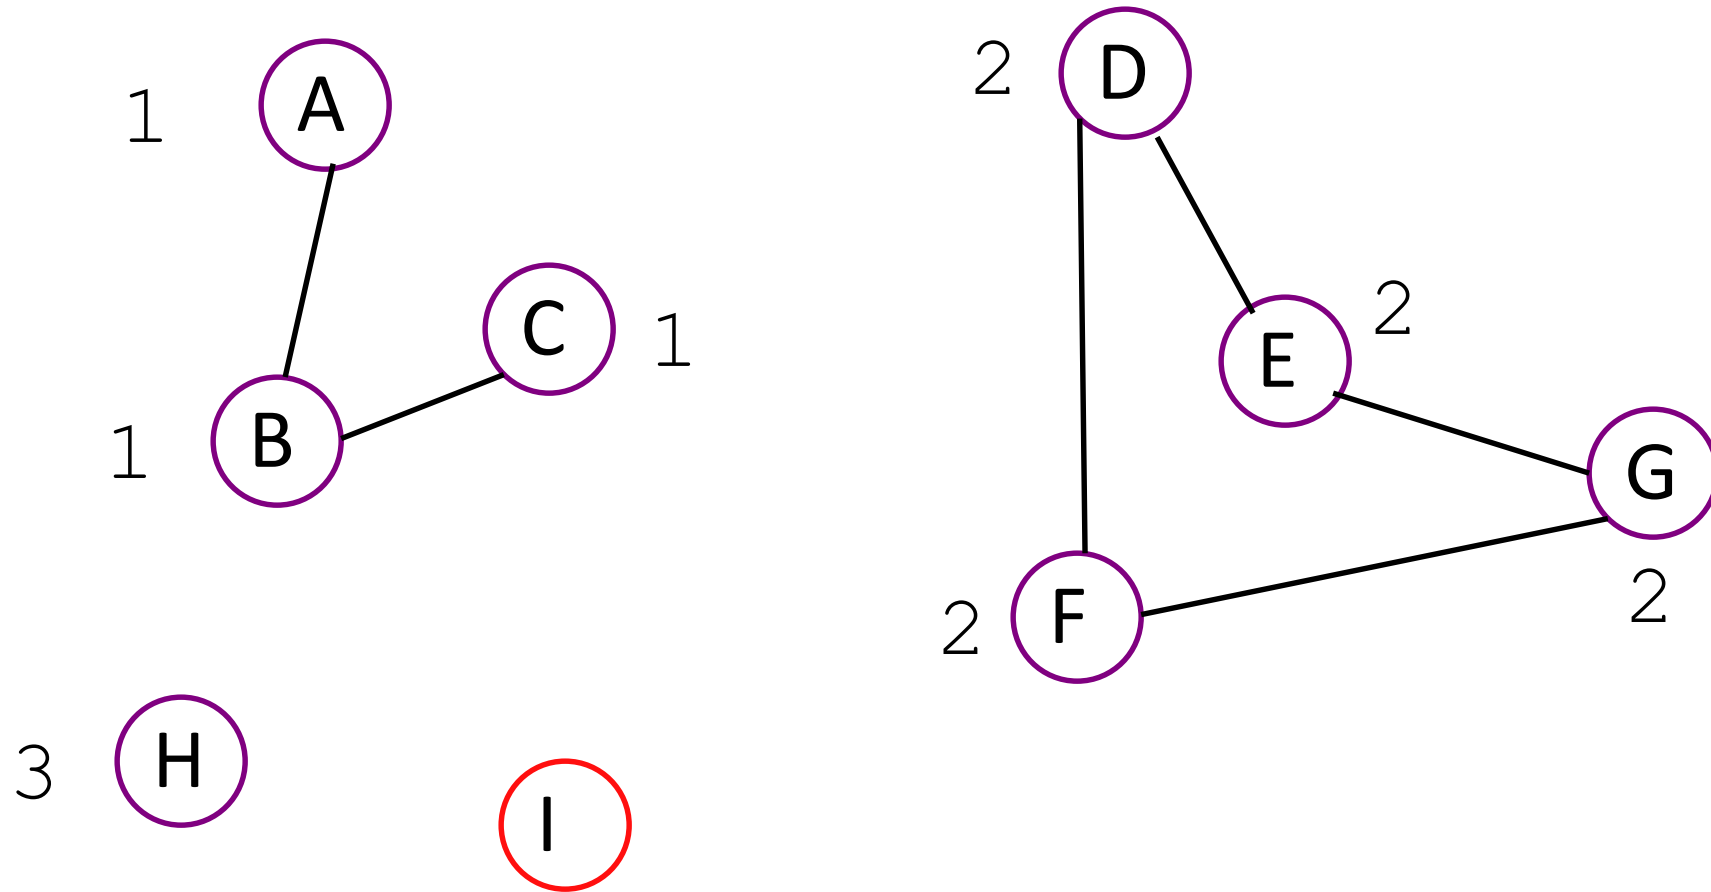
Example

CCNum: 3



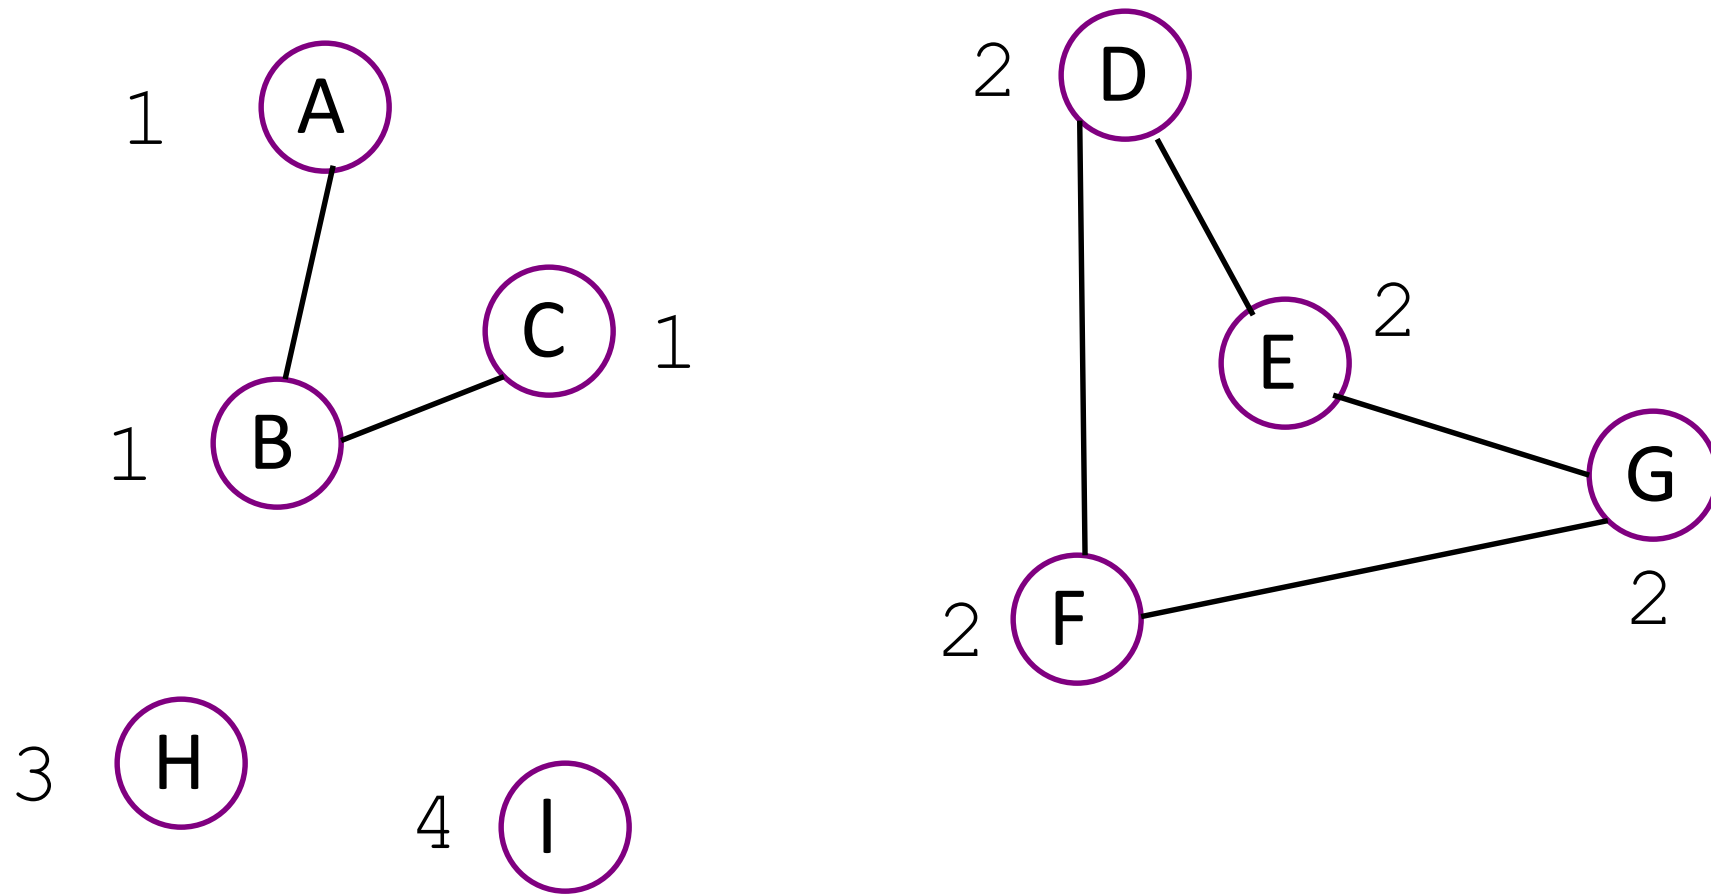
Example

CCNum: 4



Example

CCNum: 4



Discussion about DFS

What does DFS actually *do*?

Discussion about DFS

What does DFS actually *do*?

- No output.
- Marks all vertices as visited.
- Easier ways to do this.

Discussion about DFS

What does DFS actually *do*?

- No output.
- Marks all vertices as visited.
- Easier ways to do this.

However, DFS also is a useful way to explore the graph. By *augmenting* the algorithm a bit (like we did with the connected components algorithm), we can learn useful things.

Pre- and Post- Orders

Augment how?

- Keep track of what algorithm does & in what order.
- Have a “clock” and note time whenever:
 - Algorithm visits a new vertex for the first time.
 - Algorithm finishes processing a vertex.
- Record values as $v.pre$ and $v.post$.

Computing Pre- & Post- Orders

```
ConnectedComponents (G)
```

```
  For  $v \in G$ 
```

```
     $v.\text{visited} \leftarrow \text{false}$ 
```

```
  For  $v \in G$ 
```

```
    If not  $v.\text{visited}$ 
```

```
      explore( $v$ )
```

```
explore( $v$ )
```

```
   $v.\text{visited} \leftarrow \text{true}$ 
```

```
  For each edge  $(v, w)$ 
```

```
    If not  $w.\text{visited}$ 
```

```
      explore( $w$ )
```

Computing Pre- & Post- Orders

```
ConnectedComponents (G)
```

```
  clock  $\leftarrow$  1
```

```
  For v  $\in$  G
```

```
    v.visited  $\leftarrow$  false
```

```
  For v  $\in$  G
```

```
    If not v.visited
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

Computing Pre- & Post- Orders

```
ConnectedComponents (G)
```

```
  clock  $\leftarrow$  1
```

```
  For v  $\in$  G
```

```
    v.visited  $\leftarrow$  false
```

```
  For v  $\in$  G
```

```
    If not v.visited
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  v.pre  $\leftarrow$  clock
```

```
  clock++
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

Computing Pre- & Post- Orders

```
ConnectedComponents (G)
```

```
  clock  $\leftarrow$  1
```

```
  For v  $\in$  G
```

```
    v.visited  $\leftarrow$  false
```

```
  For v  $\in$  G
```

```
    If not v.visited
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  v.pre  $\leftarrow$  clock
```

```
  clock++
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

```
      explore(w)
```

```
  v.post  $\leftarrow$  clock
```

```
  clock++
```


Computing Pre- & Post- Orders

```
ConnectedComponents (G)
```

```
  clock  $\leftarrow$  1
```

```
  For v  $\in$  G
```

```
    v.visited  $\leftarrow$  false
```

```
  For v  $\in$  G
```

```
    If not v.visited
```

```
      explore(v)
```

```
explore(v)
```

```
  v.visited  $\leftarrow$  true
```

```
  v.pre  $\leftarrow$  clock
```

```
  clock++
```

```
  For each edge (v,w)
```

```
    If not w.visited
```

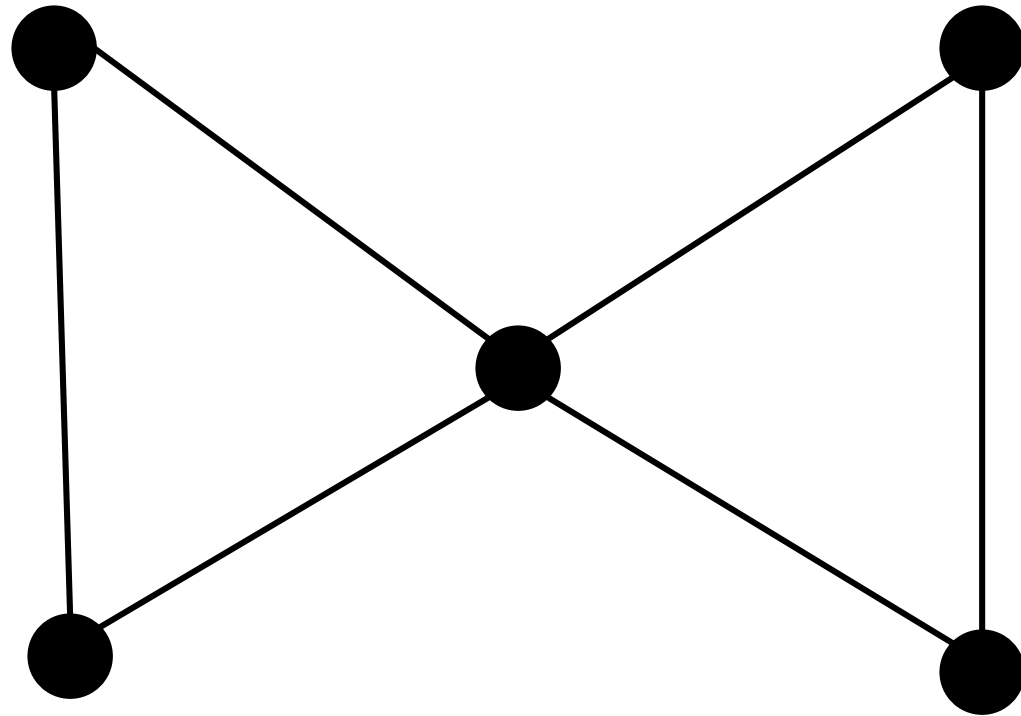
```
      explore(w)
```

```
  v.post  $\leftarrow$  clock
```

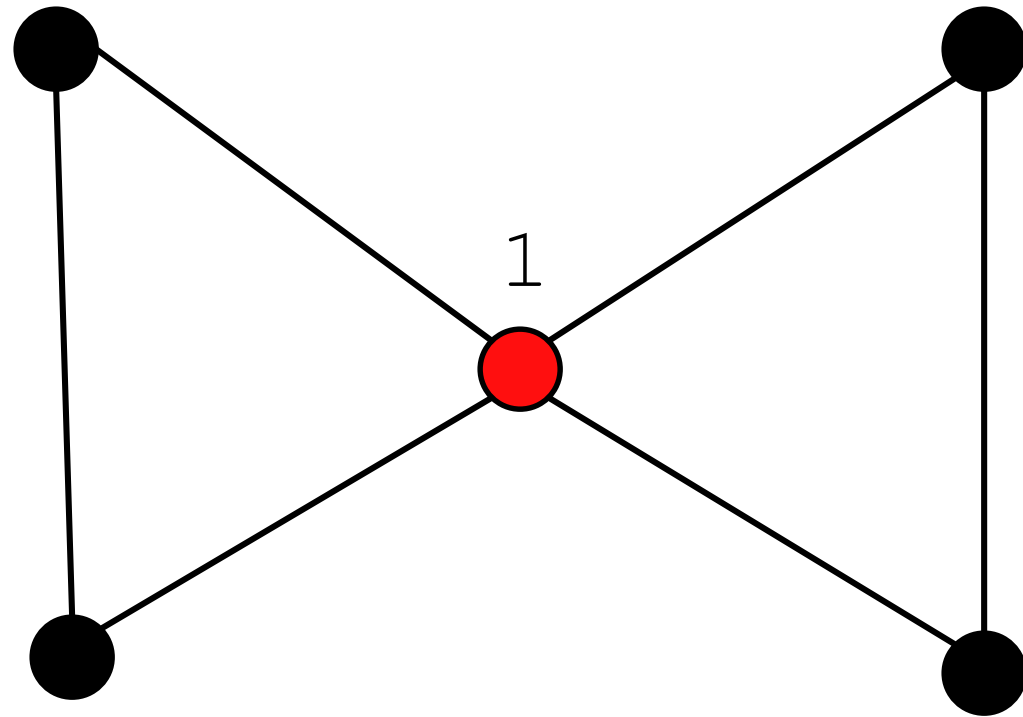
```
  clock++
```

Runtime $O(|V|+|E|)$.

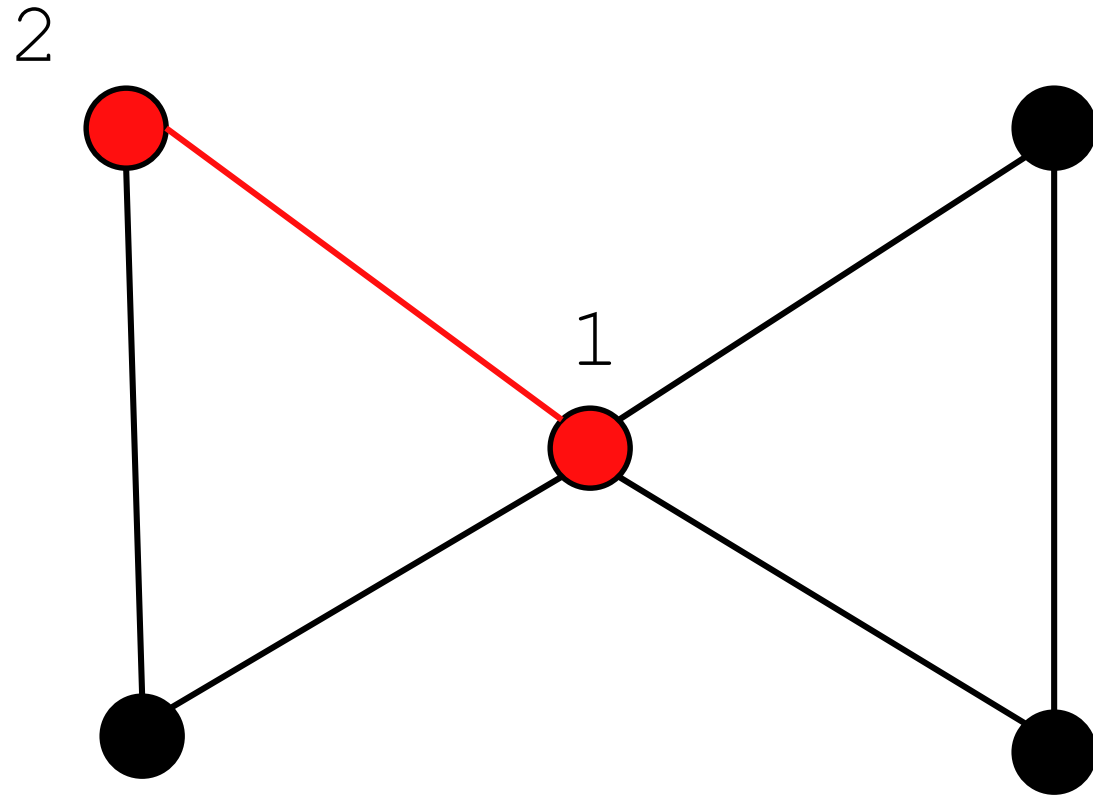
Example



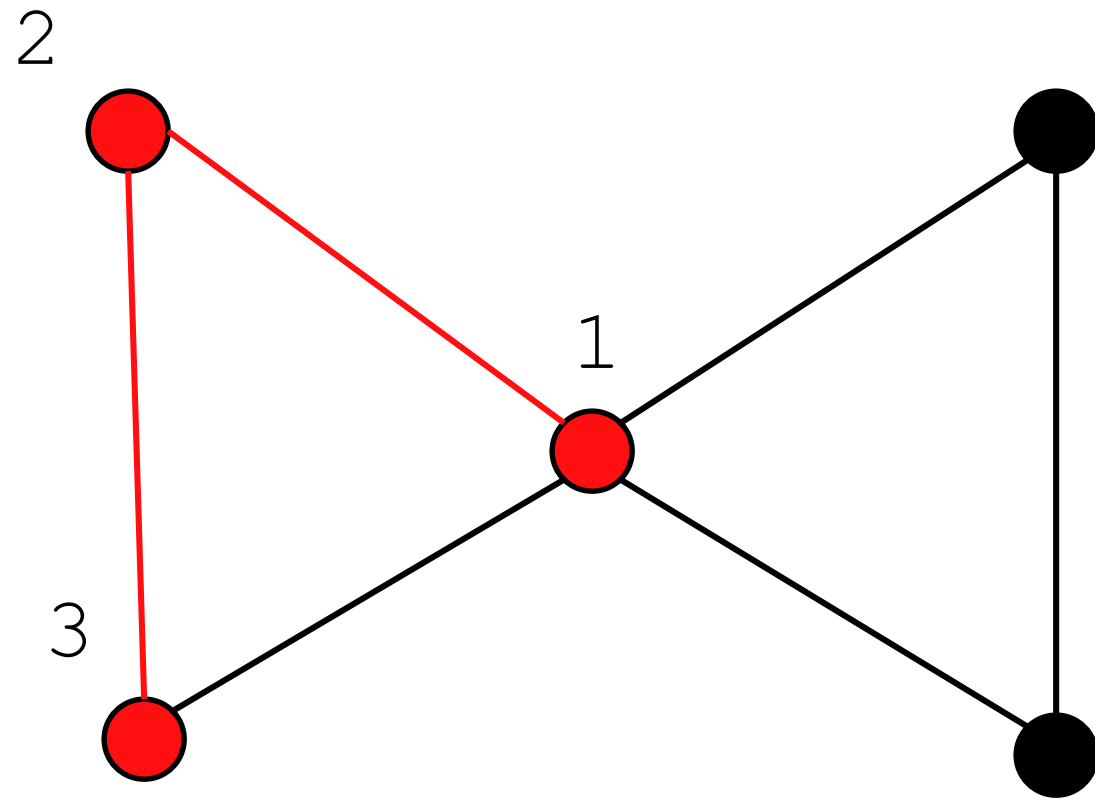
Example



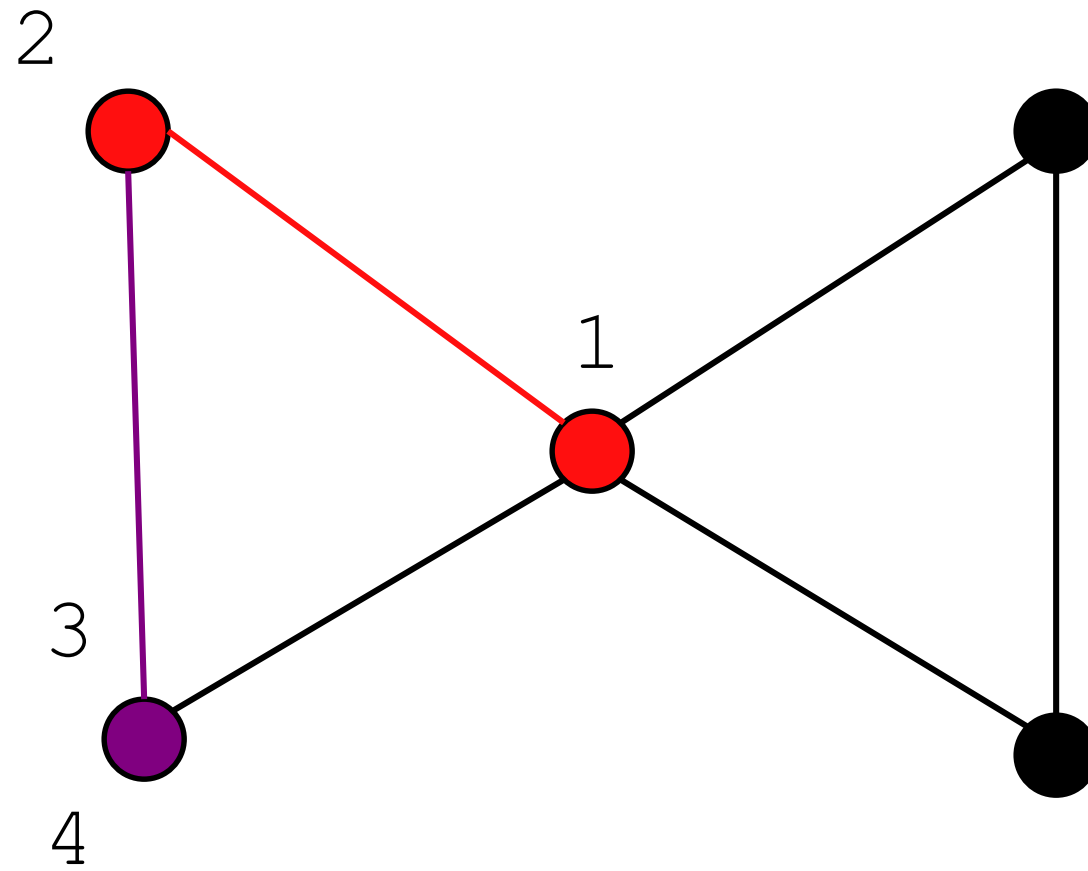
Example



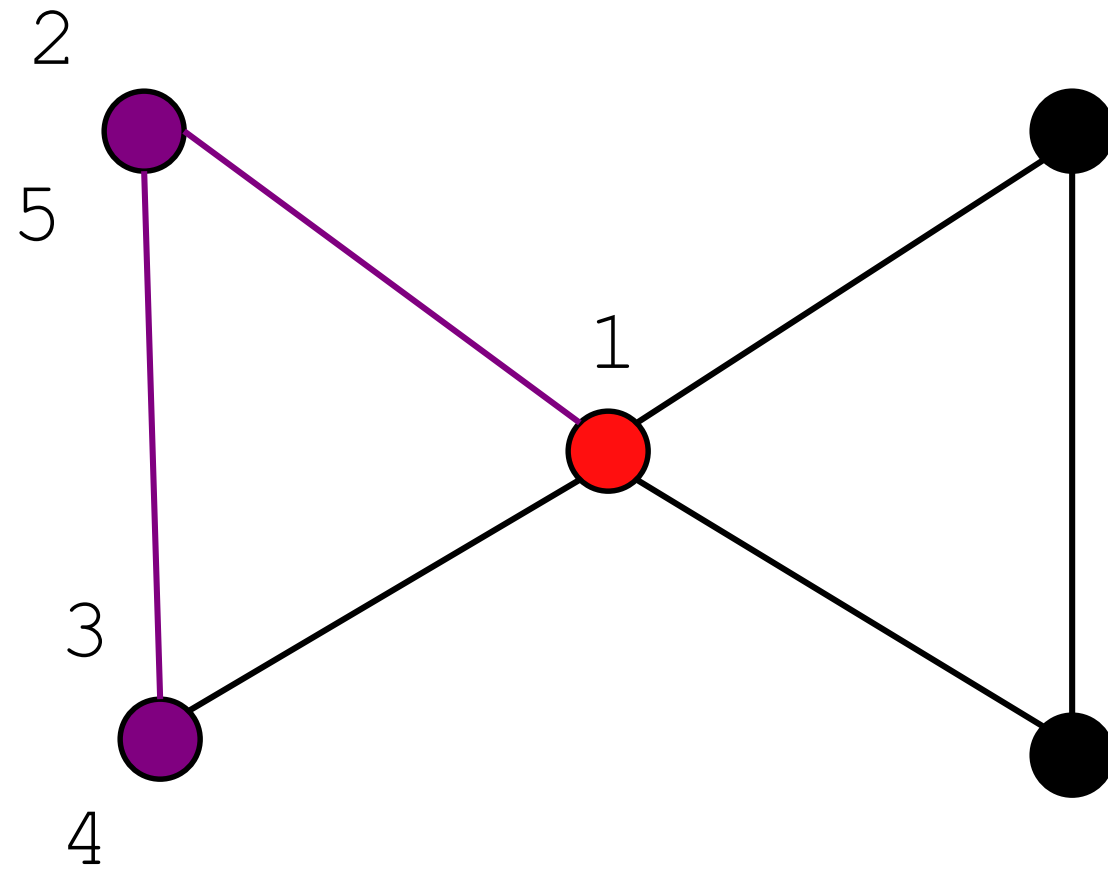
Example



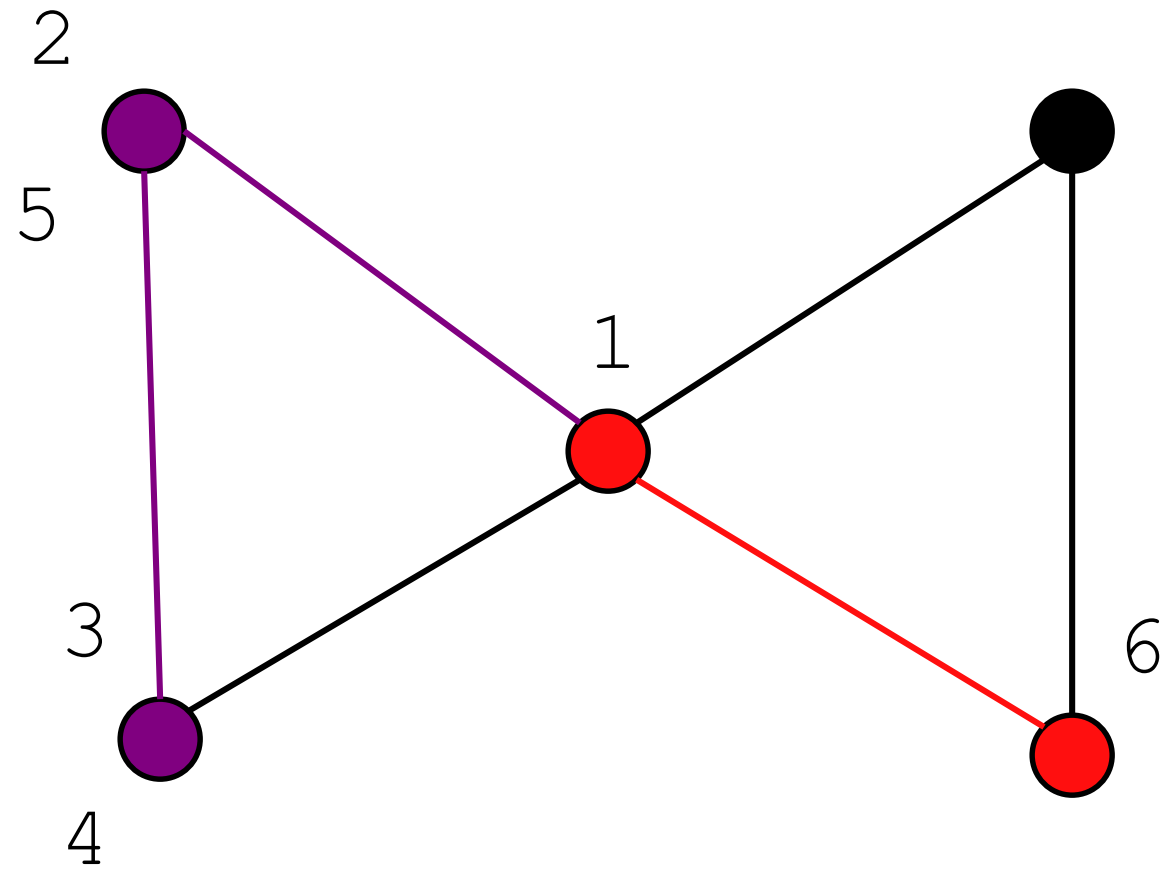
Example



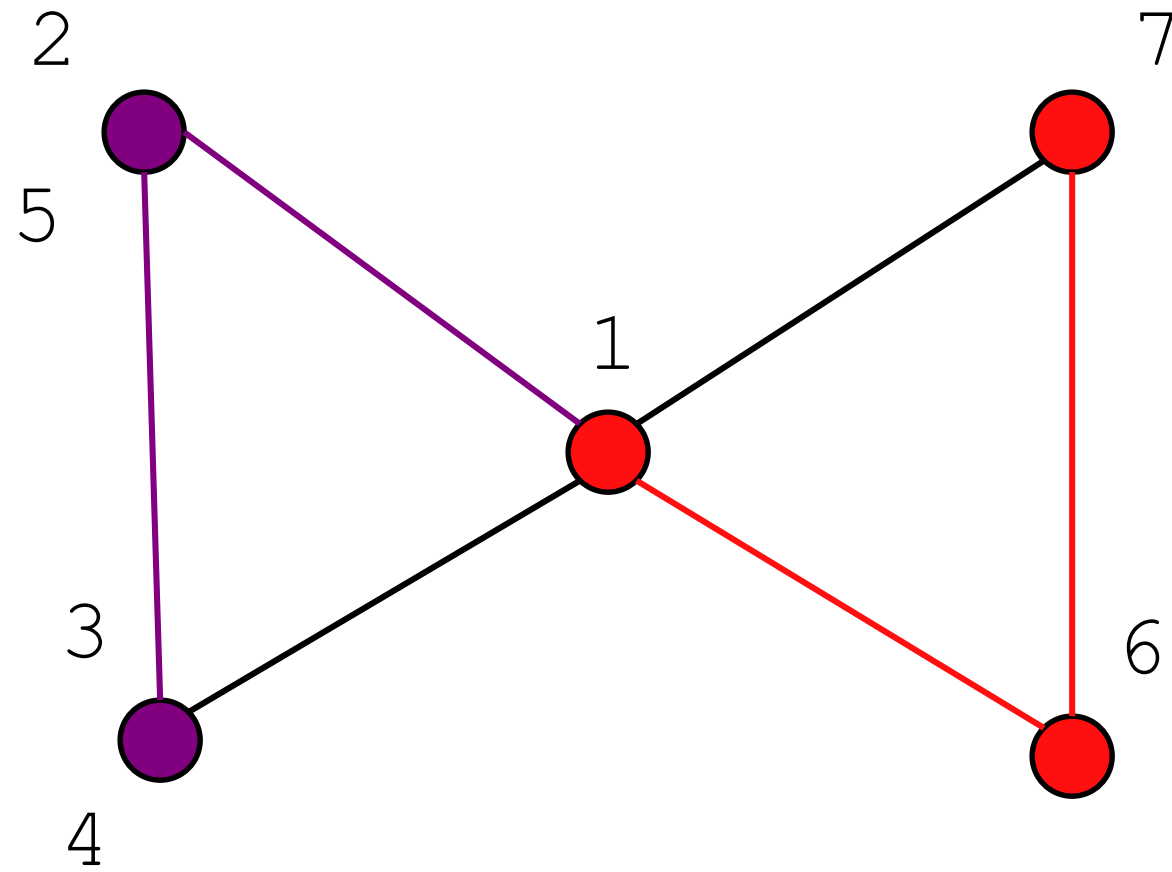
Example



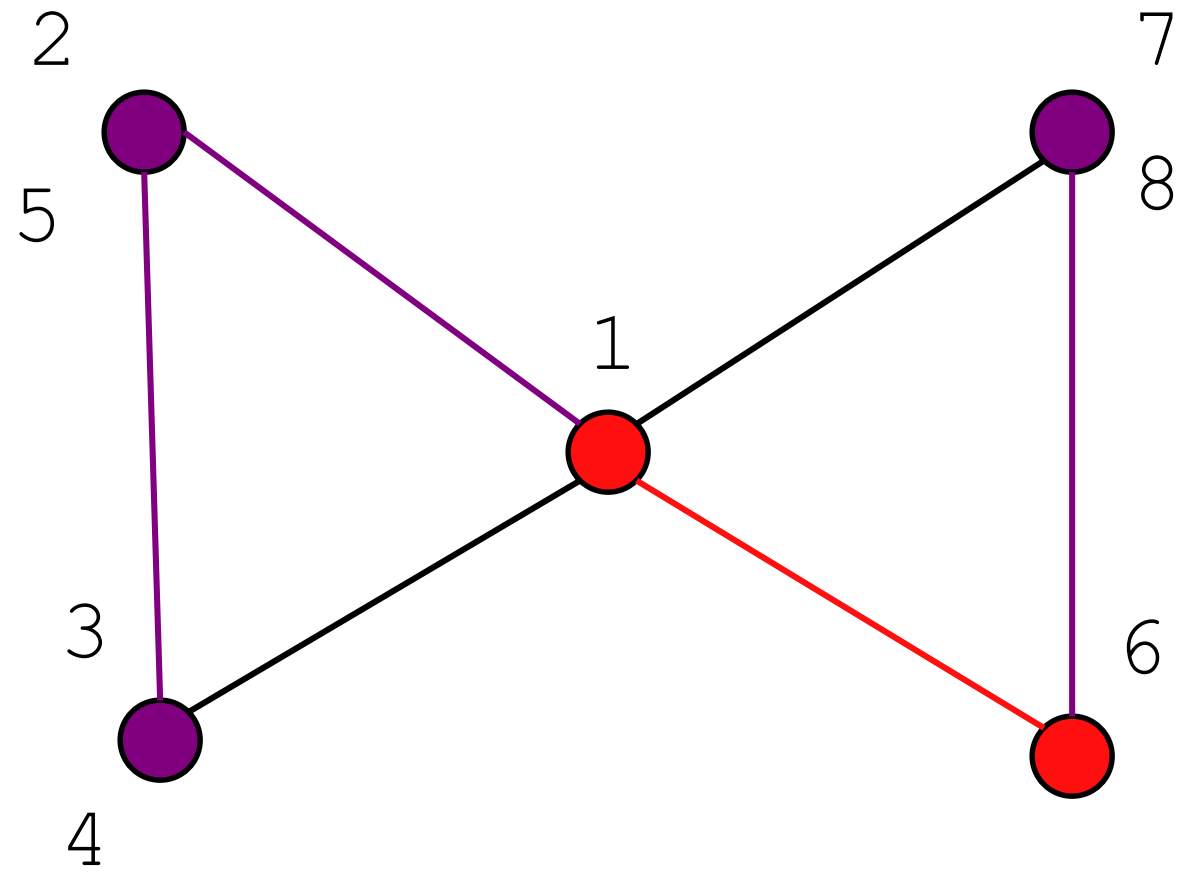
Example



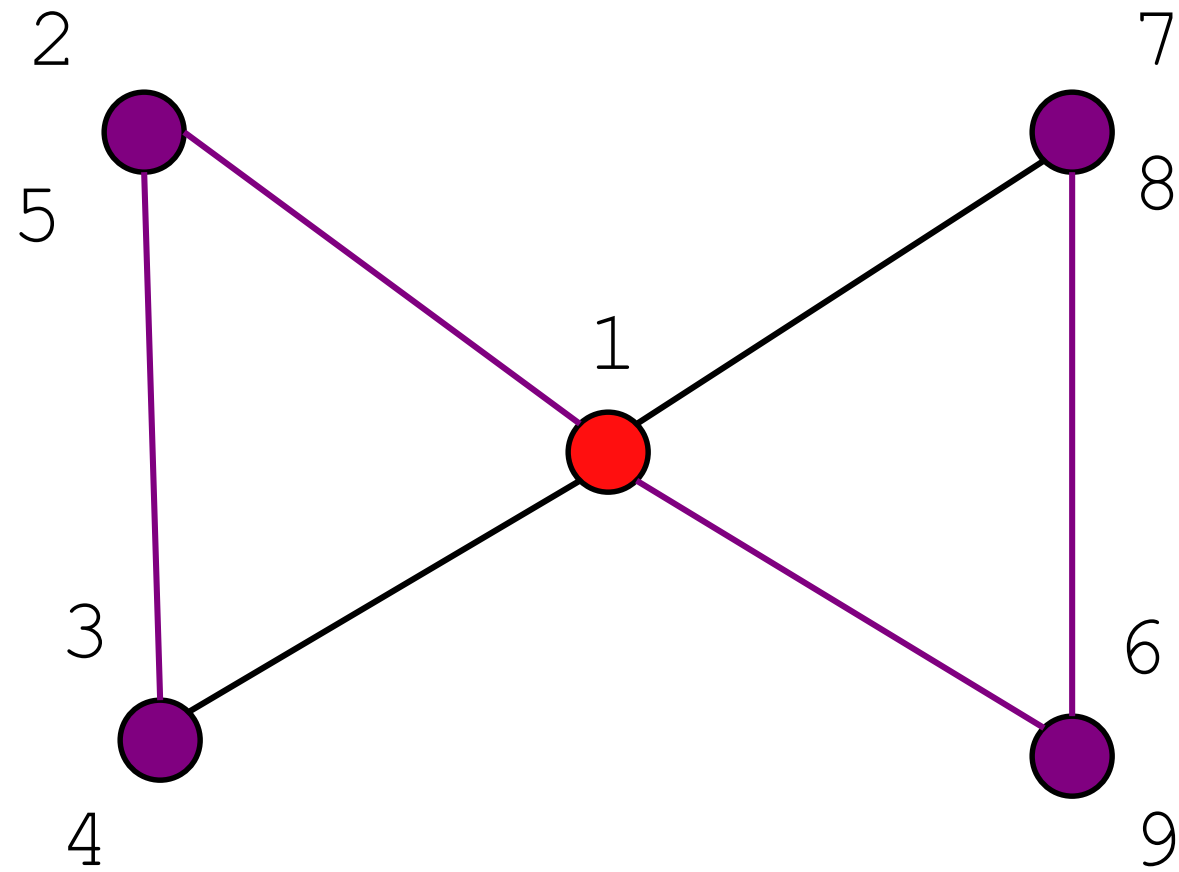
Example



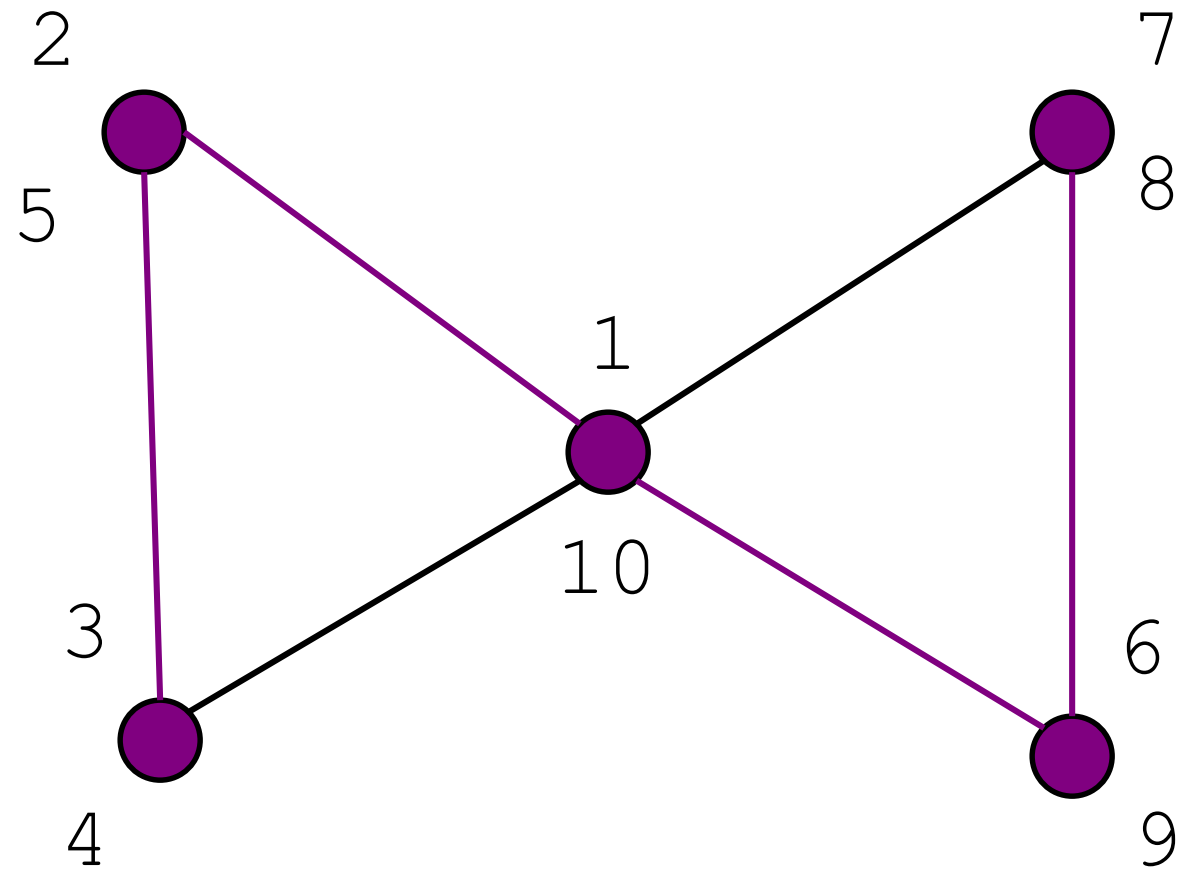
Example



Example

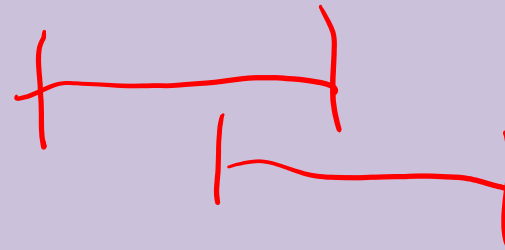


Example



What do these orders tell us?

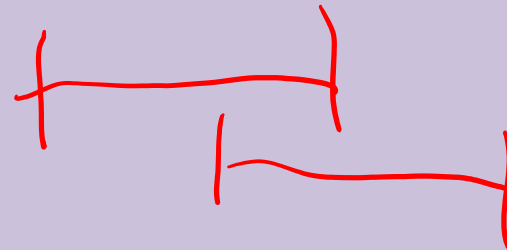
Prop: For vertices v, w consider intervals $[v.pre, v.post]$ and $[w.pre, w.post]$.
These intervals:



What do these orders tell us?

Prop: For vertices v, w consider intervals $[v.pre, v.post]$ and $[w.pre, w.post]$. These intervals:

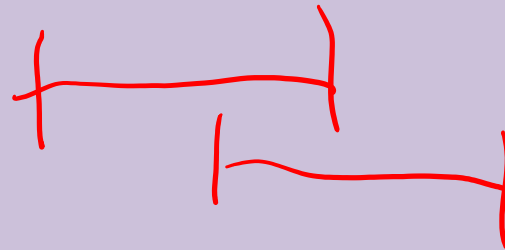
1. Contain each other if v is an ancestor/descendant of w in the DFS tree.



What do these orders tell us?

Prop: For vertices v, w consider intervals $[v.pre, v.post]$ and $[w.pre, w.post]$. These intervals:

1. Contain each other if v is an ancestor/descendant of w in the DFS tree.
2. Are disjoint if v and w are cousins in the DFS tree.



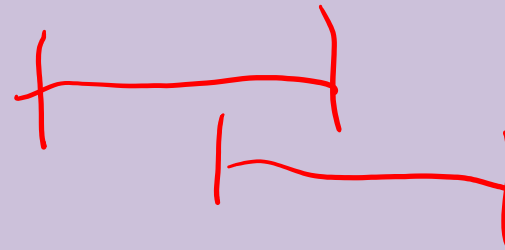
What do these orders tell us?

Prop: For vertices v, w consider intervals $[v.pre, v.post]$ and $[w.pre, w.post]$. These intervals:

1. Contain each other if v is an ancestor/descendant of w in the DFS tree.
2. Are disjoint if v and w are cousins in the DFS tree.

3. Never interleave

$$(v.pre < w.pre < v.post < w.post)$$



Proof

- Assume algorithm finds v before w
($v.pre < w.pre$)

Proof

- Assume algorithm finds v before w
($v.pre < w.pre$)
- If algorithm discovers w *after* fully processing v :
 - $v.post < w.pre$
 - Intervals disjoint
 - v and w are cousins

Proof

- Assume algorithm finds v before w
($v.pre < w.pre$)
- If algorithm discovers w *after* fully processing v :
 - $v.post < w.pre$
 - Intervals disjoint
 - v and w are cousins
- If algorithm discovers w *before* fully processing v :
 - Algorithm finishes processing w before it finishes v
 - $v.pre < w.pre < w.post < v.post$
 - Nested intervals
 - v is ancestor of w

Question: Possible Intervals

Which pairs of pre-post intervals are not possible for DFS?
(multiple correct answers)

A) [1,2] & [3,4]

B) [1,3] & [2,4]

C) [1,4] & [2,3]

D) [1,5] & [2,4]

E) [1,6] & [2,5]

Question: Possible Intervals

Which pairs of pre-post intervals are not possible for DFS?
(multiple correct answers)

A) [1,2] & [3,4]

B) [1,3] & [2,4]

C) [1,4] & [2,3]

D) [1,5] & [2,4]

E) [1,6] & [2,5]

Directed Graphs

Often an edge makes sense both ways, but sometimes streets are one directional.

Directed Graphs

Often an edge makes sense both ways, but sometimes streets are one directional.

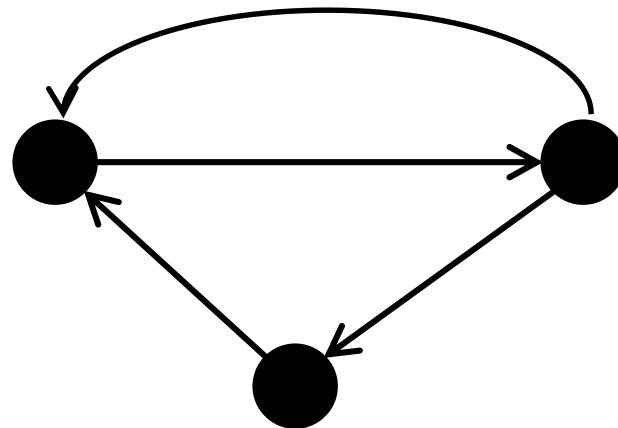
Definition: A directed graph is a graph where each edge has a direction. Goes *from* v *to* w .

Directed Graphs

Often an edge makes sense both ways, but sometimes streets are one directional.

Definition: A directed graph is a graph where each edge has a direction. Goes *from* v to w .

Draw edges with arrows to denote direction.



Question: Directed Graphs

Which of the following does NOT need to be modeled as a directed rather than undirected graph:

- A) The Internet (links connecting webpages)
- B) Facebook (friendships connecting people)
- C) Twitter (followings connecting people)
- D) Maps (roads connecting intersections)

Question: Directed Graphs

Which of the following does NOT need to be modeled as a directed rather than undirected graph:

- A) The Internet (links connecting webpages)
- B) Facebook (friendships connecting people)
- C) Twitter (followings connecting people)
- D) Maps (roads connecting intersections)

DFS on Directed Graphs

DFS on Directed Graphs

- Same code

DFS on Directed Graphs

- Same code
- Only follow *directed* edges from v to w .

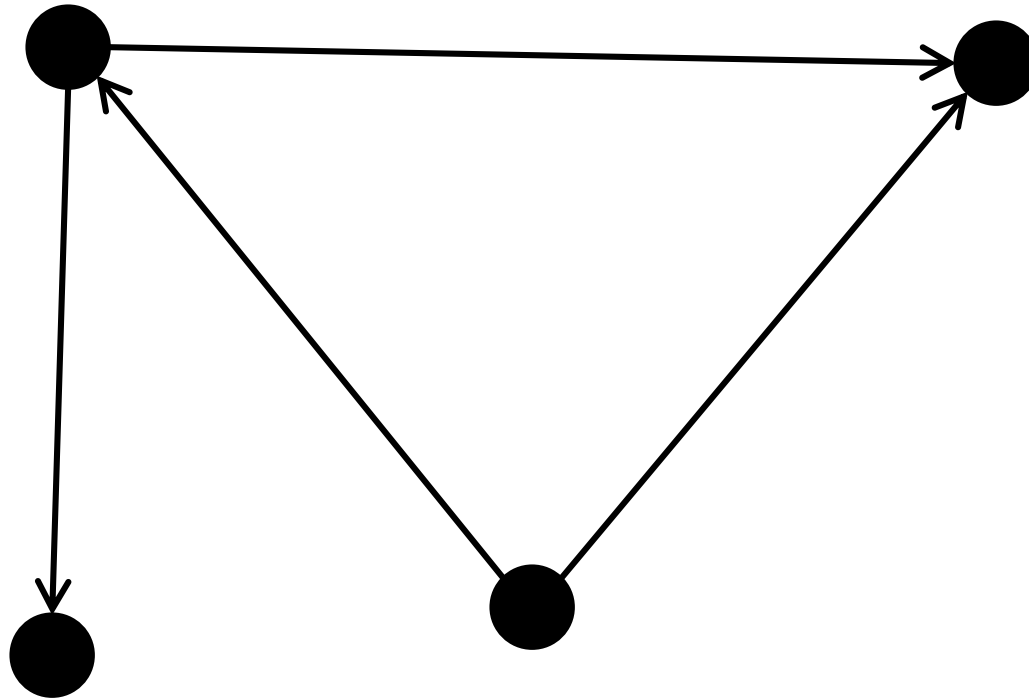
DFS on Directed Graphs

- Same code
- Only follow *directed* edges from v to w .
- Runtime still $O(|V| + |E|)$

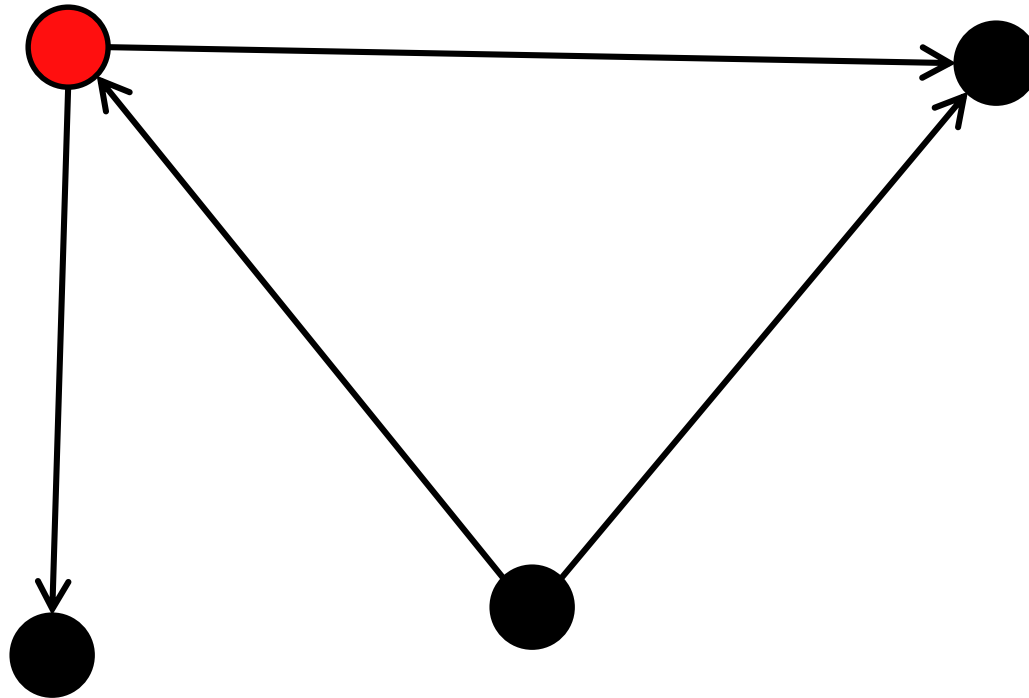
DFS on Directed Graphs

- Same code
- Only follow *directed* edges from v to w .
- Runtime still $O(|V|+|E|)$
- `explore(v)` discovers all vertices reachable from v following only directed edges.

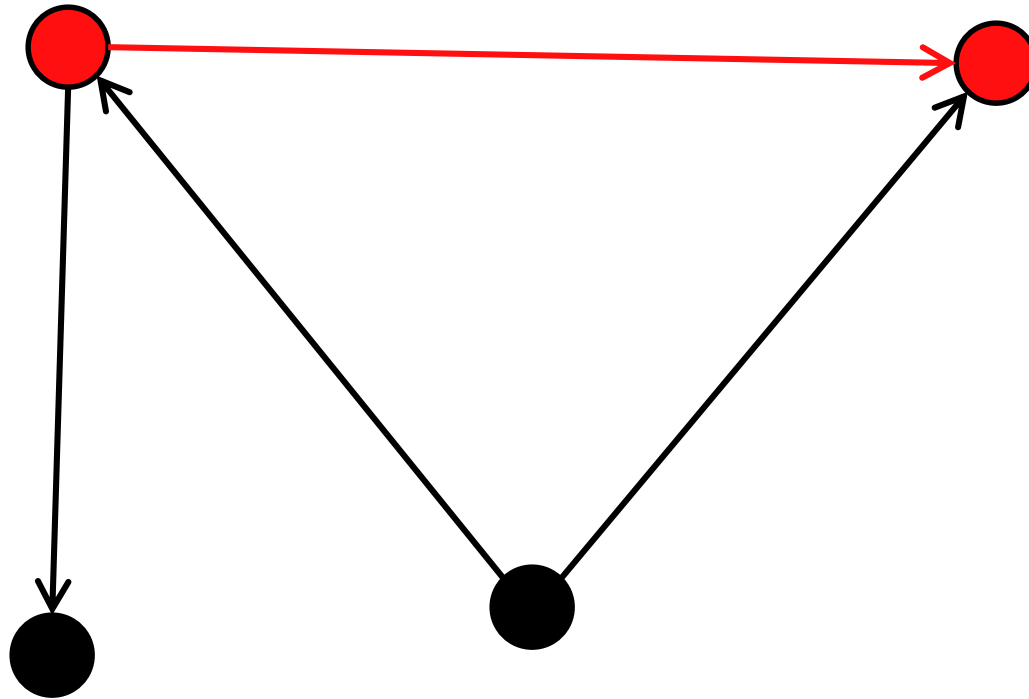
Example



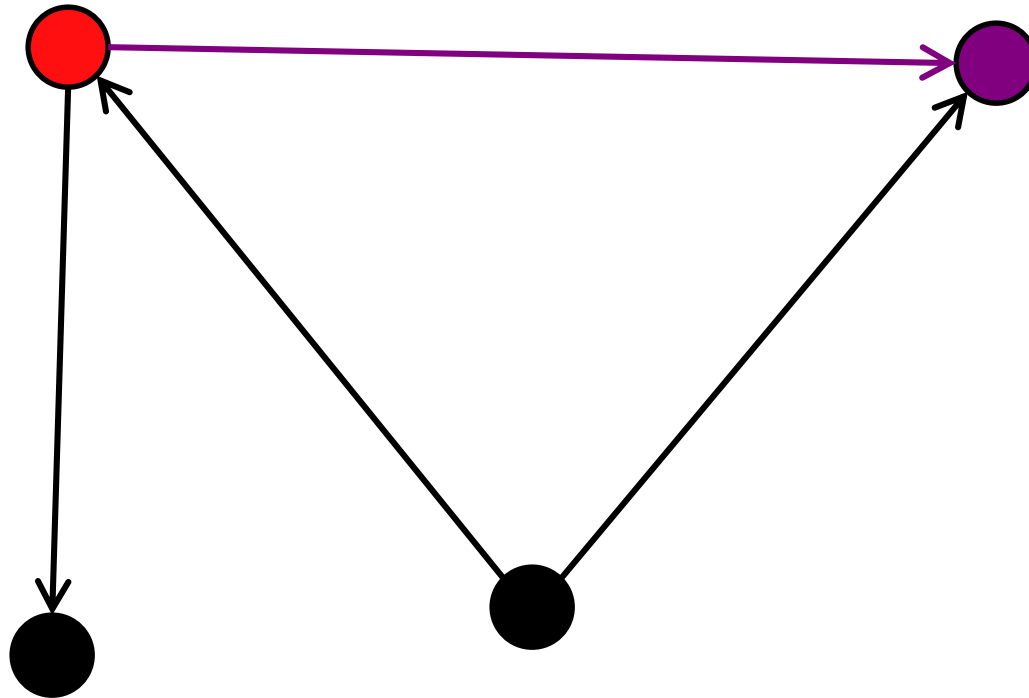
Example



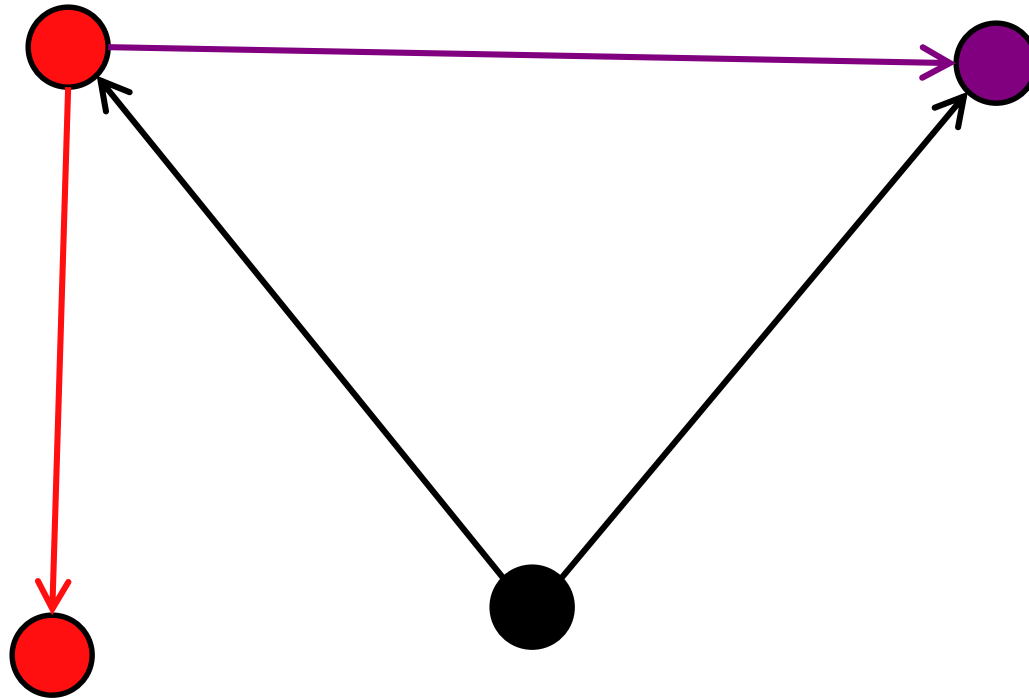
Example



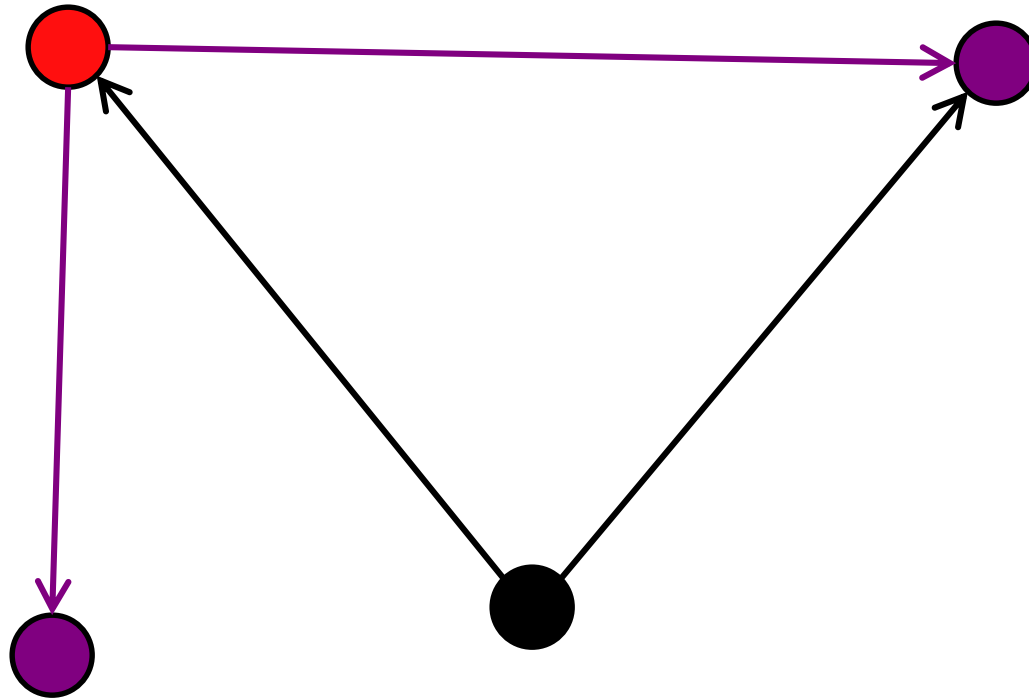
Example



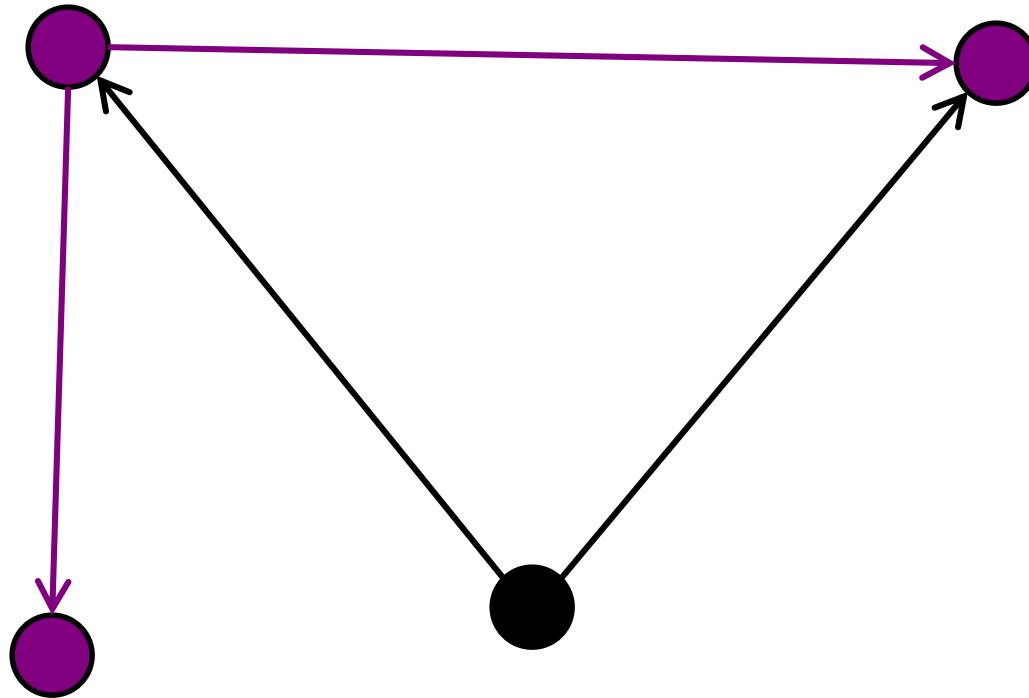
Example



Example



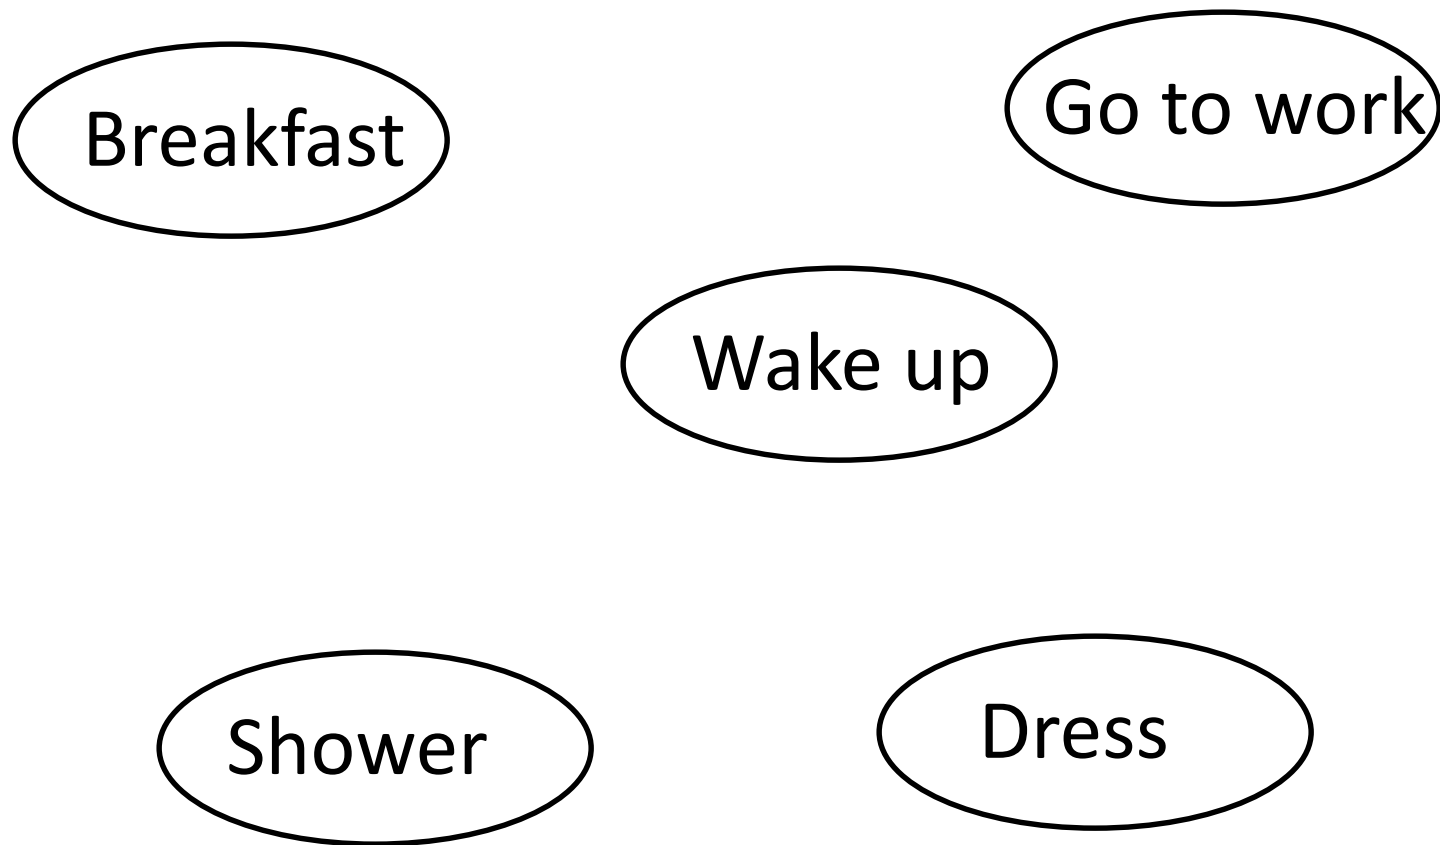
Example



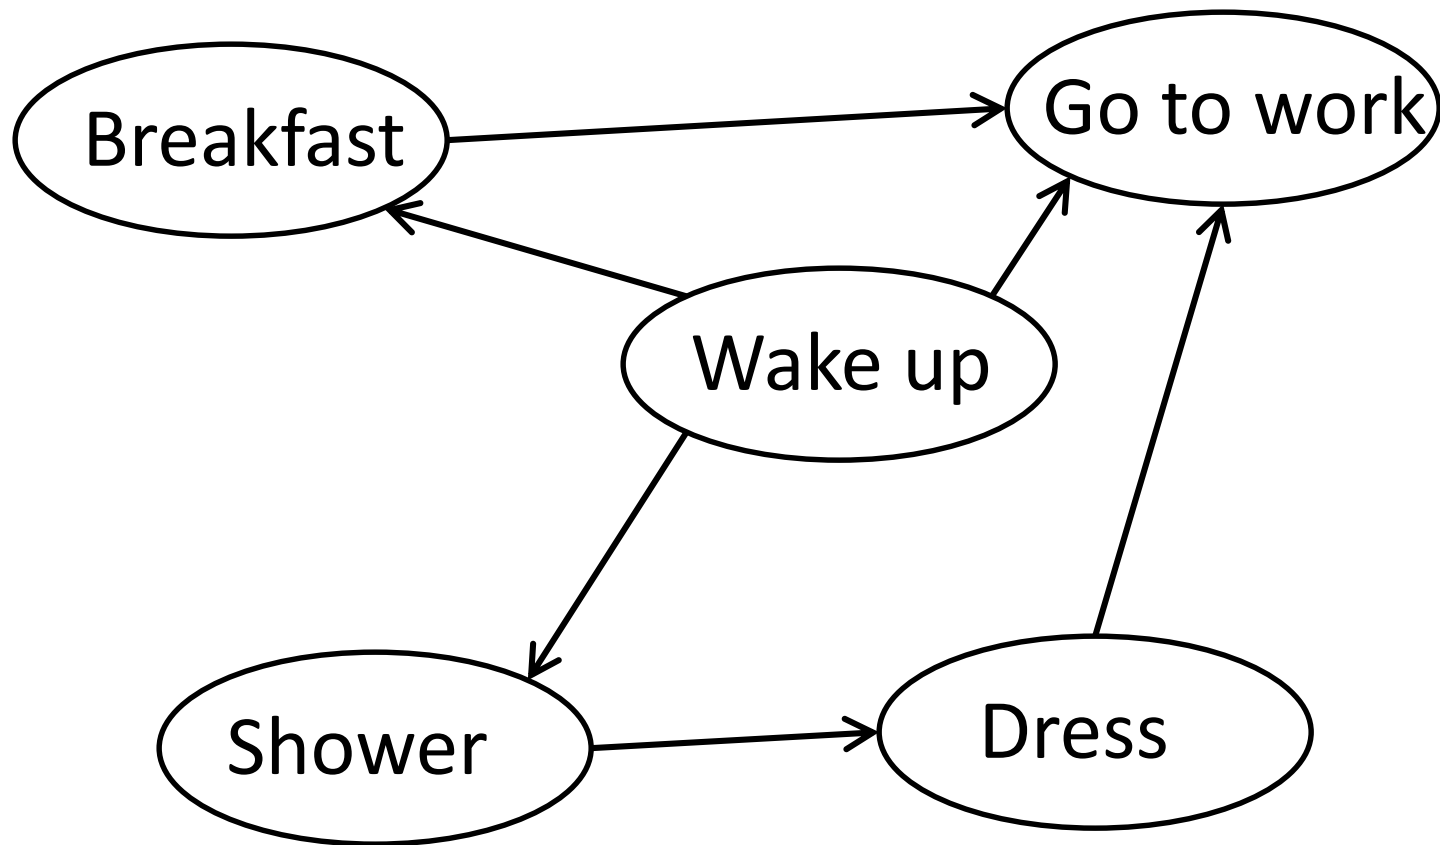
Directed Acyclic Graphs

- Directed graphs as dependencies
- Linear orderings
- DAGs definition
- Topological sort

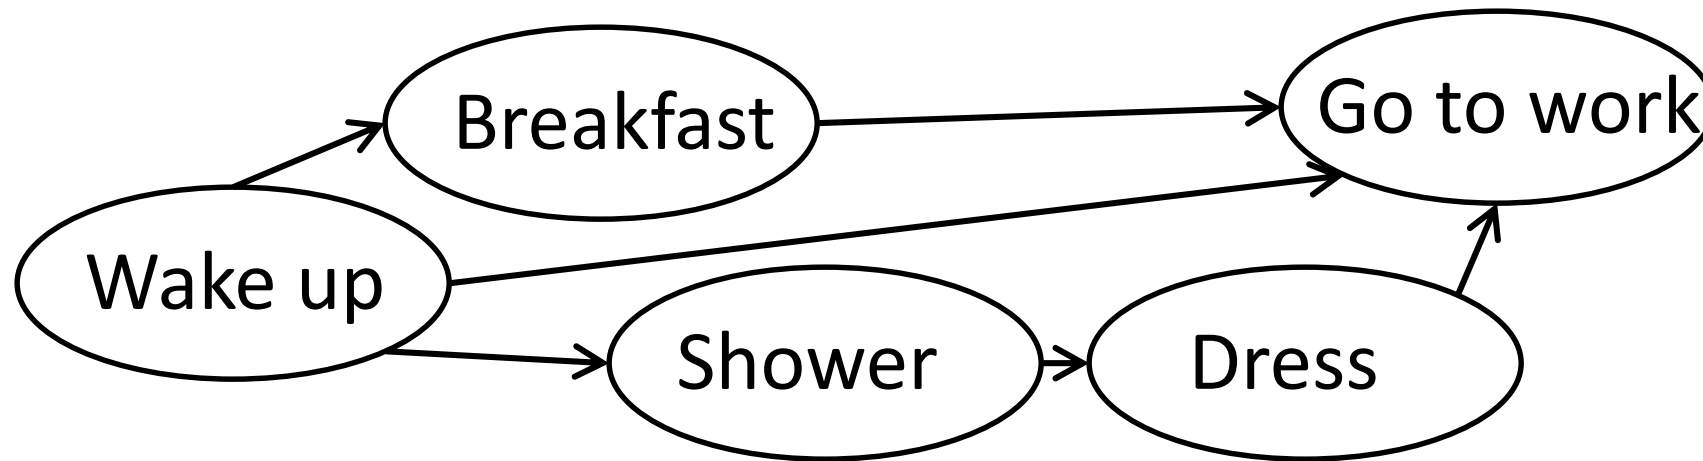
Dependency Graphs



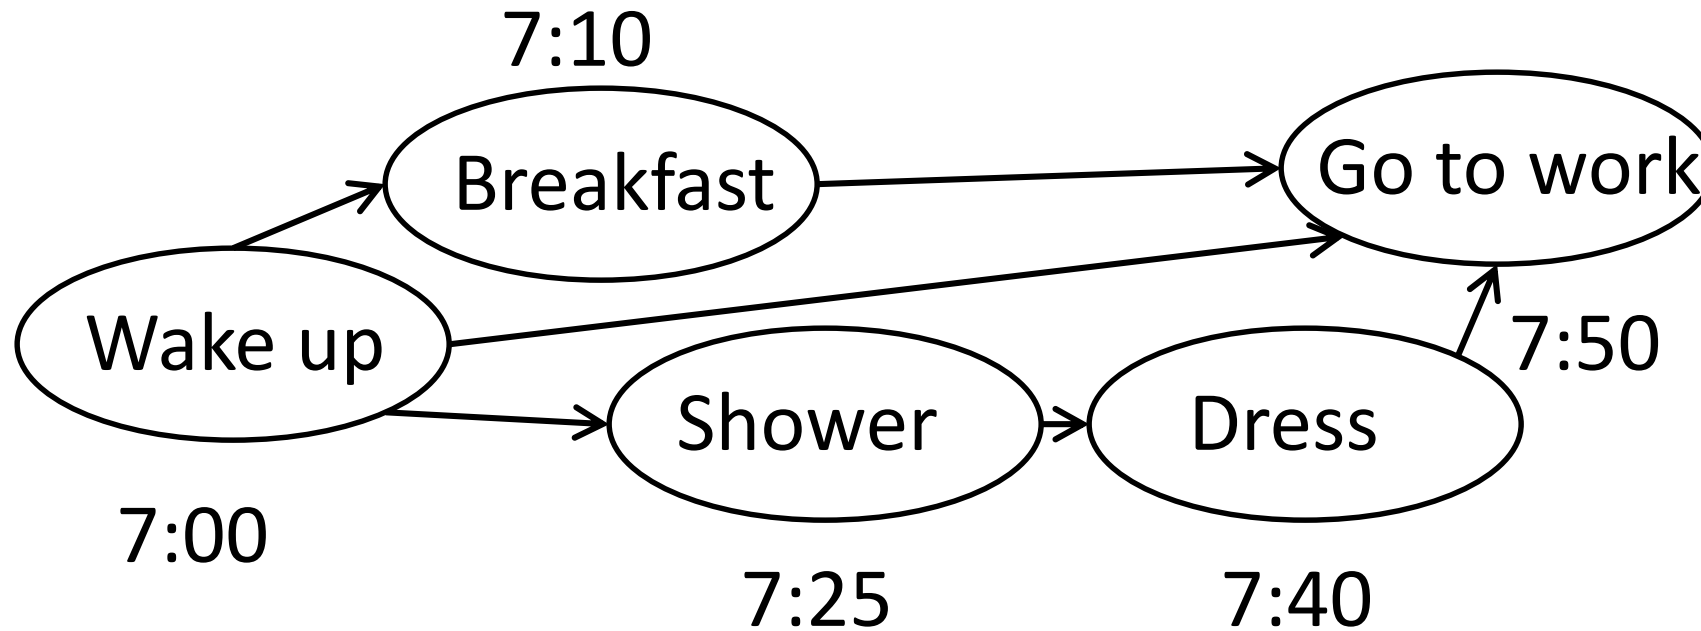
Dependency Graphs



Dependency Graphs



Dependency Graphs



Dependency Graphs

A directed graph can be thought of as a graph of dependencies.
Where an edge $v \rightarrow w$ means that v should come before w .

Dependency Graphs

A directed graph can be thought of as a graph of dependencies. Where an edge $v \rightarrow w$ means that v should come before w .

Definition: A topological ordering of a directed graph is an ordering of the vertices so that for each edge (v,w) , v comes before w in the ordering.

Question: Existence of Orderings

Does every directed graph have a topological ordering?

A) Yes

B) No

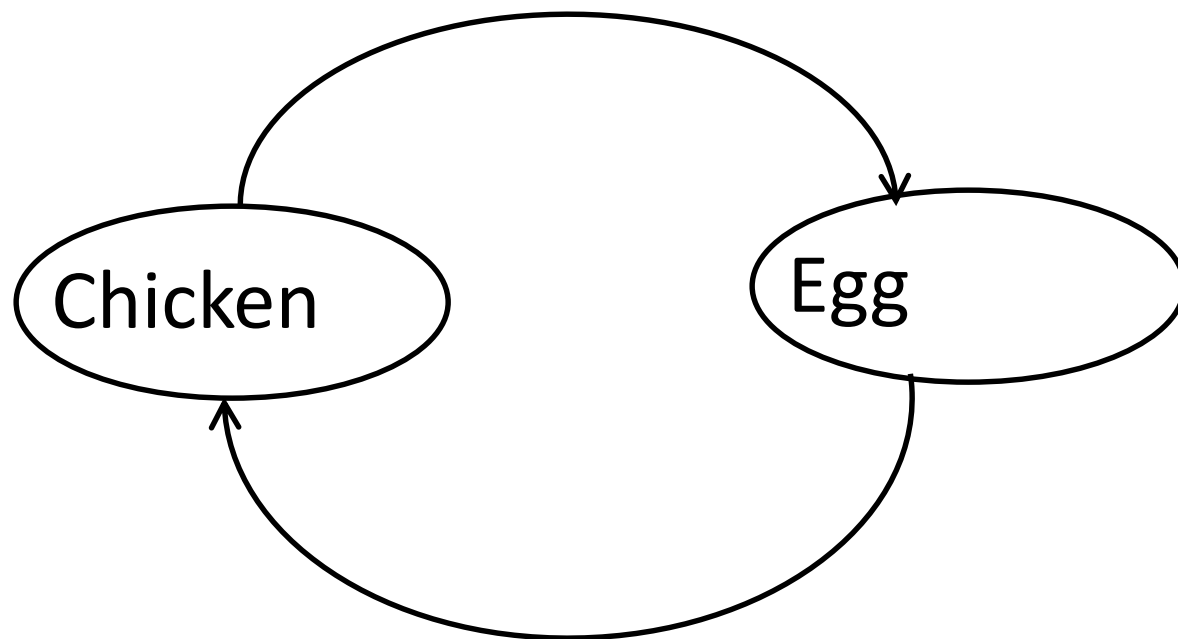
Question: Existence of Orderings

Does every directed graph have a topological ordering?

A) Yes

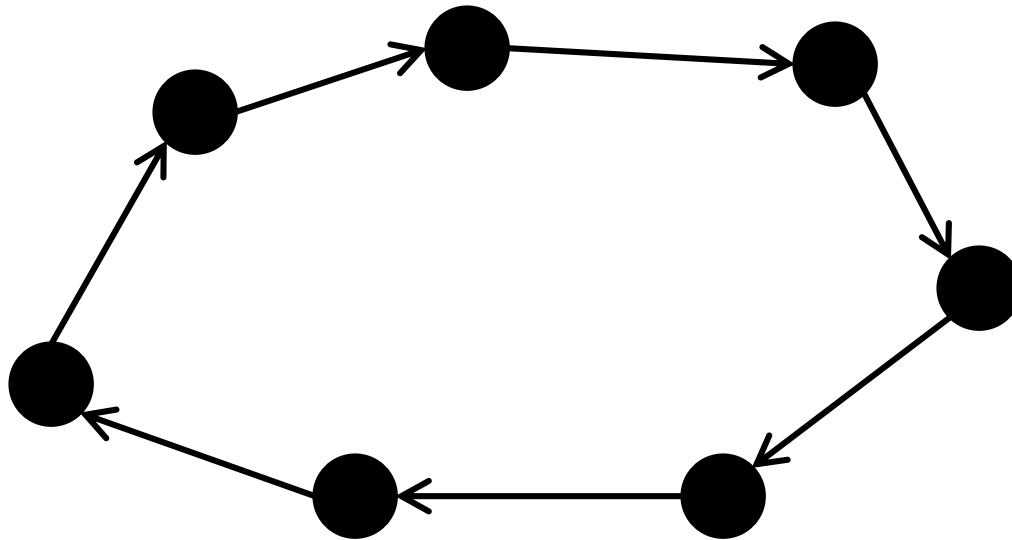
B) No

Counterexample



Cycles

Definition: A cycle in a directed graph is a sequence of vertices $v_1, v_2, v_3, \dots, v_n$ so that there are edges $(v_1, v_2), (v_2, v_3), \dots, (v_{n-1}, v_n), (v_n, v_1)$



Obstacle

Proposition: If G is a directed graph with a cycle, then G has no topological ordering.

Obstacle

Proposition: If G is a directed graph with a cycle, then G has no topological ordering.

Proof:

- Have cycle v_1, v_2, \dots, v_n .

Obstacle

Proposition: If G is a directed graph with a cycle, then G has no topological ordering.

Proof:

- Have cycle v_1, v_2, \dots, v_n .
- AFSOC we have an ordering.

Obstacle

Proposition: If G is a directed graph with a cycle, then G has no topological ordering.

Proof:

- Have cycle v_1, v_2, \dots, v_n .
- AFSOC we have an ordering.
- Find earliest v_i in the ordering.

Obstacle

Proposition: If G is a directed graph with a cycle, then G has no topological ordering.

Proof:

- Have cycle v_1, v_2, \dots, v_n .
- AFSOC we have an ordering.
- Find earliest v_i in the ordering.
- Note that v_i comes before v_{i-1} , in contradiction to the order property.

DAGs

Definition: A Directed Acyclic Graph (DAG) is a directed graph which contains no cycles.

DAGs

Definition: A Directed Acyclic Graph (DAG) is a directed graph which contains no cycles.

Previous result said that *only* DAGs can be topologically ordered.
Is the reverse true?

DAGs

Definition: A Directed Acyclic Graph (DAG) is a directed graph which contains no cycles.

Previous result said that *only* DAGs can be topologically ordered.
Is the reverse true?

Surprisingly, yes.

Existence of Orderings

Theorem: Let G be a (finite) DAG. Then G has a topological ordering.

Existence of Orderings

Theorem: Let G be a (finite) DAG. Then G has a topological ordering.

Proof:

- Consider the *last* vertex in the ordering.

Existence of Orderings

Theorem: Let G be a (finite) DAG. Then G has a topological ordering.

Proof:

- Consider the *last* vertex in the ordering.
 - Must be a *sink* (vertex with no outgoing edges).

Existence of Orderings

Theorem: Let G be a (finite) DAG. Then G has a topological ordering.

Proof:

- Consider the *last* vertex in the ordering.
 - Must be a *sink* (vertex with no outgoing edges).
- Idea: find a sink, put at end, order remaining.

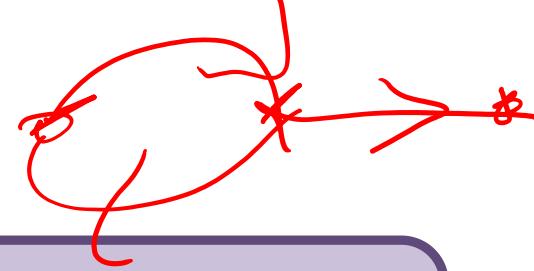
Existence of Orderings

Theorem: Let G be a (finite) DAG. Then G has a topological ordering.

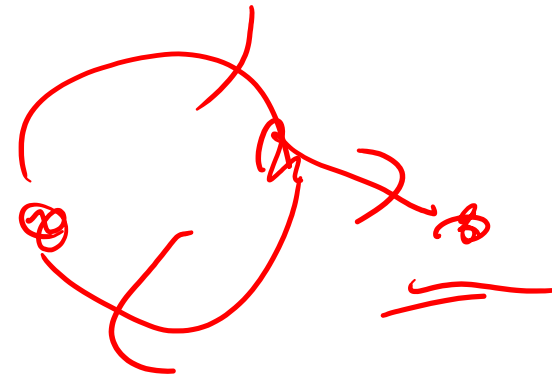
Proof:

- Consider the *last* vertex in the ordering.
 - Must be a *sink* (vertex with no outgoing edges).
- Idea: find a sink, put at end, order remaining.
- Question: Does G have a sink?

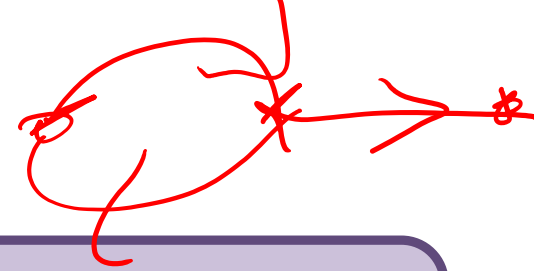
Sinks



Lemma: Every finite DAG contains at least one sink.



Sinks



Lemma: Every finite DAG contains at least one sink.

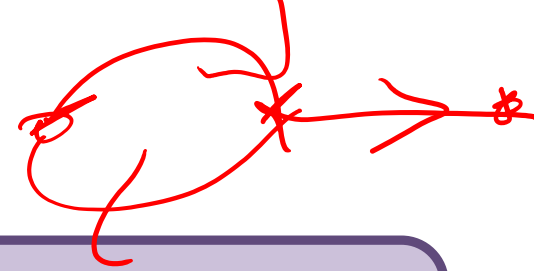
Proof:



- Start at vertex $v = v_1$

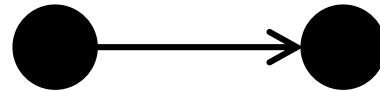


Sinks



Lemma: Every finite DAG contains at least one sink.

Proof:



- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$

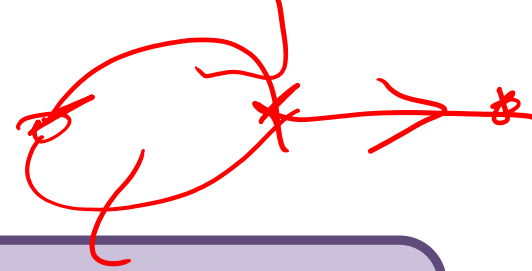
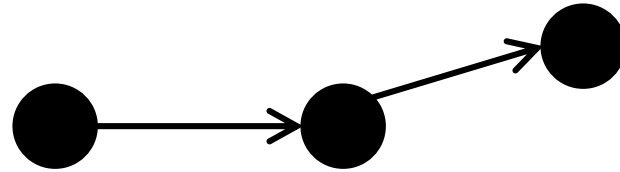


Sinks

Lemma: Every finite DAG contains at least one sink.

Proof:

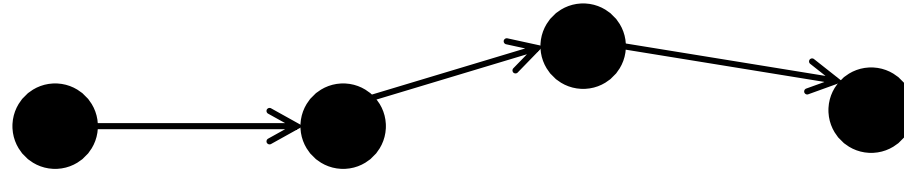
- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$



Sinks

Lemma: Every finite DAG contains at least one sink.

Proof:



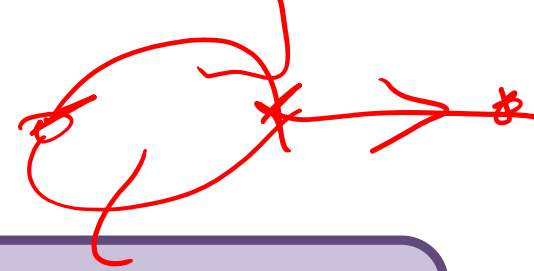
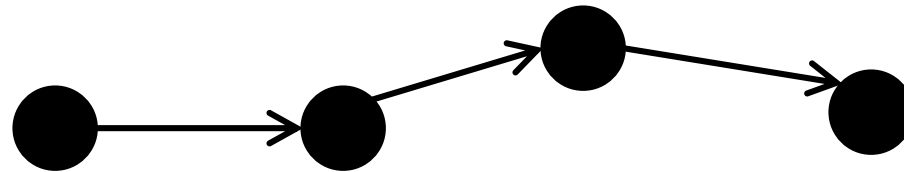
- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$

Sinks

Lemma: Every finite DAG contains at least one sink.

Proof:

- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$
- Eventually either:

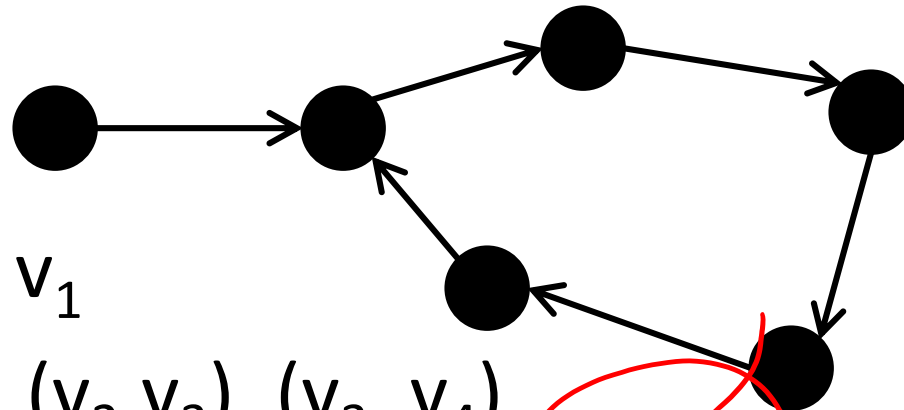


Sinks

Lemma: Every finite DAG contains at least one sink.

Proof:

- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$
- Eventually either:
 - Some vertex repeats (create cycle)

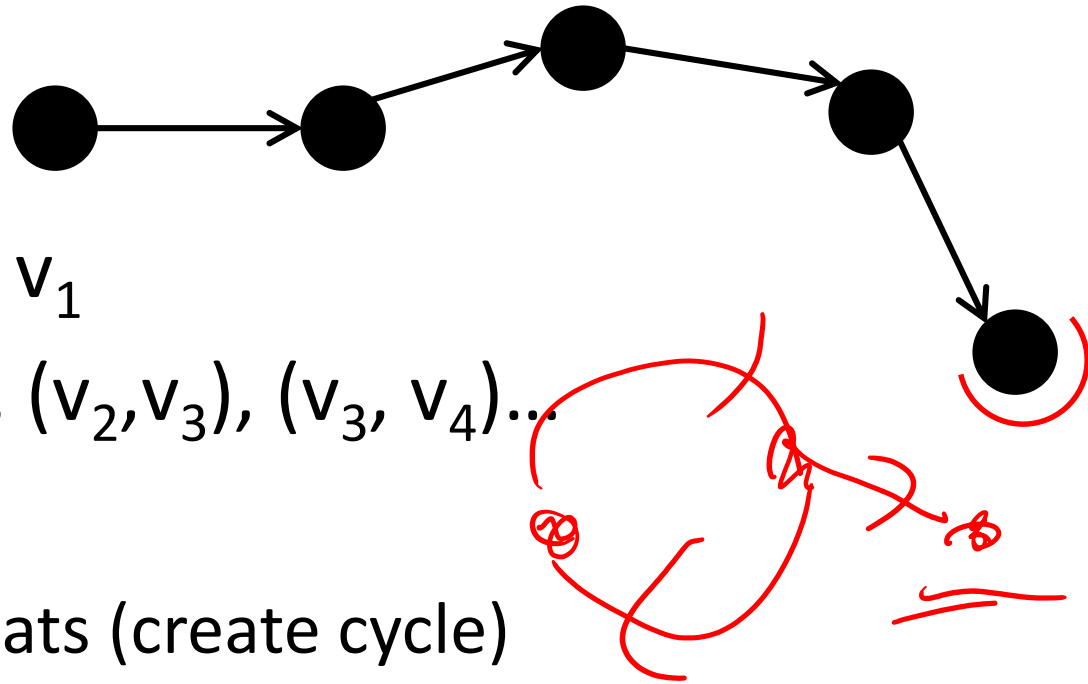


Sinks

Lemma: Every finite DAG contains at least one sink.

Proof:

- Start at vertex $v = v_1$
- Find edges $(v_1, v_2), (v_2, v_3), (v_3, v_4) \dots$
- Eventually either:
 - Some vertex repeats (create cycle)
 - Get stuck (found a sink)



Proof of Theorem

- Induction on $|G|$.

Proof of Theorem

- Induction on $|G|$.
- Find sink v .

Proof of Theorem

- Induction on $|G|$.
- Find sink v .
- Let $G' = G - v$.

Proof of Theorem

- Induction on $|G|$.
- Find sink v .
- Let $G' = G - v$.
- Inductively order G' (still a DAG).

Proof of Theorem

- Induction on $|G|$.
- Find sink v .
- Let $G' = G - v$.
- Inductively order G' (still a DAG).
- Add v to the end of the ordering.

Algorithm

Problem: Design an algorithm that given a DAG G computes a topological ordering on G .