

Cache and Main Memory

## **Review (and more)**

\* Slide title w/o “Review” are new materials

# Review: Handling Cache Hits

---


- Read hits (I\$ and D\$)
  - this is what we want – no challenges
- Write hits (D\$ only) -- **Important!**
  - Method 1: Allow cache and memory to be **inconsistent** → “**write-back**” cache
    - write the data only into the cache block (**write-back** the cache contents to the next level in the memory hierarchy when that cache block is “evicted”)
    - (need a **dirty** bit for each data cache block to tell if it needs to be written back to memory when it is evicted)
  - Method 2: Require the cache and memory to be **consistent**
    - always write the data into both the cache block and the next level in the memory hierarchy (**write-through**) so don’t need a dirty bit
    - writes run at the speed of the next level in the memory hierarchy – slow

# Review:

## Handle writes in write-back cache

### “Write-back Write-allocate”


- On write hit: write-back
- On write miss: write allocate, i.e., copy the old value from a lower level all the way up to L1



# Review:


## Handle writes in write-through cache

### “Write-through Write-no-allocate”



- On write hit: write through
- On write miss: write-no-allocate, i.e., directly write to main memory, rather than bringing the cache block into caches (only bring cache blocks into caches on read misses)

# Review: Concrete Memory Hierarchy




- 0th level: **Registers**
- 1st level: **Primary caches**
  - Split instruction (I\$) and data (D\$)
  - Typically 8KB to 64KB each
- 2nd level: **2<sup>nd</sup> and 3<sup>rd</sup> cache** (L2, L3)
  - On-chip, typically made of SRAM
  - 2<sup>nd</sup> level typically ~256KB to 512KB
  - “Last level cache” typically 4MB to 16MB
- 3rd level: **main memory**
  - Made of DRAM (“Dynamic” RAM)
  - Typically 1GB to 4GB for desktops/laptops
    - Servers can have >1 TB
- 4th level: **disk (swap and files)**
  - Uses magnetic disks or flash drives

# Review: Designing a cache hierarchy

---

- For any memory component:  $t_{\text{hit}}$  vs.  $\%_{\text{miss}}$  tradeoff
- Upper components (I\$, D\$) emphasize low  $t_{\text{hit}}$ 
  - **Frequent access** →  $t_{\text{hit}}$  **important**
  - $t_{\text{miss}}$  is low →  $\%_{\text{miss}}$  less important
  - Lower capacity and lower associativity (to reduce  $t_{\text{hit}}$ )
  - Small-medium block-size (to reduce conflicts)
  - **Split instruction & data cache to allow simultaneous access**
- Moving down (L2, L3) emphasis turns to  $\%_{\text{miss}}$ 
  - **Infrequent access** →  $t_{\text{hit}}$  **less important**
  - $t_{\text{miss}}$  is high →  $\%_{\text{miss}}$  important
  - High capacity, associativity, and block size (to reduce  $\%_{\text{miss}}$ )
  - **Unified insn+data caches to dynamically allocate capacity**

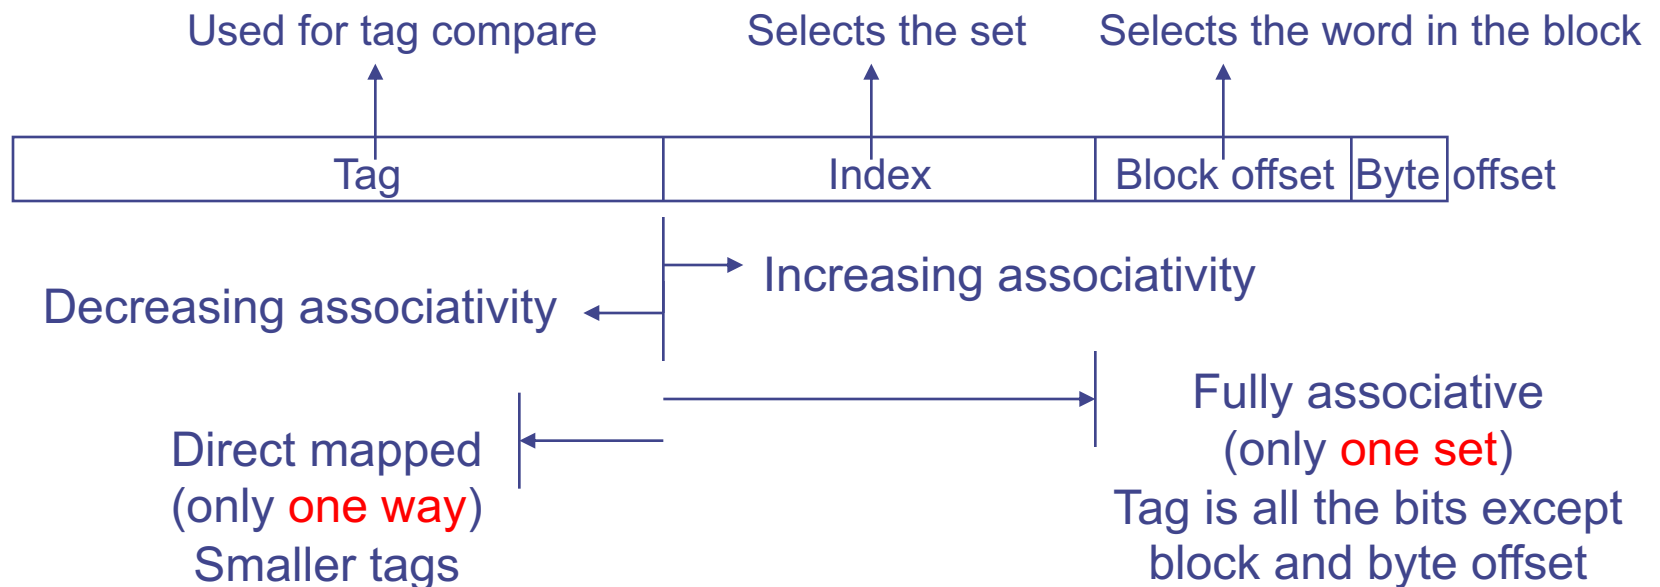
# Review: Hierarchy Performance -- AMAT



$$\begin{aligned} t_{avg} &= t_{avg-M1} \\ &= t_{hit-M1} + (\%_{miss-M1} * t_{miss-M1}) \\ &= t_{hit-M1} + (\%_{miss-M1} * t_{avg-M2}) \\ &= t_{hit-M1} + (\%_{miss-M1} * (t_{hit-M2} + (\%_{miss-M2} * t_{miss-M2}))) \\ &= t_{hit-M1} + (\%_{miss-M1} * (t_{hit-M2} + (\%_{miss-M2} * t_{avg-M3}))) \\ &\dots \end{aligned}$$

# Review: Range of Set Associative Caches

- For a fixed size cache, each increase by a factor of two in associativity doubles the number of blocks per set (i.e., the number or ways) and halves the number of sets – decreases the size of the index by 1 bit and increases the size of the tag by 1 bit





# Review: 3C Model


---

- Divide cache misses into three categories
  - **Compulsory (cold)**: never seen this address before
    - **Would miss even in infinite cache**
  - **Capacity**: miss caused because cache is too small
    - **Would miss even in fully associative cache**
    - Identify? Consecutive accesses to block separated by access to at least N other distinct blocks (N is number of frames in cache)
  - **Conflict**: miss caused because cache associativity is too low
    - Identify? **All other misses**
  - **(Coherence)**: miss due to external invalidations
    - Only in shared memory multiprocessors (later)
- Calculated by multiple simulations
  - Simulate infinite cache, fully-associative cache, normal cache
  - Subtract to find each count

# Review: How to calculate tag overhead

---


- Tag overhead =  $(K / N) * 100\%$



# Review: Main Memory technology -- DRAM

- Dynamic random access memory
- Capacitor charge state indicates stored value

- 1T1C
  - 1 access transistor
  - 1 capacitor
    - Whether the capacitor is charged or discharged indicates storage of 1 or 0



- Capacitor **leaks** through the RC path
  - DRAM cell loses charge over time
  - DRAM cell needs to be **refreshed**

# DRAM Types


| <i>Segment</i> | <i>DRAM Standards &amp; Architectures</i>   |
|----------------|---|
| Commodity      | DDR3 (2007) [14]; DDR4 (2012) [18]  |
| Low-Power      | LPDDR3 (2012) [17]; LPDDR4 (2014) [20]  |
| Graphics       | GDDR5 (2009) [15]   |
| Performance    | eDRAM [28], [32]; RLDram3 (2011) [29]   |
| 3D-Stacked     | WIO (2011) [16]; WIO2 (2014) [21]; MCDRAM (2015) [13];<br>HBM (2013) [19]; HMC1.0 (2013) [10]; HMC1.1 (2014) [11]   |
| Academic       | SBA/SSA (2010) [38]; Staged Reads (2012) [8]; RAIDR (2012) [27];<br>SALP (2012) [24]; TL-DRAM (2013) [26]; RowClone (2013) [37];<br>Half-DRAM (2014) [39]; Row-Buffer Decoupling (2014) [33];<br>SARP (2014) [6]; AL-DRAM (2015) [25] |

Table 1. Landscape of DRAM-based memory


Kim et al., “Ramulator: A Fast and Extensible DRAM Simulator,” IEEE Comp Arch Letters 2015.

# Review: Main memory organization

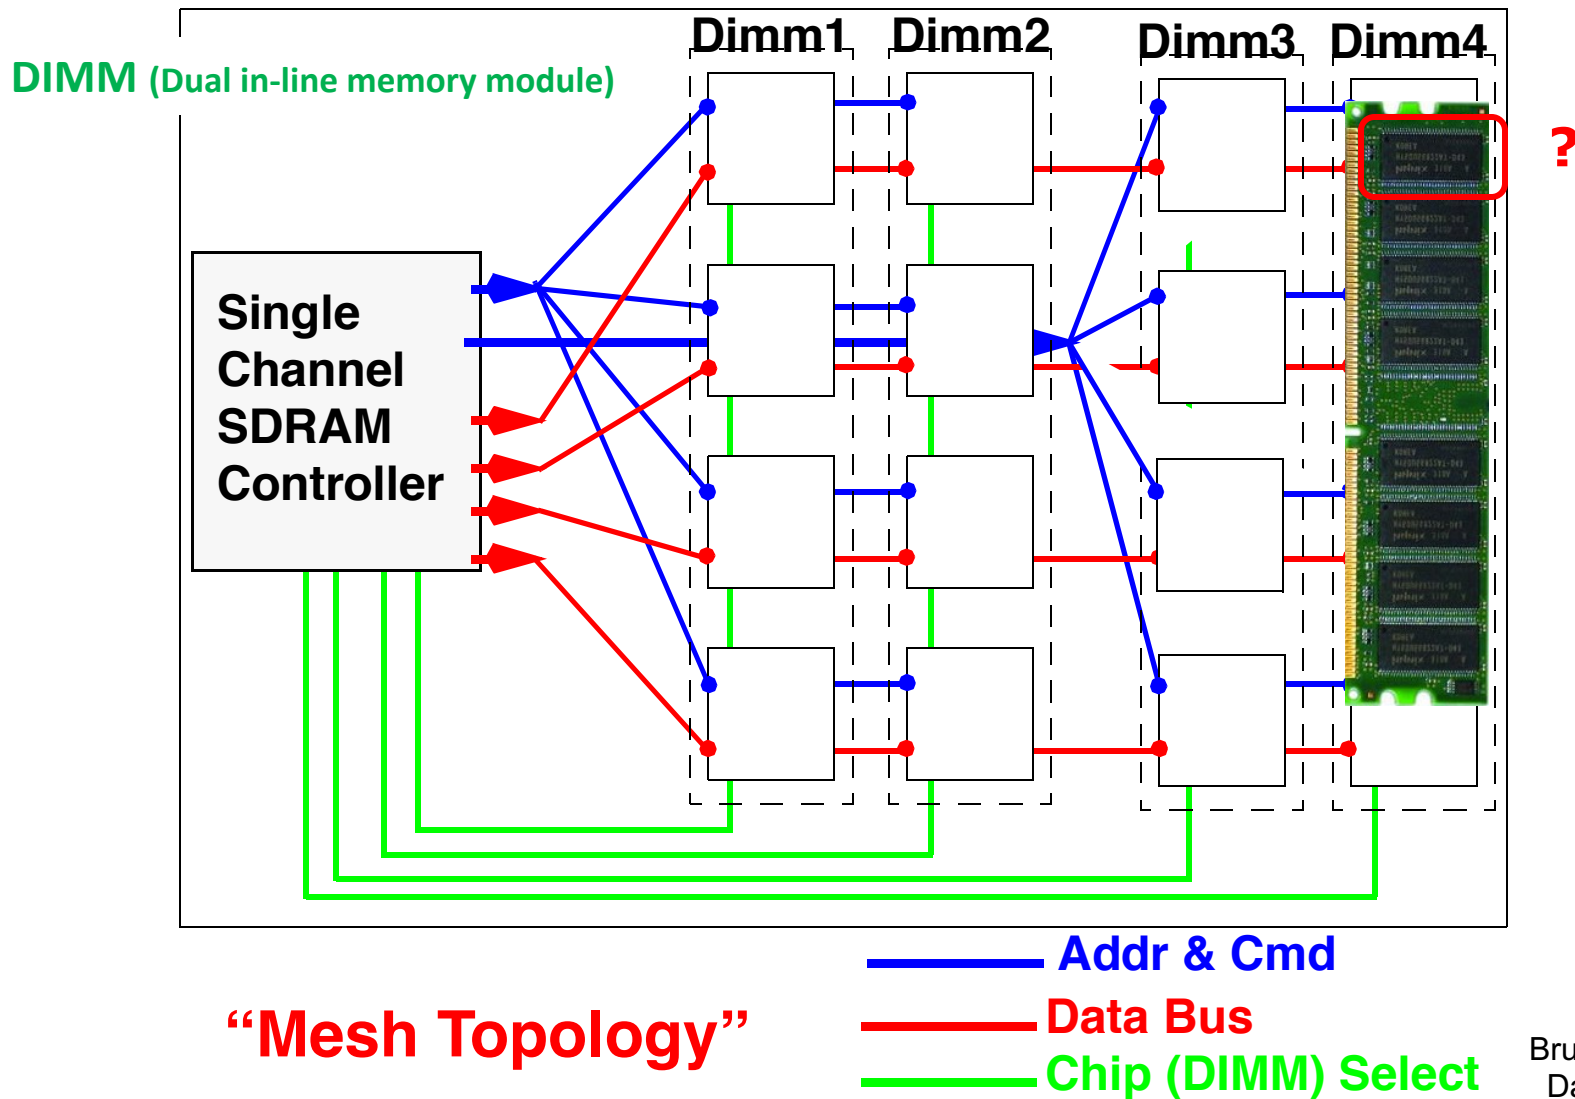
- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/column



# Review: Test yourself



# Review: Test yourself



Source:  
Bruce Jacob &  
David Wang

# Review: Test yourself

## DIMM (Dual in-line memory module)





Side view

SIDE

4.00

Front of DIMM


Back of DIMM



?




# Review: Test yourself



# Review: Rows and columns


Column size = DRAM chip data bus width



# Review: Burst length

---

- Column: Smallest addressable quantity of DRAM chip
  - 1 column: Column size == DRAM chip data bus width (4, 8, 16, 32)
    - (i.e., 4 bits, 8 bits, 16 bits, 32 bits)
  - Get “n” columns per access.  $n = (1, 2, 4, 8)$ 
    - “Burst length” ==  $n$




# Review: Memory bandwidth

---

- Memory performance: Latency -- we already know
- What about memory throughput? → called “memory bandwidth”
  - How many bytes main memory can send per second?
  - Unit: GB/s or MB/s
  - Read bandwidth & write bandwidth
- “Peak memory bandwidth”
  - The maximum bandwidth a main memory device can achieve
  - This is what vendors like to advertise about

# Review: Peak memory bandwidth calculation

Peak memory bandwidth = memory bus width \* memory bus frequency



# Review: Page Mode DRAM

---

- [Known] A DRAM bank is a 2D array of cells:
  - rows x columns
- A “DRAM row” is also called a “DRAM page”
- [Known] “Sense amplifiers” also called “row buffer”
- **A closed row vs. an open row**  
(analogy: write-through vs. write-back cache)
  - A closed row: data are not buffered in the row buffer after each access (similar to “write-through”)
  - An open row: data are buffered in the row buffer after access (similar to “write-back”)
  - Can only have one open row at a time ← only one row buffer in a bank

# Review: Open row vs. closed row

---

- Access to an “open row” – data in the row buffer
  - **Read/write** command reads/writes column in the row buffer
- Access to a “closed row” – data not in the row buffer
  - **Activate** command opens row (placed into row buffer)
  - **Read/write** command reads/writes column in the row buffer
  - **Precharge** command closes the row and prepares the bank for next access

# DRAM Bank Operation (open row policy)

Access Address:  
(Row 0, Column 0)  
(Row 0, Column 1)  
(Row 0, Column 85)  
(Row 1, Column 0)

Row address 0

Row decoder

Columns

Rows

Row 1

Row Buffer ~~CONFLICT!~~

Column address 85

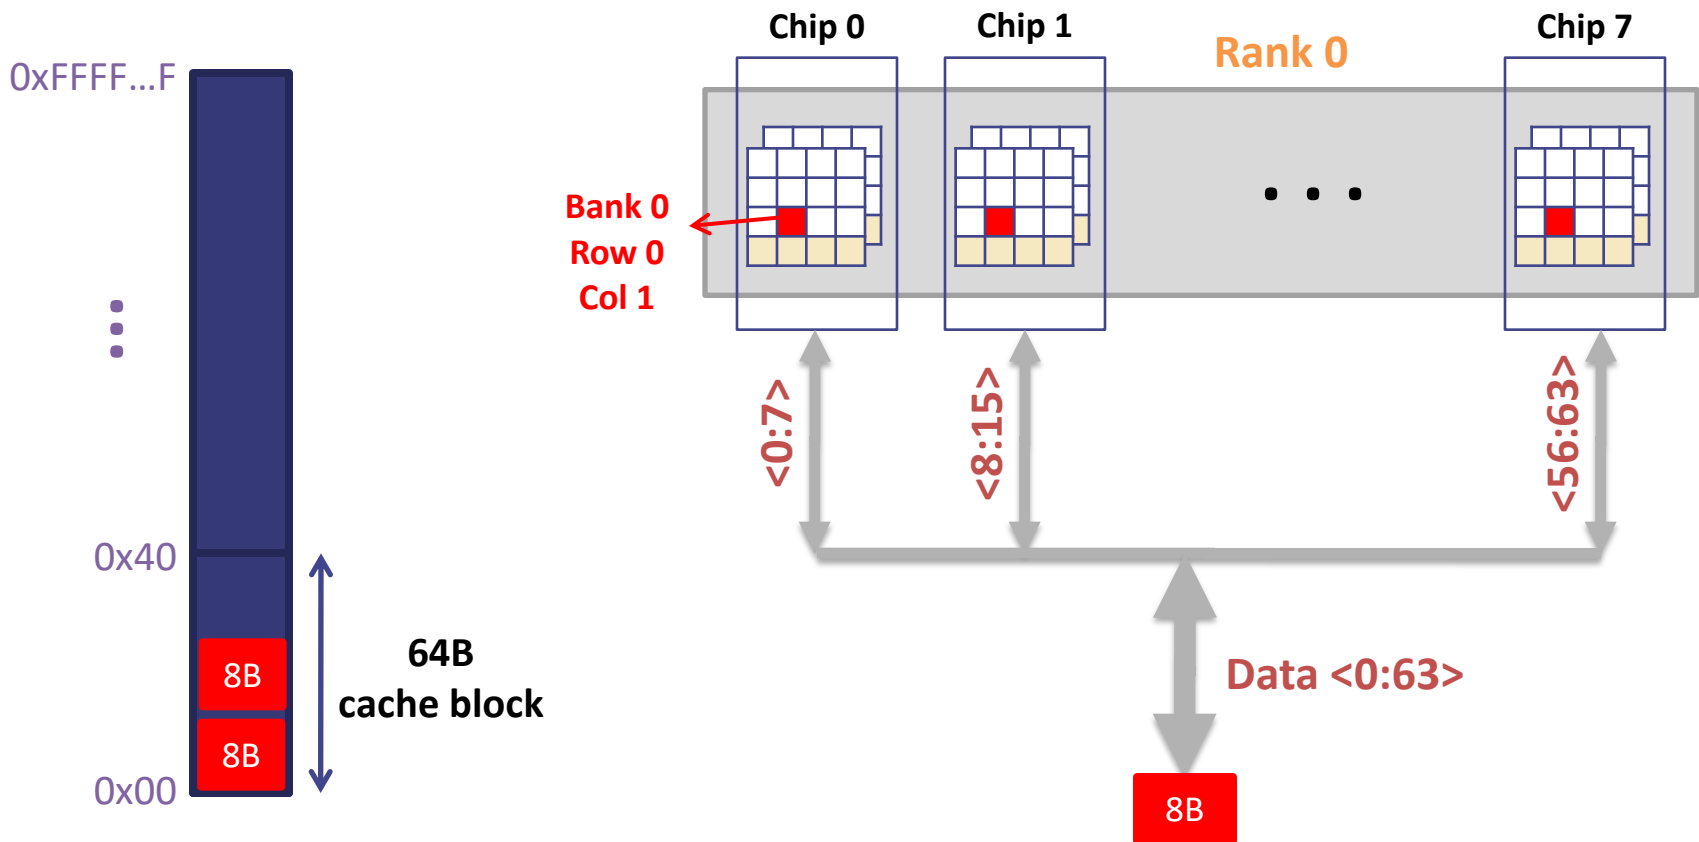
Column mux

Data



# Review: Transferring a cache block

## Physical memory space




**A 64B cache block takes 8 I/O cycles to transfer.  
During the process, 8 columns are read sequentially.**


# Review: Test yourself

---

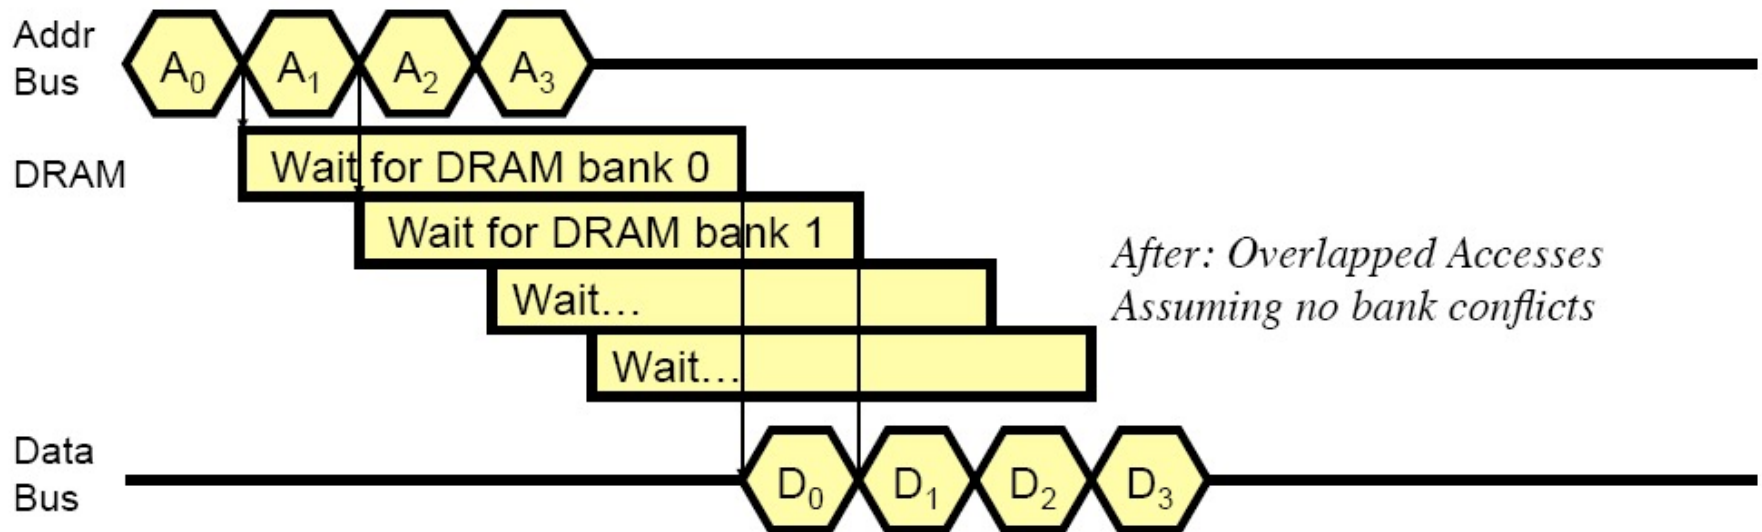
- List the methods you've learned for improving parallelism of memory access
  - Multiple banks
  - Multiple channels



# Review: How Multiple Banks Help

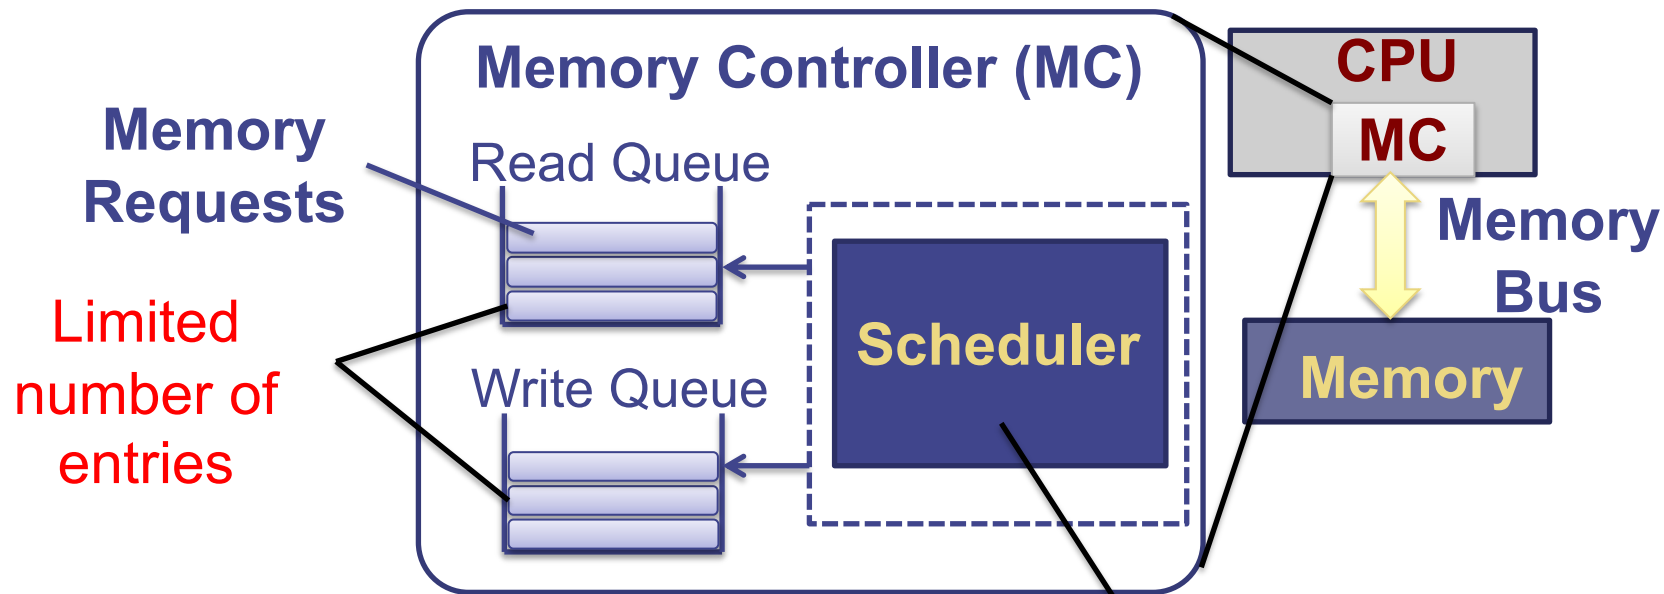


*Before: No Overlapping  
Assuming accesses to different DRAM rows*



*After: Overlapped Accesses  
Assuming no bank conflicts*

# Memory controller

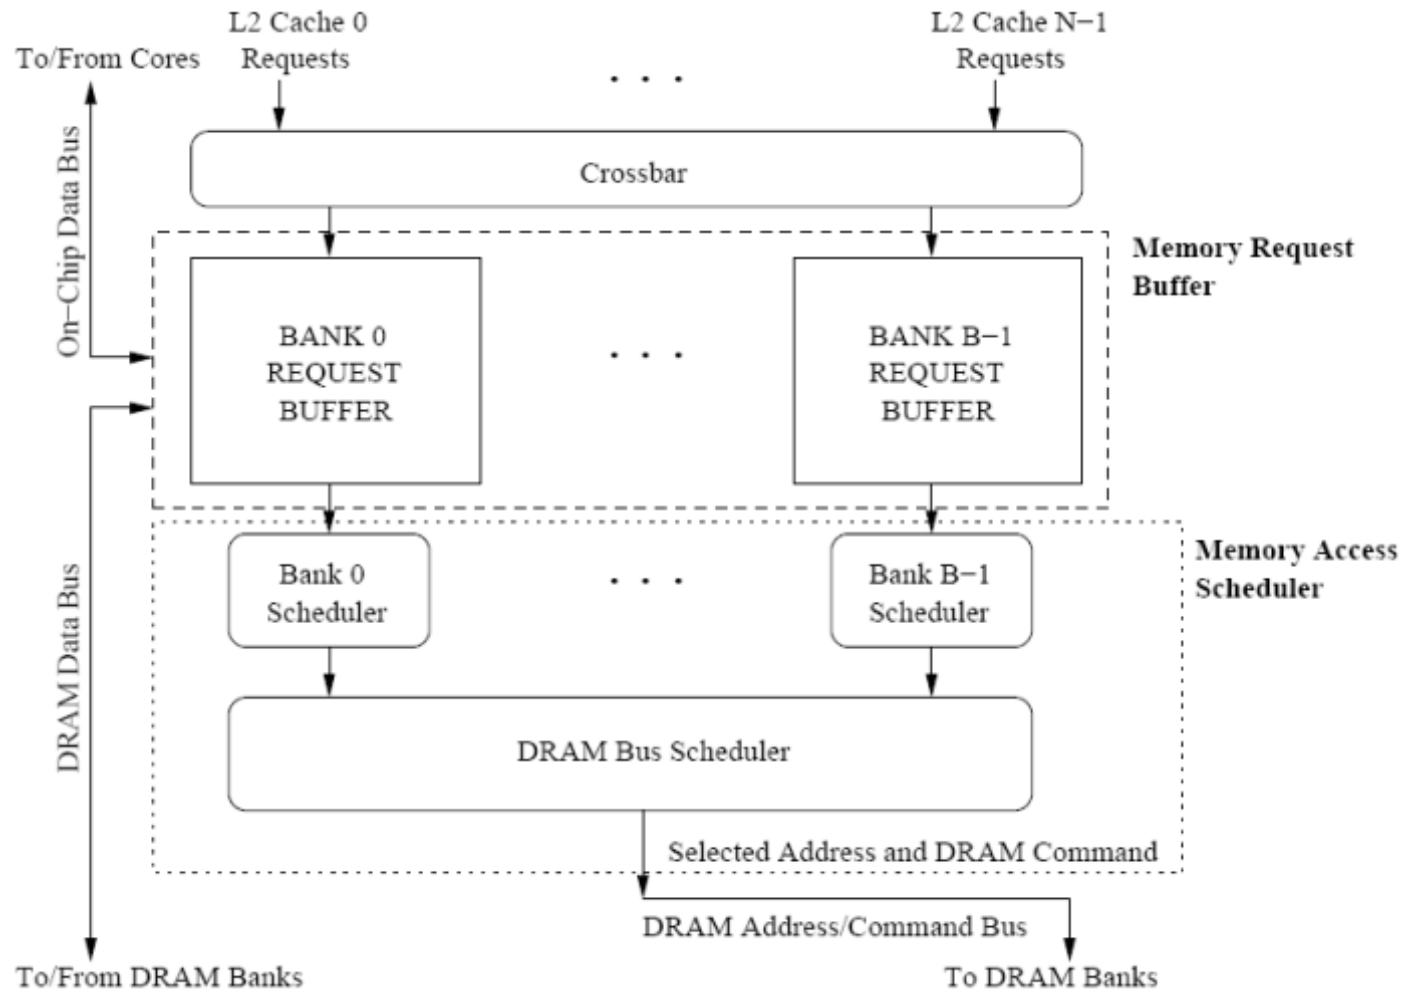


If one of queues overflows, need to **drain** it by stalling current memory service

- Interrupting the on-going memory service
  - Bus turnaround overhead

Determine which requests can be sent on the memory bus to be serviced

# A Modern DRAM Controller



# Scheduling Policy for Single-Core Systems

---


- A row-conflict memory access takes significantly longer than a row-hit access
- Current controllers take advantage of the row buffer
- **FR-FCFS** (first ready, first come first served) scheduling policy
  1. Row-hit first
  2. Oldest first

Goal 1: Maximize row buffer hit rate → maximize DRAM throughput

Goal 2: Prioritize older requests → ensure forward progress

- Is this a good policy in a multi-core system?
  - Mutlu and Moscibroda, Stall-Time Fair Memory Access Scheduling for Chip Multiprocessors, MICRO'07.
  - Mutlu and Moscibroda, Parallelism-Aware Batch Scheduling: Enhancing both Performance and Fairness of Shared DRAM Systems, ISCA'08.

# **How do we encode the address: Address mapping**






# Address mapping

- Single-channel system with 8-byte memory bus
  - 2GB memory, 8 banks, 16K rows & 2048 columns per bank, 64B cache blocks



- Row interleaving


- Ensure consecutive rows of memory mapped to consecutive (i.e., different) banks



- Accesses to consecutive rows serviced in a pipelined manner

- Cache block interleaving

- Ensure consecutive cache block addresses mapped to consecutive banks



- Accesses to consecutive cache blocks serviced in pipelined manner

# Example/Exercise: Address mapping



- Single-channel, 8-byte memory bus, 2GB memory, 8 banks, 16K rows & 2048 columns (i.e., 16K \* 2KB) per bank, 64B cache blocks
  - How many bits needed in address? Answer: 31 bits
  - How many bits needed in address to identify **each byte in one bus transfer**? Answer: 3 bits
  - How many bits needed in address to identify **each bus transfer for transferring one row** of data? Answer: 11 bits
  - How many bits needed in address to identify **each row in one bank**? Answer: 14 bits
  - How many bits needed in address to identify each bank? Answer: 3 bits
- A program has the following physical address access sequence:
  - 0x0000 0001      *Bank 0, Row 0, Col. 0, Byte 1*
  - 0x0000 4001      *Bank 0, Row 1, Col. 0, Byte 1*
  - 0x0000 8002      *Bank 0, Row 2, Col. 0, Byte 2*



## No address mapping

| Bank ID  | Row ID  | Col. ID   | Byte in one bus transfer   |
|--|---|---|--|
| <span style="border: 1px solid black; padding: 2px;">3 bits</span> | <span style="border: 1px solid black; padding: 2px;">14 bits</span> | <span style="border: 1px solid black; padding: 2px;">11 bits</span> | <span style="border: 1px solid black; padding: 2px;">3 bits</span> |

# Example/Exercise: Address mapping

- Single-channel, 8-byte memory bus, 2GB memory, 8 banks, 16K rows & 2048 columns (i.e., 16K \* 2KB) per bank, 64B cache blocks
  - How many bits needed in address? Answer: 31 bits
  - How many bits needed in address to identify **each byte in one bus transfer**? Answer: 3 bits
  - How many bits needed in address to identify **each bus transfer for transferring one row** of data? Answer: 11 bits
  - How many bits needed in address to identify **each row in one bank**? Answer: 14 bits
  - How many bits needed in address to identify each bank? Answer: 3 bits
- A program has the following physical address access sequence:
  - 0x0000 0001      *Row 0, Bank 0, Col. 0, Byte 1*
  - 0x0000 4001      *Row 0, Bank 1, Col. 0, Byte 1*
  - 0x0000 8002      *Row 0, Bank 2, Col. 0, Byte 2*
  - 0x0A00 800A      *Row 0x500, Bank 2, Col. 1, Byte 2*

**No row conflicts!**

## No address mapping

| Bank ID | Row ID  | Col. ID | Byte in one bus transfer |
|---------|---------|---------|--------------------------|
| 3 bits  | 14 bits | 11 bits | 3 bits                   |

## Have address mapping @row interleaving

| Row ID  | Bank ID | Col. ID | Byte in one bus transfer |
|---------|---------|---------|--------------------------|
| 14 bits | 3 bits  | 11 bits | 3 bits                   |

35

# Exercise: Address mapping

- Single-channel, 8-byte memory bus, 2GB memory, 8 banks, 16K rows & 2048 columns (i.e., 16K \* 2KB) per bank, 64B cache blocks
  - How many bits needed in address? Answer: 31 bits
  - How many bits needed in address to identify each byte in one bus transfer? Answer: 3 bits
  - How many bits needed in address to identify each bus transfer for transferring one row of data? Answer: 11 bits
  - How many bits needed in address to identify each row in one bank? Answer: 14 bits
  - How many bits needed in address to identify each bank? Answer: 3 bits
- A program has the following physical address access sequence:
  - Come up with an address sequence yourself to evaluate cache block interleaving

**Have address mapping @cache block interleaving**

| Row ID  | Col. ID-H | Col. ID-L | Byte in one bus transfer |
|---------|-----------|-----------|--------------------------|
| 14 bits | 8 bits    | 3 bits    | 3 bits                   |

↑ Let bank ID start at bit 6  
Bank ID  
3 bits

**No address mapping**

| Bank ID | Row ID  | Col. ID | Byte in one bus transfer |
|---------|---------|---------|--------------------------|
| 3 bits  | 14 bits | 11 bits | 3 bits                   |