

What we learned: Idea 1

CPU



GPU



What we learned: Idea 2

Each GPU core has many ALUs



What we learned: Idea 3

Interleave groups of threads




**What we learned:
processing data**

**What about
moving data to processors?**




NVIDIA GTX 680 (2012)

NVIDIA Kepler GK104 architecture

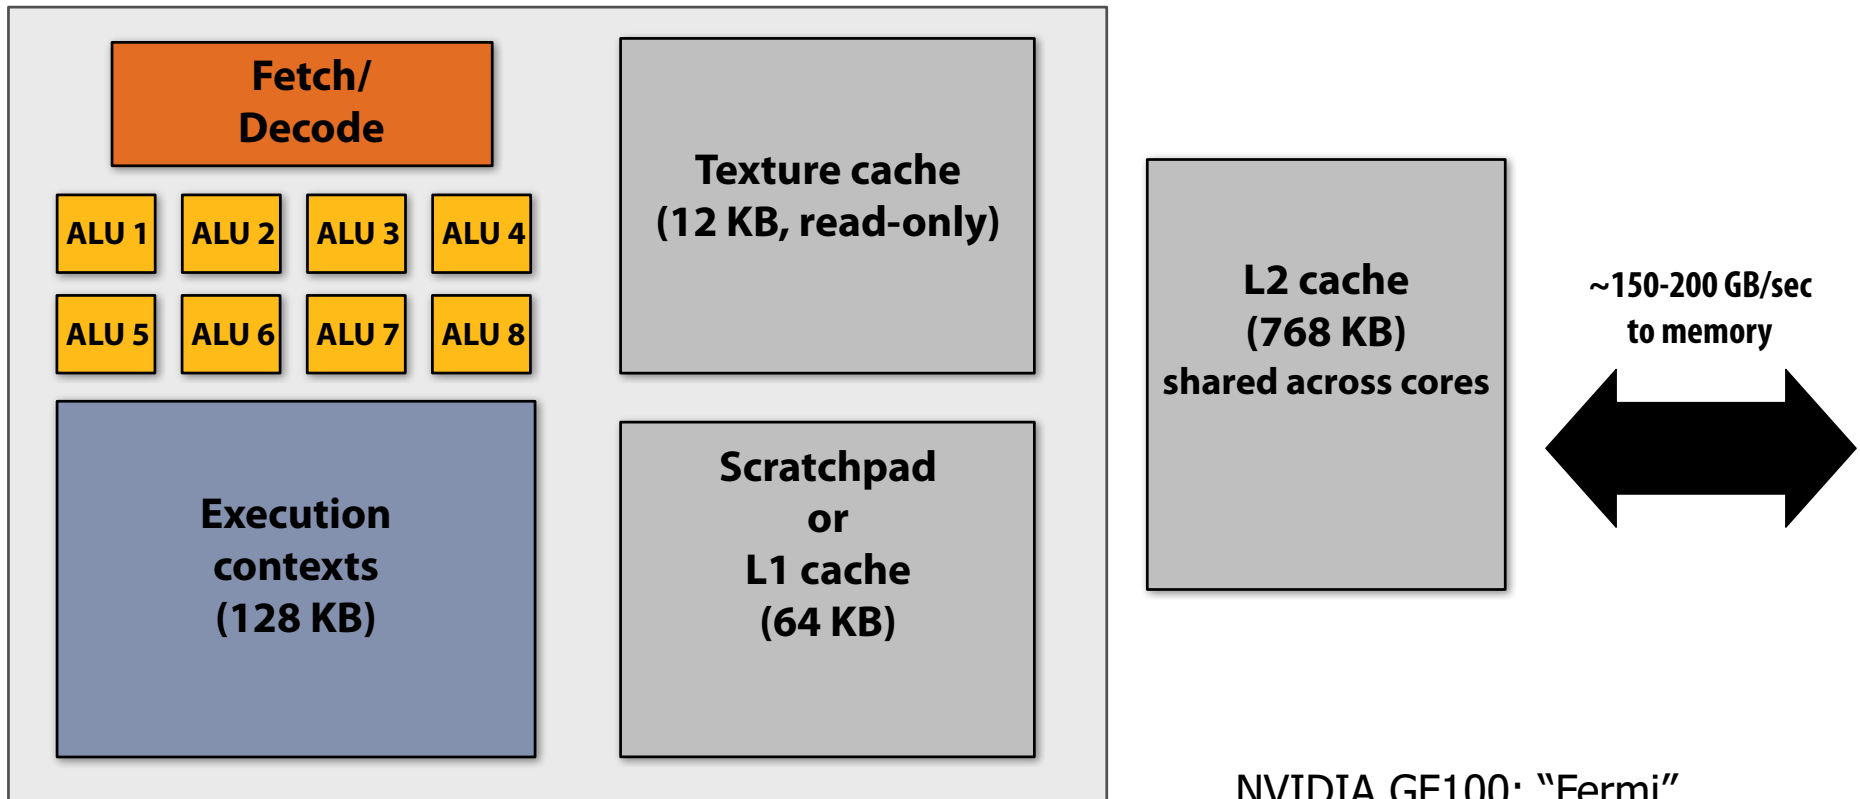


- 1 GHz clock
- Eight SMX cores per chip
- $8 \times 192 = 1,536$ SIMD mul-add ALUs
= 3 TFLOPs
- Up to 512 interleaved warps per chip
(16,384 CUDA threads/chip)
- TDP: 195 watts



"GPU-style" memory hierarchy

Processing Core (many per chip)

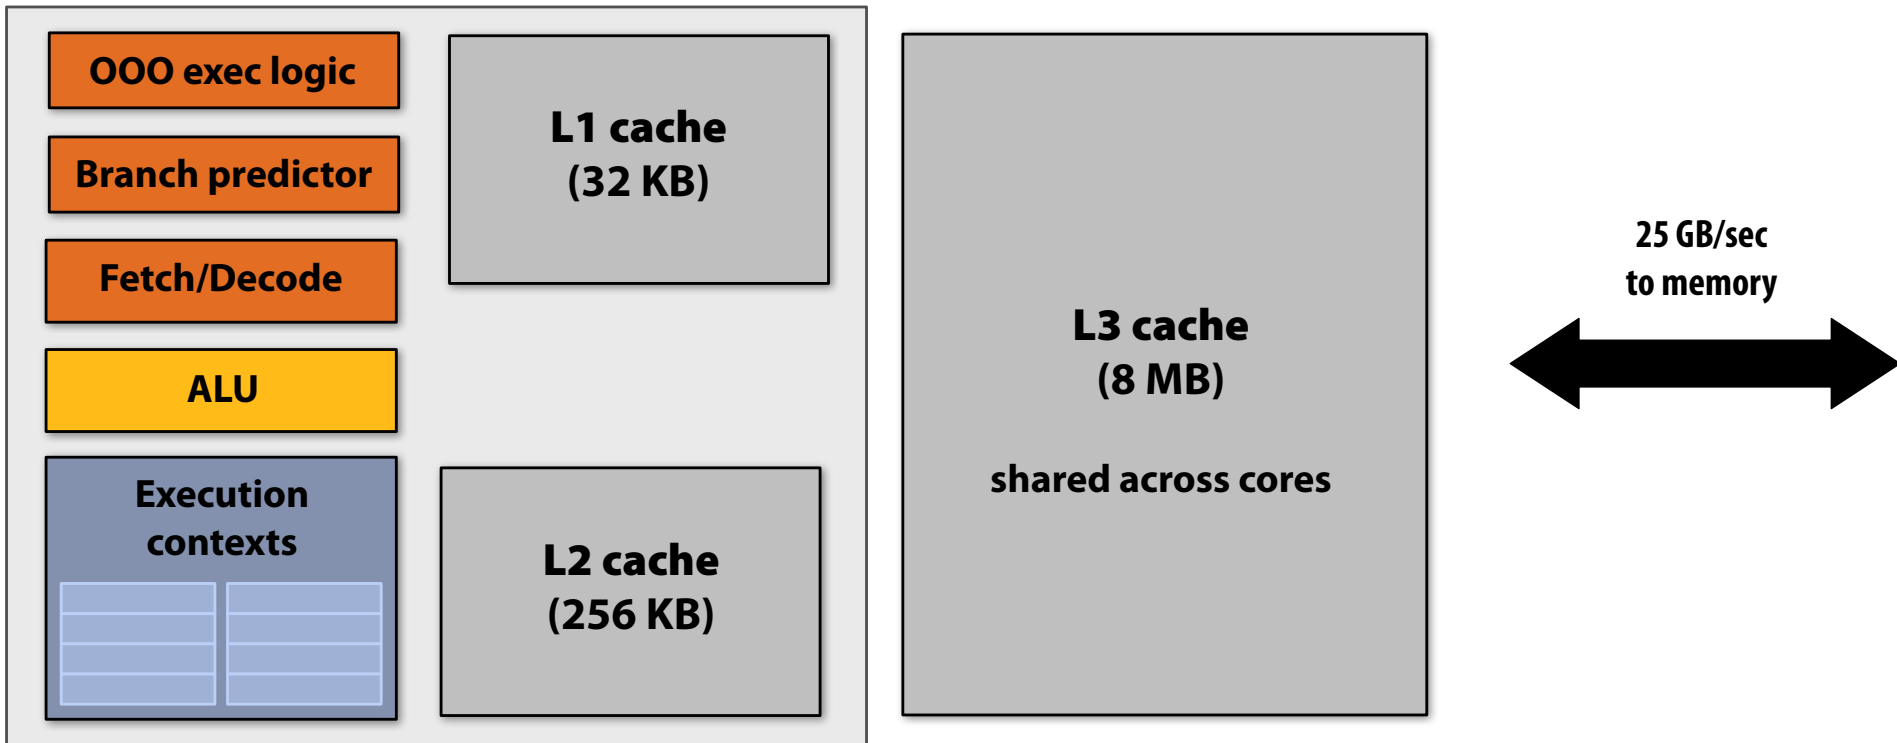


NVIDIA GF100: "Fermi"

More cores, more ALUs, no large traditional cache hierarchy (use threads to tolerate latency) Require high-bandwidth connection to memory

Recall: "CPU-style" memory hierarchy

Processing Core (several per chip)




CPU cores run efficiently when data is resident in cache (caches reduce latency, provide high bandwidth)

GPUs are memory bandwidth hungry

Task: element-wise multiplication of two vectors A and B

1. Load input A[i]
2. Load input B[i]
3. Load input C[i]
4. Compute $A[i] \times B[i] + C[i]$
5. Store result into D[i]



Four memory operations (16 bytes) for every MUL-ADD

Radeon HD 5870 can do 1600 MUL-ADDs per clock

Need ~20 TB/sec of bandwidth to keep functional units busy

Less than 1% efficiency... but 6x faster than CPU!

Bandwidth is a critical resource

- A high-end GPU (e.g., Radeon HD 5870) has...
 - Over **twenty times** (2.7 TFLOPS) the compute performance of quad-core CPU
 - No large cache hierarchy to absorb memory requests
- GPU memory systems are designed for throughput
 - GDDR: today's version is GDDR5
 - Wide memory bus & high memory frequency
 - E.g., 384-bit memory bus
 - Bandwidth can achieve 150-200 GB/sec
 - Still, this is only **six-to-eight times** the bandwidth available to CPU



Bandwidth limited!

If processors request data at too high a rate, the memory system cannot keep up.

No amount of latency hiding helps this.

Overcoming bandwidth limits are a common challenge for application developers on throughput-optimized systems.




Deal with the bandwidth challenge


- Use available bandwidth well
- Fetch data from memory less often (share/reuse data)
- Request data less often (instead, do more math: it's "free")
 - "arithmetic intensity" : ratio of math to data access
- New memory technologies beyond GDDRx
 - HBM (3D/2.5D integrated memory), e.g., AMD Fury X, NVIDIA Pascal and Volta

Idea 1: Using available bandwidth well

Processors generate memory requests



Memory request queue



Memory controller



384 bit memory bus

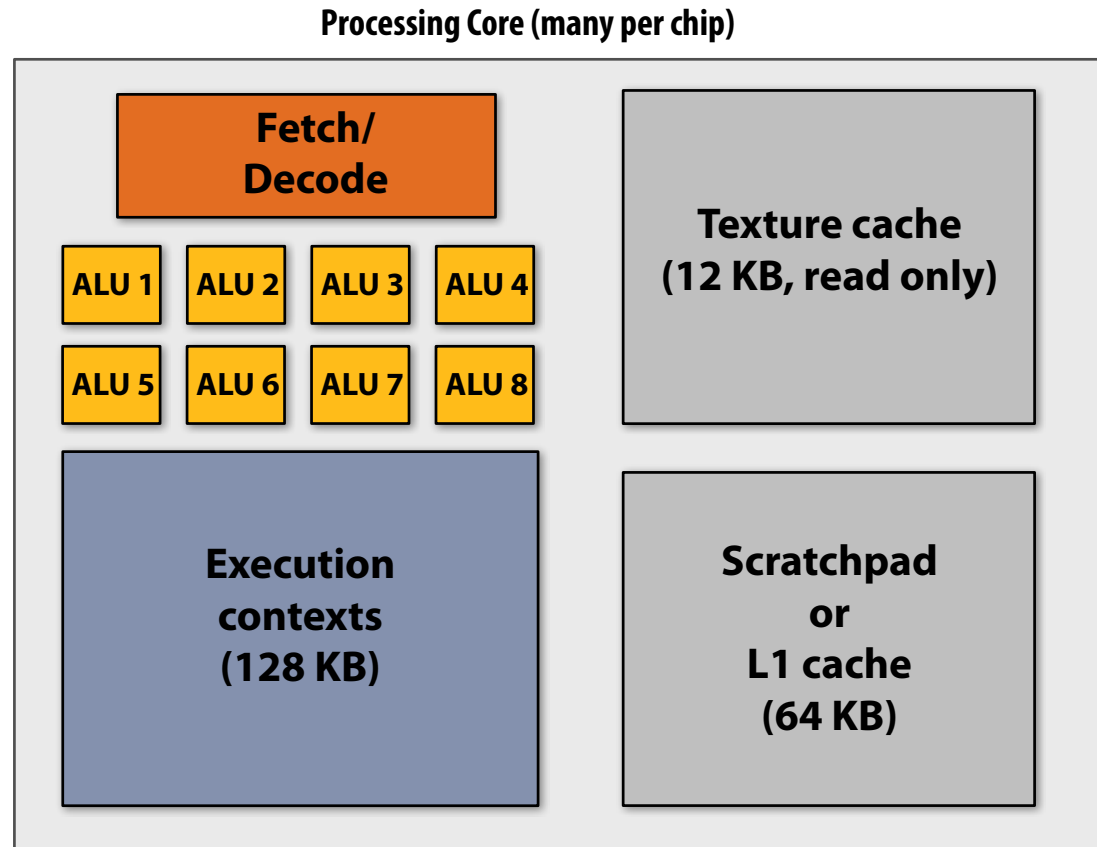
Memory (DDR5)



- GPUs feature sophisticated **memory request reordering** logic
- Repack/reorder/interleave many buffered memory requests to maximize memory utilization

Idea 2: Fetch data from memory less often

- Scratchpad for reuse known at compile-time
 - Intra-fragment reuse
 - Cross-fragment reuse
- Compression



Load-data into scratchpad (LD addr -> scratchpad addr)
Many fragments reuse data loaded into scratchpad once

Deal with the bandwidth challenge

- Use available bandwidth well
- Fetch data from memory less often (share/reuse data)
- Request data less often (instead, do more math: it's "free")
 - "arithmetic intensity" : ratio of math to data access
- New memory technologies



Shading often has high arithmetic intensity

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 ks;  
float shinyExp;  
float3 lightDir;  
float3 viewDir;  
  
float4 phongShader(float3 norm, float2 uv)  
{  
    float result;  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    float spec = dot(viewDir, 2 * dot(-lightDir, norm) * norm + lightDir);  
    result = kd * clamp(dot(lightDir, norm), 0.0, 1.0);  
    result += ks * exp(spec, shinyExp);  
    return float4(result, 1.0);  
}
```



Image credit: <http://caig.cs.nctu.edu.tw/course/CG2007>

3 scalar float operations + 1 exp()
8 float3 operations + 1 clamp()
1 texture access (highlighted in red)

**Vertex processing often has higher arithmetic intensity than
fragment processing (less use of texturing)**

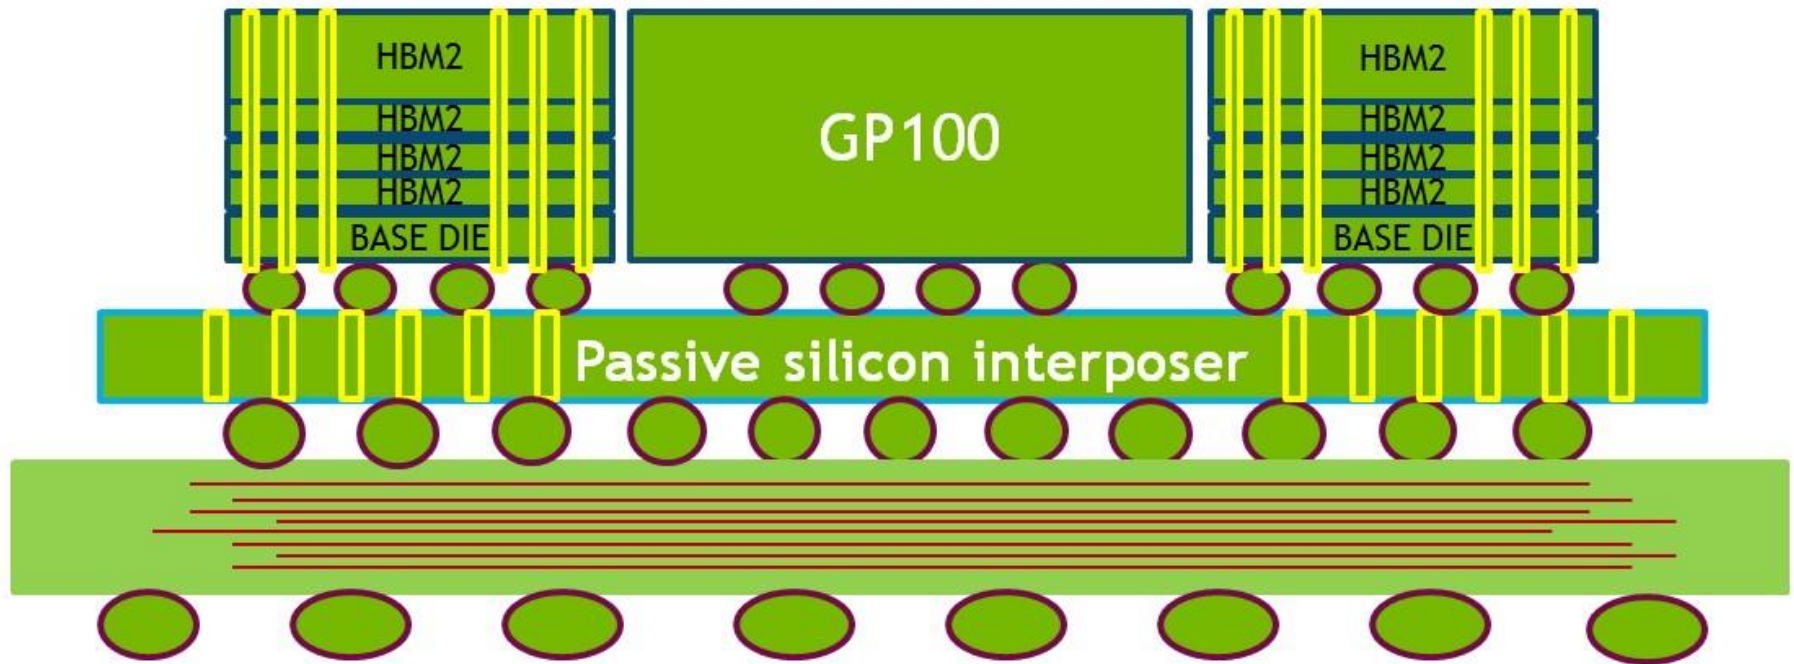


Deal with the bandwidth challenge

- Use available bandwidth well
- Fetch data from memory less often (share/reuse data)
- Request data less often (instead, do more math: it's "free")
 - "arithmetic intensity" : ratio of math to data access
- New memory technologies
 - E.g., 3D integration – main memory can be "on-chip"



Idea 4: NVIDIA Tesla P100 (Pascal 2016)



Some cards feature 16 GB HBM2 in four stacks with a total of 4096bit bus with a memory bandwidth of 720 GB/s

Summary: GPU Memory Design Ideas

- Use available bandwidth well
- Fetch data from memory less often (share/reuse data)
- Request data less often (instead, do more math: it's "free")
 - "arithmetic intensity" : ratio of math to data access
- New memory technologies
 - E.g., 3D integration – main memory can be "on-chip"

