

Midterm Examination

Instructor: Jishen Zhao

Due on: May 3, 2024 @ 11:59PM (48 points)

Name: _____

PID: _____

Email: _____

Q1	4	
Q2	6	
Q3	8	
Q4	14	
Q5	16	
Total	48	

“While taking this examination, I have not witnessed any wrongdoing, nor have I personally violated any conditions of this course’s integrity policy.”

If you can honestly attest to the statement above, **write “I excel with Integrity” below and sign.**

If you do not write and sign, the instructor will contact you by e-mail to request you clarify/explain.

Signature:

Instructions:

- This exam is open book and open notes. Show your work and insert your answer in the space(s) provided. Please provide details on how you reach a result unless directed by the question as not to.
- The exam totals 48 points. It counts for 24% of your course grade. Please submit answers to the following questions as a PDF via Gradescope by May 3, 2024 at 11:59 PM. Policy of late submissions is the same homework assignments. Handwritten or typed responses are accepted.

Q1 (4 points): Single choice and short answer.

- 1A)** What is the primary problem when dealing with branches (such as beq and bne in MIPS) in pipeline designs? (Single choice, no further explanation needed)
- You need to fetch an instruction after the branch before knowing the branch outcome
 - Calculating if the branch should be taken or not taken complicates the design of the Execute (or ALU) stage
 - You can have a branch dependent on a value loaded from memory which delays completing the branch
 - Trick question – branches pose no more problem than any non-branch instruction for a pipeline.
- 1B)** A program P has an instruction count of 10 billion, an average CPI of 3, and runs on a processor with a clock rate of 2 GHZ. What is the execution time for program P? **15s (No further explanation needed)**

Q2 (6 points): We have a program with 30 billion instructions that takes 45 seconds to run on a 2GHz machine. It is given that the program consists of 25% branch instructions, and the CPI of branch instructions is 4.

- 2A)** What is the average CPI of the program?
 $CPI = (45 * 2 * 10^9) / (30 * 10^9) = 3$
- 2B)** Using a newly developed compiler, the recompiled program now uses 20 billion instructions. It is still composed of 25% branch instructions, but the CPI of the branch instructions has been reduced by a factor of 2 (CPI of the other types of instructions remains the same). What is the expected execution time speedup of the new program over the original program (on the same machine)?
 $CPI_{other} = (3 * 30 * 10^9 - 4 * 30 * 10^9 * 25\%) / (30 * 10^9 * 75\%) = 8/3$
 $New\ execution\ time = (20 * 25\% * 10^9 * 4/2 + 20 * 75\% * 10^9 * (8/3)) = 50 * 10^9\ cycles$
 $Original\ execution\ time = 45 * 2 * 10^9 = 90 * 10^9\ cycles$
 $Speedup = 90/50 = 1.8$

Q3 (16 points): Given the following assembly code:

```
L3: addu R7, R4, R3
    lw R7, (R7)
    addu R8, R5, R3
    lw R8, (R8)
    mul R7, R7, R8
    addu R2, R2, R7
    addiu R3, R3, #4
    bne R3, R6, L3
```

Where registers are written as R(reg number), e.g., R1.

In a non-pipelined CPU with the following instruction categories and execution latencies:

Instruction	Latency (Cycles)
add (addu, addiu)	4
load (lw)	8
multiply (mul)	20
branch (bne)	10

- 3A) (4 points)** What is the CPI of the loop for one iteration (the bne instruction included)? **There are a total of 8 instructions in the loop. The percentage breakdown by type is:**
add: 50%
load: 25%
multiply: 12.5%
branch: 12.5%
So the CPI is: $0.5 * 4 + 0.25 * 8 + 0.125 * 20 + 0.125 * 10 = 7.75$
- 3B) (4 points)** Which of the following optimizations would produce a bigger CPI improvement for one iteration (the bne instruction included)?

- Implement prefetching to reduce the latency of loads from 8 cycles to 6 cycles.
- Implement branch prediction to reduce the latency of branch instructions from 10 cycles to 3 cycles.

CPI for prefetching: $0.5 \cdot 4 + 0.25 \cdot 6 + 0.125 \cdot 20 + 0.125 \cdot 10 = 7.25$

CPI for branch prediction: $0.5 \cdot 4 + 0.25 \cdot 8 + 0.125 \cdot 20 + 0.125 \cdot 3 = 6.875$

So, branch prediction results in more CPI improvement

3C) (8 points) How many cycles will one iteration through the loop take? Please fill in the provided pipeline diagram (with D in the stall cycles and arrows showing where bypassing is taking place).

Assumptions:

- The pipeline is an in-order pipeline, i.e., no out-of-order execution
- The pipeline stalls on true data dependencies that cannot be bypassed (i.e., all kinds of data bypassing are available).
- The pipeline also stalls if a structural hazard would occur.
- mult has 4 pipelined execution stages P1, P2, P3, P4, which takes 4 cycles in total; add has a 1-cycle eXecution stage. mult and add use different ALUs.
- mult does not have M stage; all other instructions have M stage.
- bne comparison is done in the Decode stage.

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
L3: addu R7, R4, R3	F	D	X	M	W												
lw R7, (R7)		F	D	X	M	W											
addu R8, R5, R3			F	D	X	M	W										
lw R8, (R8)				F	D	X	M	W									
mul R7, R7, R8					F	D	D	P1	P2	P3	P4	W					
addu R2, R2, R7						F	F	D	D	D	D	X	M	W			
addiu R3, R3, #4								F	F	F	F	D	X	M	W		
bne R3, R6, L3												F	D	D	X	M	W

Q4 (10 points): Branch prediction

For the following questions, assume a branch outcome of

T T N N N T T N N N T T N N N ...

T = taken, N = not-taken

- 4A)** Fill in the table below to determine the prediction accuracy of a 1-bit prediction scheme, where the initial state is not taken. Use a * in the “Incorrect” row (as shown below) to indicate when the Prediction **incorrectly** predicts the branch outcome, while leaving the entry blank if it is a correct prediction.

Branch Outcome	T	T	N	N	N	T	T	N	N	N	T	T	N	N	N	...
State	N	T	T	N	N	N	T	T	N	N	N	T	T	N	N	...
Prediction	N	T	T	N	N	N	T	T	N	N	N	T	T	N	N	...
Incorrect	/		*			*		*			*		*			

What is the steady state prediction accuracy in percentage? **60%**

- 4B)** Fill in the table below to determine the prediction accuracy of a 2-bit, saturating counter prediction scheme where the initial prediction state is strongly not taken (N)? Use a * in the “Incorrect” row (as shown below) to indicate when the Prediction **incorrectly** predicts the branch outcome, while leaving the entry blank if it is a correct prediction.

Branch Outcome	T	T	N	N	N	T	T	N	N	N	T	T	N	N	N	...
State	N	n	t	n	N	N	n	r	n	N	N	n	t	n	N	...
Prediction	N	N	T	N	N	N	N	T	N	N	N	N	T	N	N	...
Incorrect	/	*	*			*	*	*			*	*	*			

What is the steady state prediction accuracy in percentage? **40%**

Q5 (12 points):Pipelining

Assume the following program is running on the 5-stage in-order pipeline processor shown in class. All registers are initialized to 0. Assuming only WX and WD (register file internal forwarding) forwarding, branches are resolved in **Decode** stage, and branches are always predicted **not-taken**. How many cycles will it take to execute the program, if the branch outcome is **actually taken**? Draw a pipeline diagram (table) to show the details of your work. Use arrows to indicate forwarding.

```

lw $r6 0($r10)
lw $r7 0($r11)
add $r2 $r6 $r7
beq $r2 $r3 label
sub $r6 $r8 $r4
sw $r6 0($r10)
label: lw $r1 0($r2)
      or $r4 $r2 $r1

```

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
lw \$r6 0(\$r10)	F	D	X	M	W										
lw \$r7 0(\$r11)		F	D	X	M	W									
add \$r2 \$r6 \$r7			F	D	D	X	M	W							
beq \$r2 \$r3 label					F	D	D	D	X	M	W				
sub \$r6 \$r8 \$r4								F	&	&	&	&			
sw \$r6 0(\$r10)															
label: lw \$r1 0(\$r2)									F	D	X	M	W		
or \$r4 \$r2 \$r1										F	D	D	X	M	W