

CSE 152A: Computer Vision

Manmohan Chandraker

Lecture 6: Edges and Corners



Overall goals for the course

- Introduce fundamental concepts in computer vision
- Enable one or all of several such outcomes
 - Pursue higher studies in computer vision
 - Join industry to do cutting-edge work in computer vision
 - Gain appreciation of modern computer vision technologies
- Engage in discussions and interaction
- This is a great time to study computer vision!

Course Details

Course details


- Class webpage:
 - <https://cseweb.ucsd.edu/~mkchandraker/classes/CSE152A/Winter2024/>
- Instructor email:
 - mkchandraker@ucsd.edu
- Grading
 - 35% final exam
 - 40% homework assignments
 - 20% mid-term
 - 5% self-study exercise
 - Ungraded quizzes
- Aim is to learn together, discuss and have fun!

Course details


- TAs
 - Nicholas Chua: nchua@ucsd.edu
 - Tarun Kalluri: sskallur@ucsd.edu
 - Sreyas Ravichandran: srravichandran@ucsd.edu
- Tutors
 - Kun Wang, Kevin Chan, Zixian Wang: [{kuw010, tsc003, ziw081}@ucsd.edu](mailto:{kuw010,tsc003,ziw081}@ucsd.edu)
- Discussion section: M 3-3:50pm
- TA office hours and tutor hours to be posted on webpage
- Piazza for questions and discussions:
 - <https://piazza.com/ucsd/winter2024/cse152a>

Recap

Photometric: Modeling appearance



Geometric: Modeling projection



- The coordinate system
 - We will use the pinhole model as an approximation
 - Put the optical center (**C**enter **O**f **P**rojection) at the origin
 - Put the image plane (**P**rojection **P**lane) in front of the COP
 - The camera looks down the negative z axis.

Geometric: Camera projection

- To project a point (x,y,z) in *world* coordinates into a camera
- First transform (x,y,z) into *camera* coordinates
- Need to know
 - Camera position (in world coordinates)
 - Camera orientation (in world coordinates)
- Then project into the image plane
 - Need to know camera *intrinsics*
- These can all be described with matrices.

Projection matrix


$$\mathbf{\Pi} = \underset{\text{intrinsics}}{\mathbf{K}} \underset{\text{projection}}{\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}} \underset{\text{rotation}}{\begin{bmatrix} \mathbf{R} & \begin{matrix} 0 \\ 0 \\ 0 \end{matrix} \\ 0 & 0 & 0 & 1 \end{bmatrix}} \underset{\text{translation}}{\begin{bmatrix} \mathbf{I}_{3 \times 3} & -\mathbf{c} \\ 0 & 0 & 0 & 1 \end{bmatrix}}$$



$$\mathbf{\Pi} = \mathbf{K} \left[\mathbf{R} \mid -\mathbf{R}\mathbf{c} \right]$$

Denote this by \mathbf{t} .

Projection matrix



Ideal points and the line at infinity

- Consider two parallel lines in the 2D image:

$$ax + by + c = 0$$

$$ax + by + c' = 0$$


- In homogeneous coordinates, the lines are:

$$\mathbf{l} = (a, b, c)^\top \text{ and } \mathbf{l}' = (a, b, c')^\top$$

- Their point of intersection is given by:

$$\mathbf{x}_\infty = \mathbf{l} \times \mathbf{l}' = (c - c')(-b, a, 0)^\top \sim (-b, a, 0)^\top$$


- To de-homogenize involves a division by 0
 - This is a point at “infinity”, called an ideal point
- Which line contains all ideal points $\mathbf{x}_\infty = (x, y, 0)^\top$?
 - Line at infinity: $\mathbf{l}_\infty = (0, 0, 1)^\top$.



This is a 2D image of 3D space, just for visualization. Equations here are for 2D space.

Correspondence estimation

- Motivation: panorama stitching



Extract features
Match features
Align images




CSE 152A, W124: Manmohan Chandrasekaran

[Images: Rick Szeliski]


Feature detection

Local measure of feature uniqueness


- How does the window change when you shift it?
- Shifting the window in *some direction* causes a *big change*



“flat” region:
no change in all
directions



“edge”: large change
perpendicular to the
edge direction

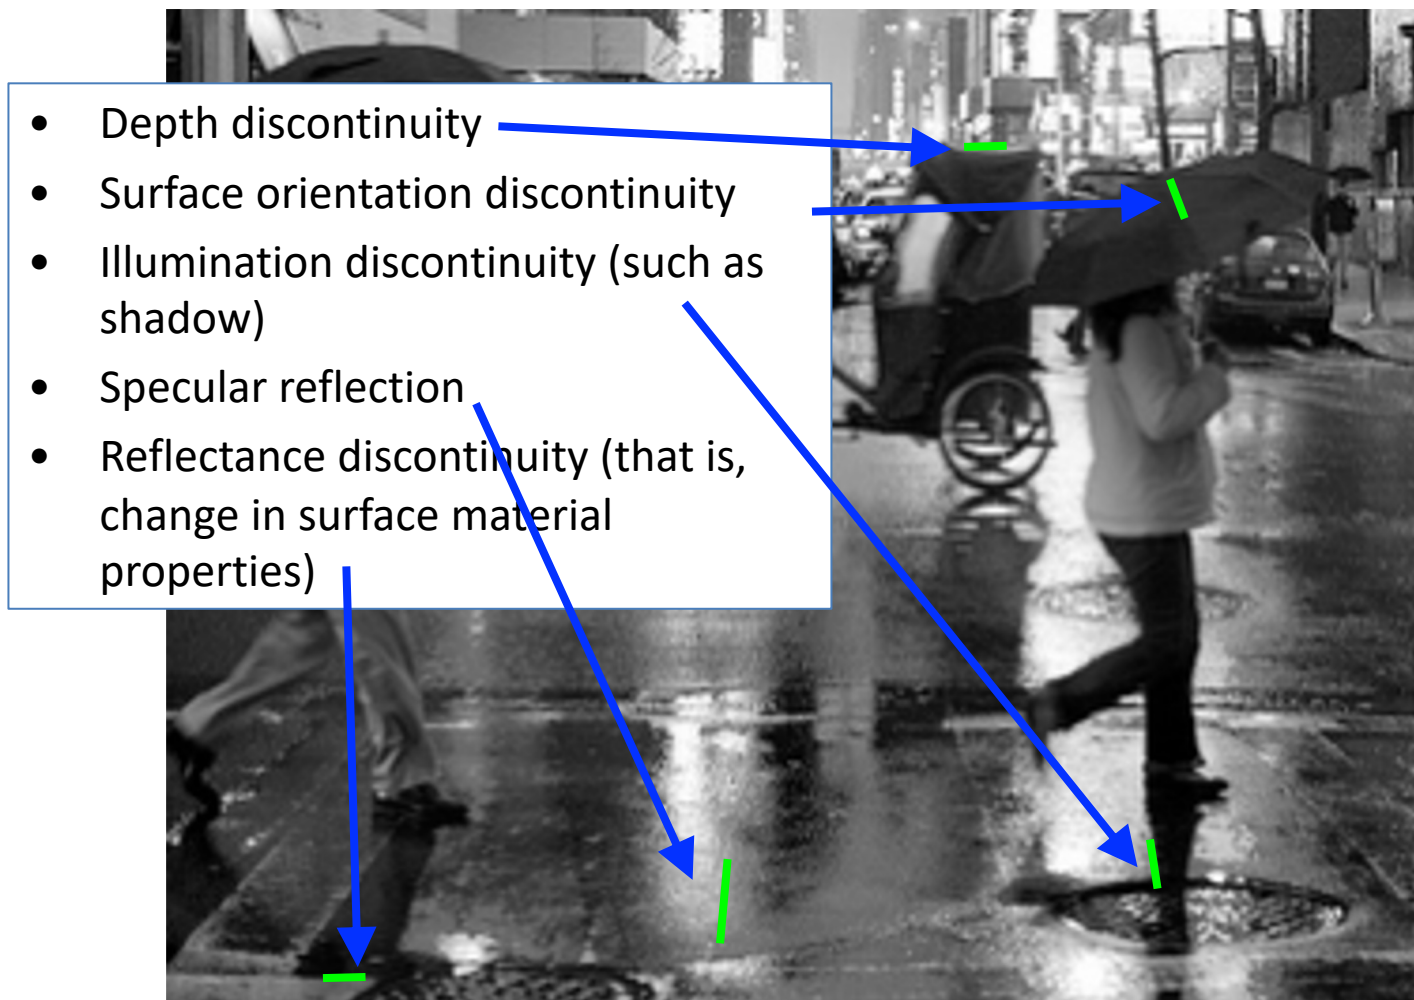


“corner”: large change
in all directions

[Darya Frolova, Denis Simakov, Weizmann Institute]


Edges

Edges in Natural Images



Source: Photografr.com


How Can We Find Edges?



Find regions where magnitude of gradient is large.


Edge is Where Change Occurs: 1-D

- Change is measured by derivative in 1D



- Biggest change: first derivative has maximum magnitude

Numerical Derivatives of Sampled Signal



Take Taylor series expansion of $f(x)$ about x_0

$$f(x) = f(x_0) + f'(x_0)(x-x_0) + \frac{1}{2} f''(x_0)(x-x_0)^2 + \dots$$

Consider samples taken at increments of h and first two terms

$$f(x_0+h) = f(x_0) + f'(x_0)h$$

$$f(x_0-h) = f(x_0) - f'(x_0)h$$

Rearranging the above two yields:

$$f'(x_0) = \frac{f(x_0+h) - f(x_0-h)}{2h} = \begin{bmatrix} -1 & 0 & 1 \\ 2h \end{bmatrix} \cdot [f(x_0-h) \quad f(x_0) \quad f(x_0+h)]$$

Numerical Derivatives

$$f'(x_0) = \frac{f(x_0 + h) - f(x_0 - h)}{2h} = \begin{bmatrix} \frac{-1}{2h} & 0 & \frac{1}{2h} \end{bmatrix} \cdot [f(x_0 - h) \quad f(x_0) \quad f(x_0 + h)]$$

- With images, units of h is pixels, so $h=1$
 - Compute x-derivative at (x_0, y_0) :




$$\frac{I(x_0 + 1, y_0) - I(x_0 - 1, y_0)}{2}$$

- Compute y-derivative at (x_0, y_0) :

$$\frac{I(x_0, y_0 + 1) - I(x_0, y_0 - 1)}{2}$$

Edge Detection with Image Gradients

- Given a function $f(x,y)$, for example, image intensity is f
- Expression for gradient: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- Represents direction of most rapid change in intensity



- Gradient direction: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
- The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Finding derivatives

Is this dI/dx or dI/dy ?



Convolution




Image (I)


*

1	2	1
0	0	0
-1	-2	-1

Kernel (K)

Note: Typically, kernel is relatively small in vision applications.


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$


Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

Convolution: $R = K * I$



Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

Numerical Derivatives

$$f'(x_0) = \frac{f(x_0 + h) - f(x_0 - h)}{2h} = \begin{bmatrix} -1 & 0 & 1 \\ 2h \end{bmatrix} \cdot [f(x_0 - h) \quad f(x_0) \quad f(x_0 + h)]$$

- With images, units of h is pixels, so $h=1$
 - Operator for derivative: $[-1/2 \ 0 \ 1/2]$
- When computing derivatives in the x and y directions, use these operators:

$$\frac{d}{dx} = [-1/2 \quad 0 \quad 1/2] \qquad \frac{d}{dy} = \begin{bmatrix} -1/2 \\ 0 \\ 1/2 \end{bmatrix}$$

- For convolution, can use square kernels

$$\mathbf{K}_x = \frac{1}{2} \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{bmatrix} \qquad \mathbf{K}_y = \frac{1}{2} \begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

Numerical Derivatives

- In practice, Sobel operator is often used
- For image \mathbf{A} , derivative images \mathbf{G}_x and \mathbf{G}_y are

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$


- Can be viewed as smoothing and derivative

$$\mathbf{G}_x = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * ([+1 \quad 0 \quad -1] * \mathbf{A}) \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 \\ 0 \\ -1 \end{bmatrix} * ([1 \quad 2 \quad 1] * \mathbf{A})$$

Aside: Gaussian Smoothing


$$G(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

A 1D Gaussian with mean 0, variance 1



$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

A 2D Gaussian with mean (0,0), variance 1




Aside: Gaussian Smoothing

A discrete approximation to a Gaussian kernel with variance 1

$$\frac{1}{273}$$



1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

Convolution



$$\frac{1}{273}$$


1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1




Corners

Corners more informative than lines

- A point on a line is hard to match




Corners more informative than lines




Corners

- A rapid change of direction in a curve
- A highly effective feature
 - Distinctive, reasonably invariant to viewpoint




Detection of corner-like features

- Examine a small window over an image



The wiggly arrows indicate graphically a directional response in the detector as it moves in the three areas shown

Detection of corner-like features




Intuition:

- Right at corner, gradient is ill-defined.
- Near corner, gradient has two different values.

The Harris corner detector

Compute second-moment matrix:

Sum over a small window
W around hypothetical
corner



Gradient with respect to x,
times gradient with respect to y

$$C = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

Matrix is symmetric

Slide credit: David Jacobs


Simple Case

First, consider the case where:

$$C = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

This means dominant gradient directions align with x or y axis.

If either λ close to 0, then **not** a corner, so seek locations where both large.

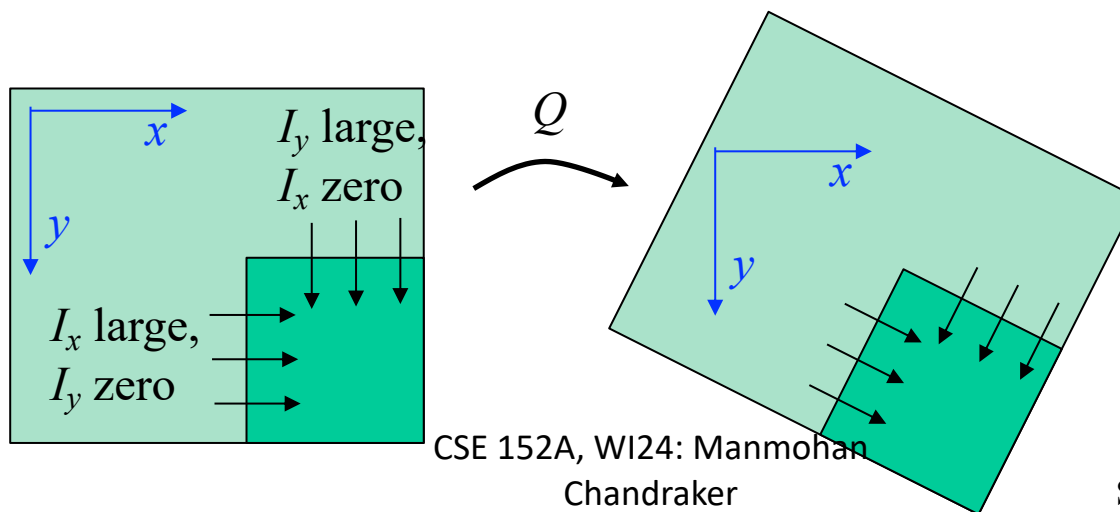


General Case

It can be shown that since C is symmetric:

$$C = Q^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} Q \quad \begin{array}{l} \swarrow \searrow \\ \text{Eigenvalues} \end{array} \quad \leftarrow \text{Rotation}$$

So every case is a rotated version of previous slide.






Simple Corner Detector: Overall Idea

- Smooth image with Gaussian filter to reduce noise
- Compute magnitude of the x and y gradients at each pixel
- Construct C in a window around each pixel
- If λ s are both big, we have a corner.

Simple Corner Detector: Implementation

- Run a small window over an image and compute spatial gradient matrix \mathbf{C} at every pixel
- Compute the minor eigenvalue of \mathbf{C} at every pixel to obtain the corner response “image” R
- Apply nonmaximal suppression to the “image” R
 - Divide into grid, choose maximum within each grid cell
 - Resulting image R' has only one corner candidate per grid cell
 - Prevents corners from being too close to each other
- Threshold resulting image R' using a global threshold T
 - Corners at pixels (x, y) corresponding to $R'(x, y) > T$

Simple Corner Detector: Outputs



Threshold=5,000