

CSE 152A: Computer Vision

Manmohan Chandraker

Lecture 7: Matching



Overall goals for the course

- Introduce fundamental concepts in computer vision
- Enable one or all of several such outcomes
 - Pursue higher studies in computer vision
 - Join industry to do cutting-edge work in computer vision
 - Gain appreciation of modern computer vision technologies
- Engage in discussions and interaction
- This is a great time to study computer vision!

Course Details

Course details


- Class webpage:
 - <https://cseweb.ucsd.edu/~mkchandraker/classes/CSE152A/Winter2024/>
- Instructor email:
 - mkchandraker@ucsd.edu
- Grading
 - 35% final exam
 - 40% homework assignments
 - 20% mid-term
 - 5% self-study exercise
 - Ungraded quizzes
- Aim is to learn together, discuss and have fun!

Course details


- TAs
 - Nicholas Chua: nchua@ucsd.edu
 - Tarun Kalluri: sskallur@ucsd.edu
 - Sreyas Ravichandran: srravichandran@ucsd.edu
- Tutors
 - Kun Wang, Kevin Chan, Zixian Wang: [{kuw010, tsc003, ziw081}@ucsd.edu](mailto:{kuw010,tsc003,ziw081}@ucsd.edu)
- Discussion section: M 3-3:50pm
- TA office hours and tutor hours to be posted on webpage
- Piazza for questions and discussions:
 - <https://piazza.com/ucsd/winter2024/cse152a>

Recap

Photometric: Modeling appearance




Geometric: Modeling projection




Feature detection

Local measure of feature uniqueness


- How does the window change when you shift it?
- Shifting the window in *some direction* causes a *big change*



“flat” region:
no change in all
directions



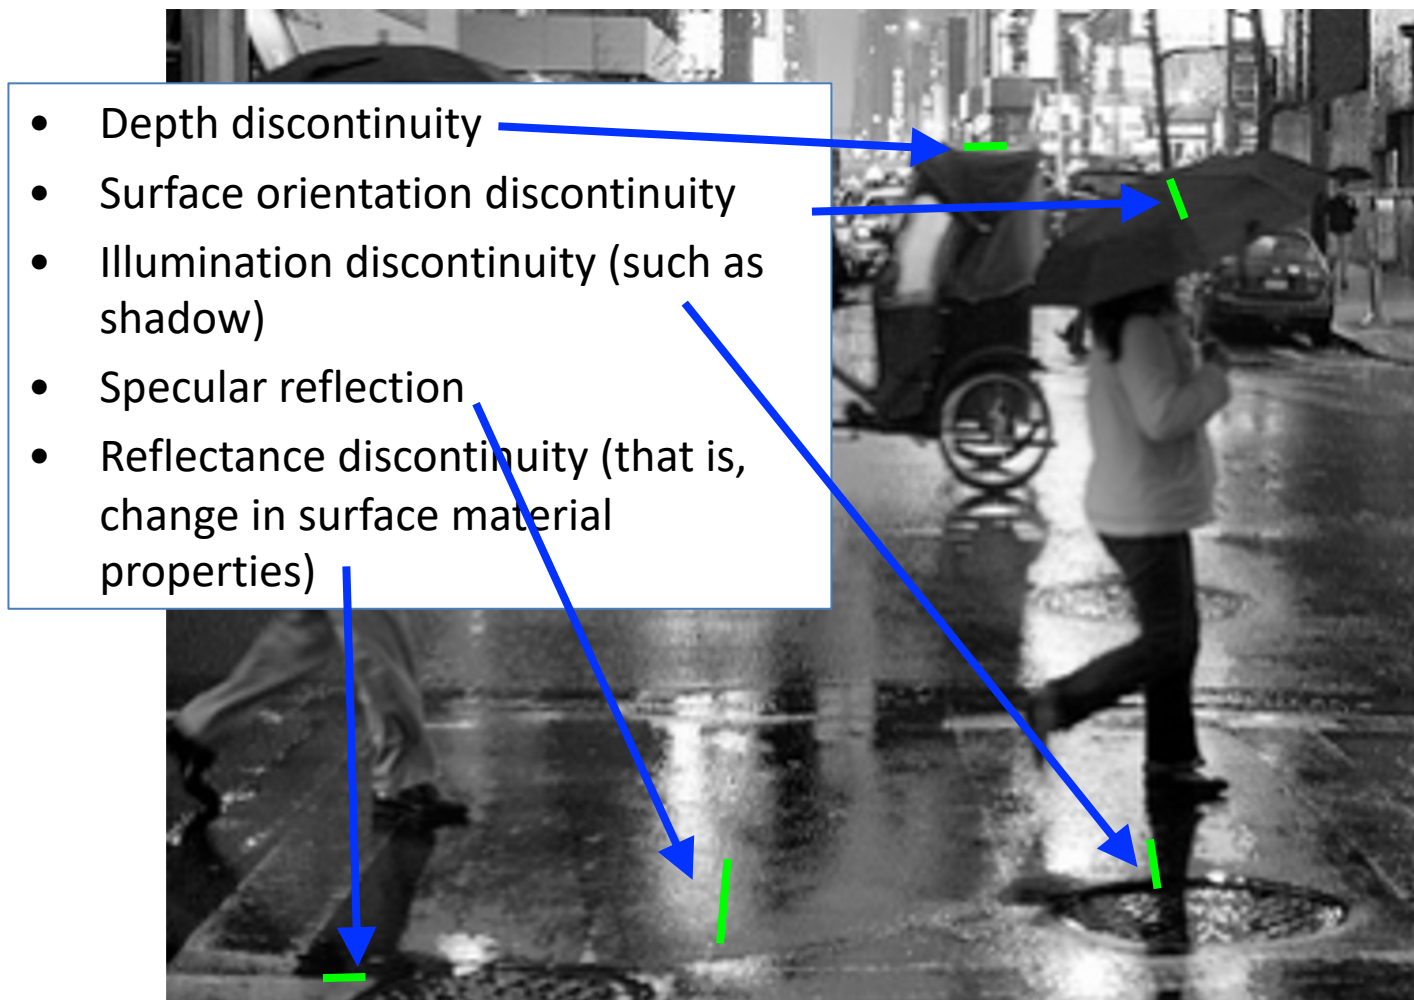
“edge”: large change
perpendicular to the
edge direction



“corner”: large change
in all directions

[Darya Frolova, Denis Simakov, Weizmann Institute]


Edges in Natural Images



Source: Photografr.com

Edge Detection with Image Gradients

- Gradient represents direction of most rapid change in intensity



- The gradient encodes *edge strength* and *edge direction* as

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2} \quad \theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

- Can efficiently compute gradient using convolutions

$$K_x = \frac{1}{2} \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{bmatrix} \quad K_y = \frac{1}{2} \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 0 \\ 0 & -1 & 0 \end{bmatrix}$$

- Sobel operator is often used in practice

$$K_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} \quad K_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Spatial filtering is convolution

Original image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0


Result

4		

$$I(x, y) * h = \sum_{i=-a}^a \sum_{j=-b}^b \underbrace{I(x-i, y-j) \cdot h(i, j)}_{\text{Inner product}}$$

Spatial filtering is convolution

Original image



Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1	1 _{x1}	1 _{x0}	0 _{x1}	0
0	1 _{x0}	1 _{x1}	1 _{x0}	0
0	0 _{x1}	1 _{x0}	1 _{x1}	1
0	0	1	1	0
0	1	1	0	0

Result

4	3	

$$I(x, y) * h = \sum_{i=-a}^a \sum_{j=-b}^b I(x-i, y-j) \cdot h(i, j)$$

Inner product

Spatial filtering is convolution

Original image



Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}

Result


4	3	4
2	4	3
2	3	4

Inner product

$$I(x, y) * h = \sum_{i=-a}^a \sum_{j=-b}^b I(x-i, y-j) \cdot h(i, j)$$

Harris Corner Detector


First, consider the second moment matrix for a simpler case:




Sum over a small window W around hypothetical corner

$$C = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

This means dominant gradient directions align with x or y axis.



In the general case, since C is symmetric, it can be shown:



$$C = Q^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} Q$$




Eigenvalues Rotation

If either λ close to 0, then **not** a corner, so seek locations where both large.

Simple Corner Detector Implementation

- Run a small window over an image and compute spatial gradient matrix \mathbf{C} at every pixel
- Compute the minor eigenvalue of \mathbf{C} at every pixel to obtain the corner response “image” \mathbf{R}
- Apply nonmaximal suppression to the “image” \mathbf{R}
 - Divide into grid, choose maximum within each grid cell
 - Resulting image \mathbf{R}' has only one corner candidate per grid cell
 - Prevents corners from being too close to each other
- Threshold resulting image \mathbf{R}' using a global threshold T
 - Corners at pixels (x, y) corresponding to $\mathbf{R}'(x, y) > T$

Simple Corner Detector: Outputs




Threshold=5,000

Matching

Feature Descriptors

We know how to detect good points


Next question: **How to match them?**



Feature Descriptors

We know how to detect good points

Next question: **How to match them?**




Lots of possibilities (this is a popular research area)

- Simple option: match square windows around the point
- State of the art: SIFT (<http://www.cs.ubc.ca/~lowe/keypoints/>)

Simple matching methods

Interest point:


- Localized position
- Informative about image content
- Repeatable under variations



Simple matching methods


Interest point:

- Localized position
- Informative about image content
- Repeatable under variations




Descriptor:

- Function applied on each W_1 and W_2 , to enable comparing them



$W_1(x, y)$: $k \times k$ pixel patch in image 1


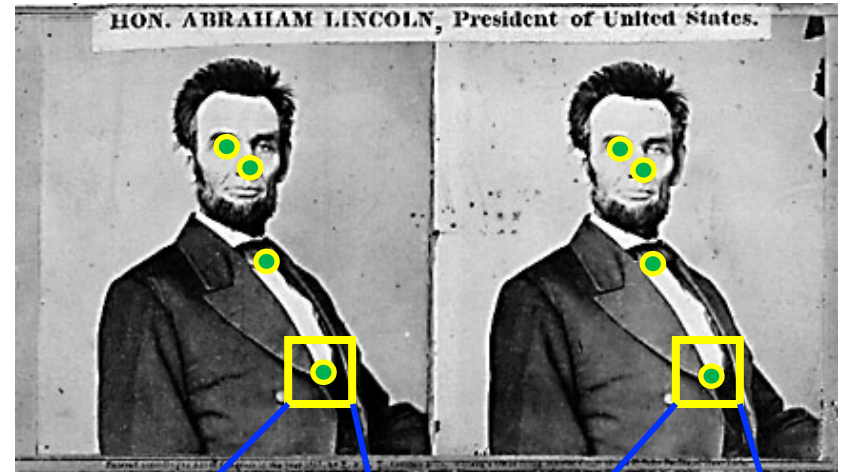


$W_2(x, y)$: $k \times k$ pixel patch in image 2


Simple matching methods

- SSD (Sum of Squared Differences)

$$\sum_{x,y} |W_1(x,y) - W_2(x,y)|^2$$



$W_1(x, y)$: $k \times k$ pixel patch in image 1



$W_2(x, y)$: $k \times k$ pixel patch in image 2

Simple matching methods

- SSD (Sum of Squared Differences)

$$\sum_{x,y} |W_1(x,y) - W_2(x,y)|^2$$


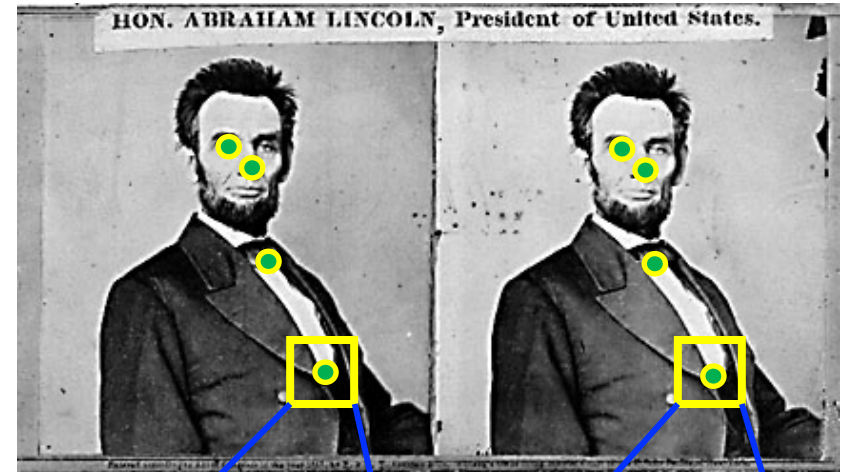
- NCC (Normalized Cross Correlation)

$$\sum_{x,y} \frac{(W_1(x,y) - \overline{W_1})(W_2(x,y) - \overline{W_2})}{\sigma_{W_1} \sigma_{W_2}}$$


$$\overline{W_i} = \frac{1}{n} \sum_{x,y} W_i, \quad \sigma_{W_i} = \sqrt{\frac{1}{n} \sum_{x,y} (W_i - \overline{W_i})^2}$$

(Mean) (Standard deviation)

- What advantages might NCC have over SSD?



$W_1(x,y)$: $k \times k$ pixel patch in image 1



$W_2(x,y)$: $k \times k$ pixel patch in image 2

Feature matching

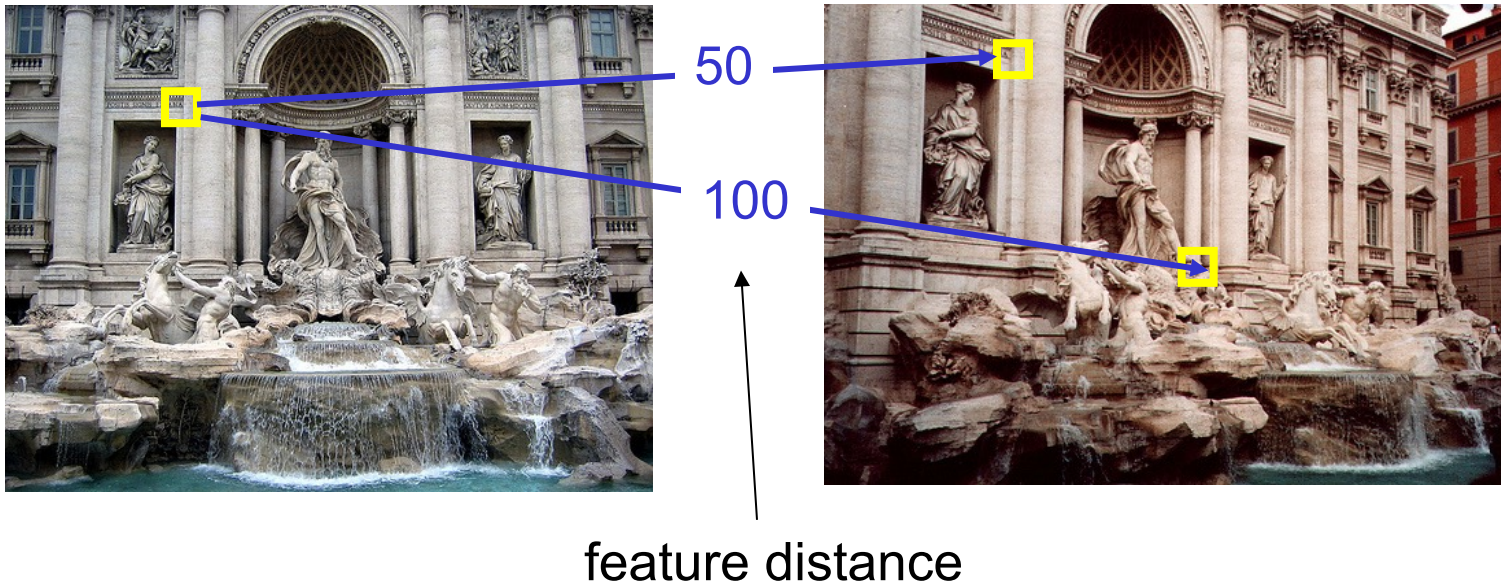
Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

Feature distance

How to define the distance function between features f_1 , f_2 ?


- Simple approach is $\text{SSD}(f_1, f_2)$
 - sum of square differences between entries of the two descriptors




Feature distance

How to define the distance function between features f_1 , f_2 ?

- Simple approach is $\text{SSD}(f_1, f_2)$
 - sum of square differences between entries of the two descriptors



I_1




I_2


Feature distance

How to define the distance function between features f_1 , f_2 ?

- Simple approach is $\text{SSD}(f_1, f_2)$
 - sum of square differences between entries of the two descriptors
 - can give good scores to ambiguous (bad) matches



I_1




I_2


Feature distance

How to define the distance function between features f_1 , f_2 ?

- Better approach: ratio distance = $\text{SSD}(f_1, f_2) / \text{SSD}(f_1, f_2')$
 - f_2 is best SSD match to f_1 in I_2
 - f_2' is 2nd best SSD match to f_1 in I_2
 - gives small values for unambiguous matches




I_1




I_2

Evaluating the results

How can we measure the performance of a feature matcher?



True or false positives




Positive match if $SSD < \text{distance threshold}$

The distance threshold affects performance

- True positives = number of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- False positives = number of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

True or false positives



Positive match if $SSD < \text{distance threshold}$

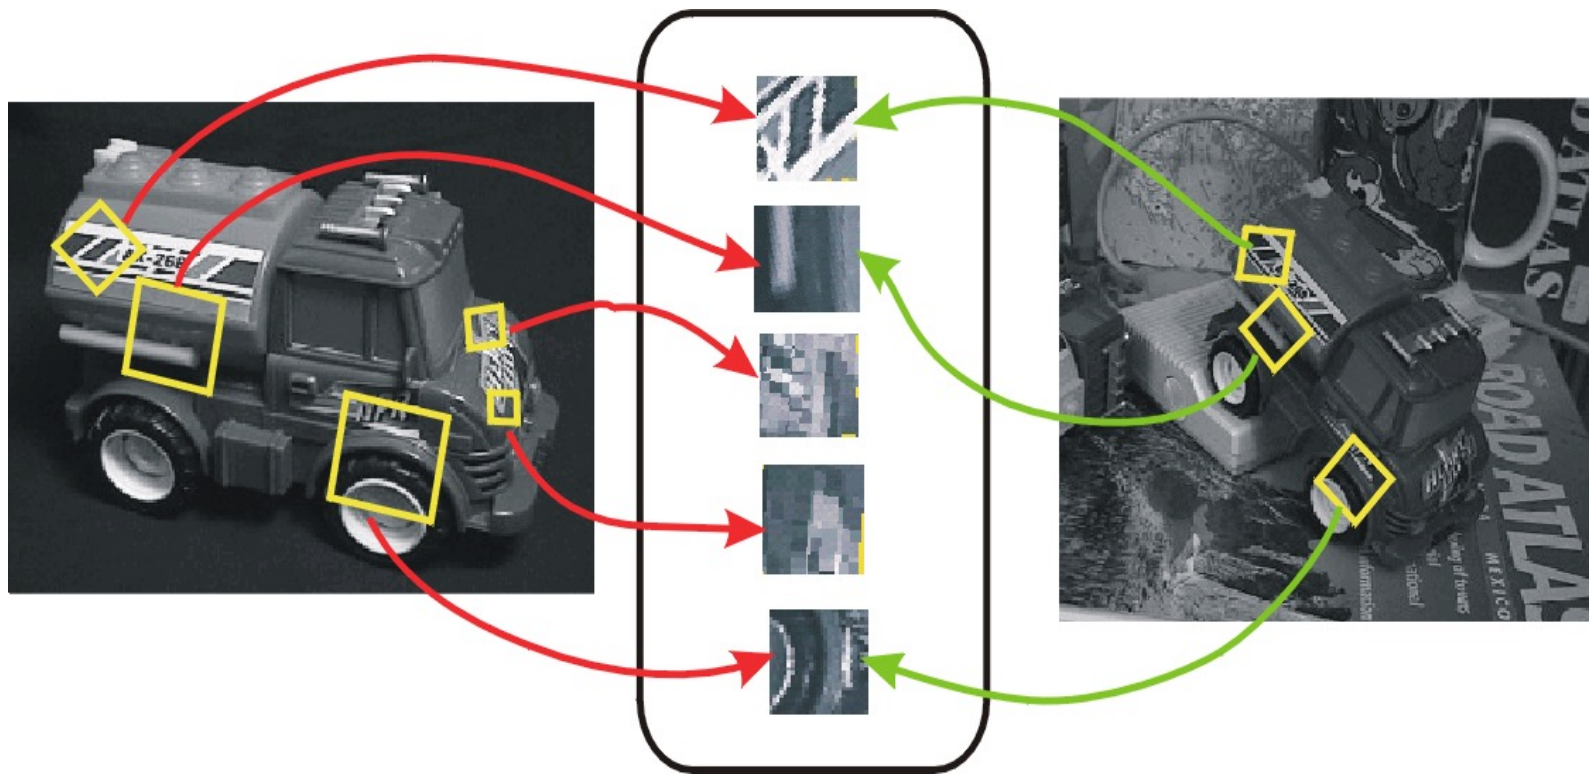
The distance threshold affects performance

- True positives = number of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
 - Increase threshold (uncertain matches are also allowed)
- False positives = number of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?
 - Decrease threshold (matches discarded unless they are very certain)

Desirable property: invariance

Find features that are invariant to transformations across two images

- geometric invariance: translation, rotation, scale
- photometric invariance: brightness, exposure, ...



Feature Descriptors

CSE 152A, W124: Manmohan Chandraker

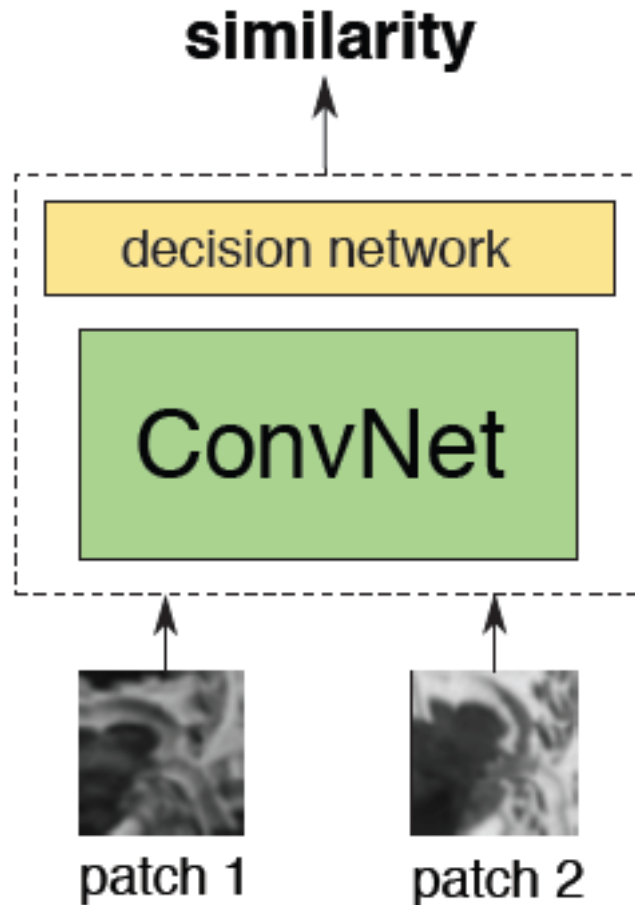
SIFT: Popular for Feature Matching

Extraordinarily robust matching technique

- Can handle changes in viewpoint
 - Up to about 30 degrees out of plane rotation
- Can handle significant changes in illumination
 - Sometimes even day vs. night (below)
- Fast and efficient—can run in real time
- Lots of code available
 - http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT



Learning correspondence





Idea:


- Siamese network to decide patch similarity
- Use intermediate activations as features.

[Zagoruyko and Komodakis, CVPR 2015]

Correspondence is a vital 3D cue



Depth from correspondence




Two measurements: X_L, X_R
Two unknowns: X_0, Z_0

Constants:

Baseline: d

Focal length: f

Depth from correspondence



Using similar triangles:

$$\frac{X_L}{f} = \frac{X_0}{Z_0} \quad \frac{X_R}{f} = \frac{X_0 - d}{Z_0}$$

Two measurements: X_L, X_R
Two unknowns: X_0, Z_0

Constants:

Baseline: d

Focal length: f

$$X_0 = \frac{d X_L}{(X_L - X_R)}$$

$$Z_0 = \frac{d f}{(X_L - X_R)}$$

Disparity: $(X_L - X_R)$

Depth is inversely proportional to disparity

Mars Exploratory Rovers: Spirit and Opportunity, 2004

Stereo camera

